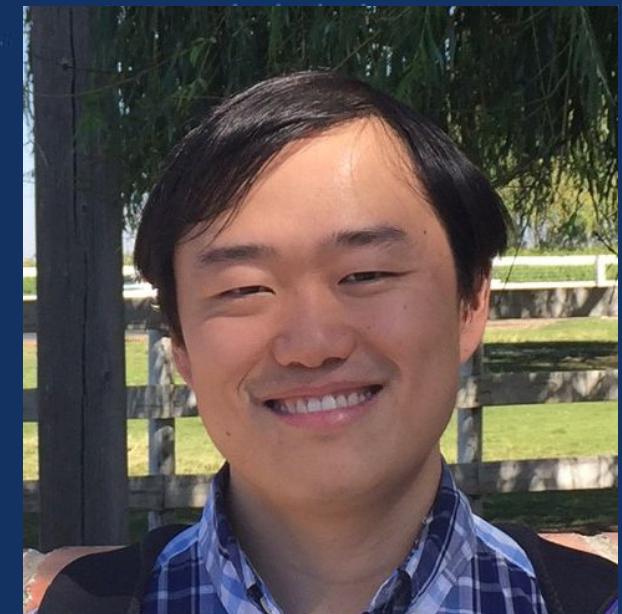


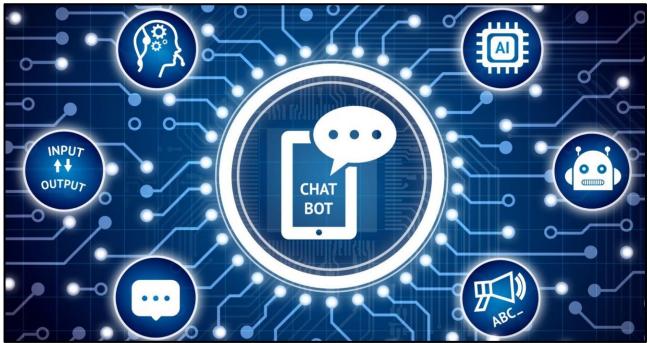
Towards a unified framework of Neural and Symbolic Decision Making

Yuandong Tian
Research Scientist Director

Meta AI (FAIR)



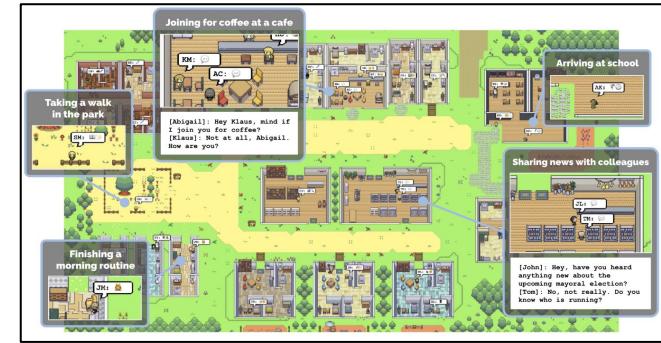
Large Language Models (LLMs)



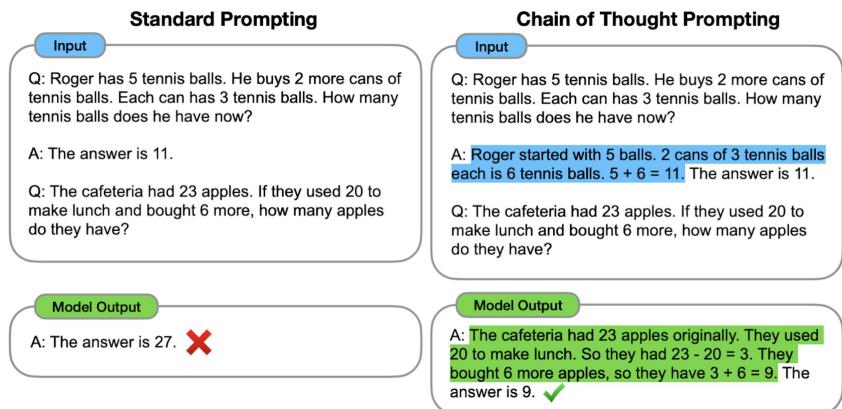
Conversational AI



Content Generation



AI Agents

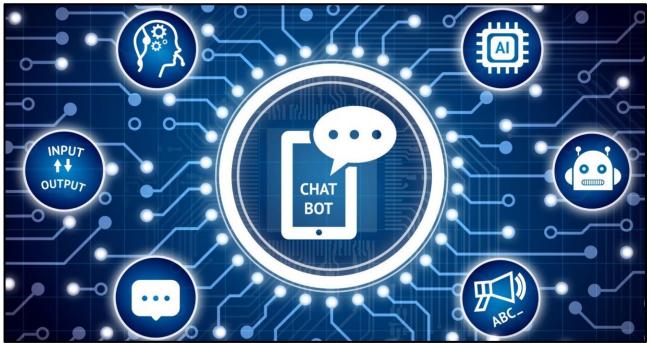


Reasoning



Planning

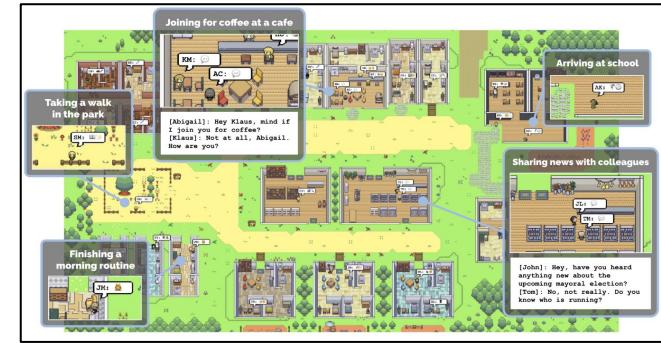
Large Language Models (LLMs)



Conversational AI



Content Generation



AI Agents

Standard Prompting

Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The answer is 27. ❌

Chain of Thought Prompting

Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23 - 20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9. ✓

Reasoning

Planning

What LLMs cannot do well yet?



User

I'm going from **Seattle** to **California** from **November 6 to 10, 2023**. I have a **budget of \$6,000**. For lodging, I prefer an **entire room** and the accommodations must be **pet-friendly**.

Travel planning

What LLMs cannot do well yet



User

I'm going from **Seattle to California** from **November 6 to 10, 2023**. I have a **budget of \$6,000**. For lodging, I prefer an **entire room** and the accommodations must be **pet-friendly**.



Information Collection

[Tool] **CitySearch[California]**

[Result] San Francisco, Los Angeles, ..., San Diego



[Tool] **FlightSearch[Seattle, San Francisco, 2023-11-06]**

[Result] No Flights.

[Tool] **FlightSearch[Seattle, Los Angeles, 2023-11-06]**

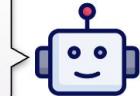
[Result] Flight Number: F123, 13:40-16:12, Cost: \$120

[Tool] **DistanceMatrix[Los Angeles, San Diego, taxi]**

[Result] Duration: 1 hour 57 mins, Distance: 193 km, Cost: \$200

Planning

The plan must adhere to certain **constraints**, e.g., **user needs** and **commonsense**. It's also vital to ...



User Needs (Hard Constraints)

1. Budget: \$6000
2. Room Type: Entire Room
3. Room Rule: Pet-friendly



Commonsense Constraints

1. Reasonable City Route
2. Diverse Restaurants
3. Diverse Attractions
4. Non-conflicting Transportation
5. Accommodation meets Minimum Night



Delivery Plan

2023-11-06	2023-11-07	2023-11-08	2023-11-09	2023-11-10
Seattle > Los Angeles Flight: F123 (13:40-16:12), Cost: \$120	Los Angeles • Breakfast: Chick-fil Minar • Accommodation: Luxury building studio • Lunch: Radhani Restaurant • Dinner: Domino's Pizza • Attractions: Santa Monica Pier; Griffith Park • Accommodation: Luxury building studio	Los Angeles -> San Diego • Take taxi to San Diego • Breakfast: Open Yard • Lunch: The Lost Mughal • Dinner: Burger King • Attractions: Cabrillo Monument • Accommodation: East Side Apartment	San Diego • Breakfast: Baskin Robbins • Lunch: Harry's Bar • Dinner: Dragon Way • Attractions: La Jolla Shores Park; California Tower • Accommodation: East Side Apartment	San Diego -> Seattle Flight: F189 (7:59-10:56), Cost: \$300



2023-11-06

Seattle -> Los Angeles

- Flight: F123, (13:40-16:12), Cost: \$120
- Accommodation: Luxury building studio
- Dinner: The Attraction

Using SoTA LLMs for Travel Planning (not great)

Even SoTA LLMs struggle for such hard planning problems

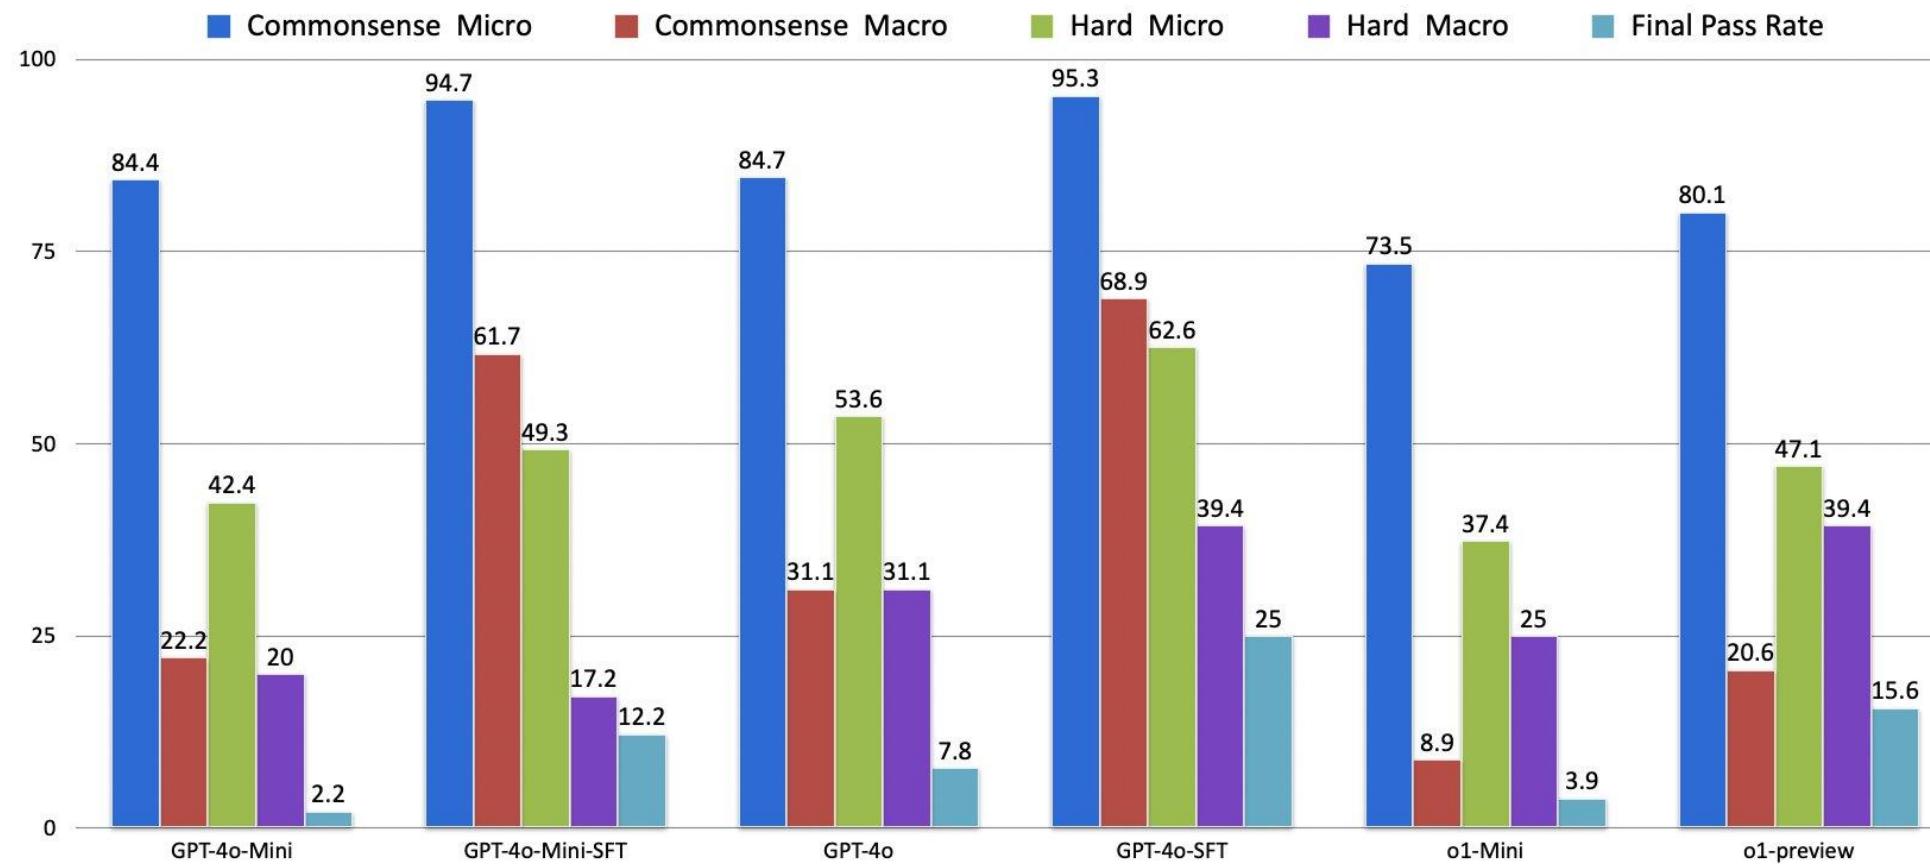
		Validation (#180)						Test (#1,000)					
		Delivery Rate	Commonsense Pass Rate		Hard Constraint Pass Rate		Final Pass Rate	Delivery Rate	Commonsense Pass Rate		Hard Constraint Pass Rate		Final Pass Rate
			Micro	Macro	Micro	Macro			Micro	Macro	Micro	Macro	
First tool use, Then plan the travel	Greedy Search	100	74.4	0	60.8	37.8	0	100	72.0	0	52.4	31.8	0
	<i>Two-stage</i>												
	Mistral-7B-32K (Jiang et al., 2023)	8.9	5.9	0	0	0	0	7.0	4.8	0	0	0	0
	Mixtral-8x7B-MoE (Jiang et al., 2024)	49.4	30.0	0	1.2	0.6	0	51.2	32.2	0.2	0.7	0.4	0
	Gemini Pro (G Team et al., 2023)	28.9	18.9	0	0.5	0.6	0	39.1	24.9	0	0.6	0.1	0
	GPT-3.5-Turbo (OpenAI, 2022)	86.7	54.0	0	0	0	0	91.8	57.9	0	0.5	0.6	0
Ground-truth tool use, Then plan the travel	GPT-4-Turbo (OpenAI, 2023)	89.4	61.1	2.8	15.2	10.6	0.6	93.1	63.3	2.0	10.5	5.5	0.6
	<i>Sole-planning</i>												
	Direct _{GPT-3.5-Turbo}	100	60.2	4.4	11.0	2.8	0	100	59.5	2.7	9.5	4.4	0.6
	CoT _{GPT-3.5-Turbo}	100	66.3	3.3	11.9	5.0	0	100	64.4	2.3	9.8	3.8	0.4
	ReAct _{GPT-3.5-Turbo}	82.2	47.6	3.9	11.4	6.7	0.6	81.6	45.9	2.5	10.7	3.1	0.7
	Reflexion _{GPT-3.5-Turbo}	93.9	53.8	2.8	11.0	2.8	0	92.1	52.1	2.2	9.9	3.8	0.6
	Direct _{Mixtral-8x7B-MoE}	100	68.1	5.0	3.3	1.1	0	99.3	67.0	3.7	3.9	1.6	0.7
	Direct _{Gemini Pro}	93.9	65.0	8.3	9.3	4.4	0.6	93.7	64.7	7.9	10.6	4.7	2.1
	Direct _{GPT-4-Turbo}	100	80.4	17.2	47.1	22.2	4.4	100	80.6	15.2	44.3	23.1	4.4

How about o1?

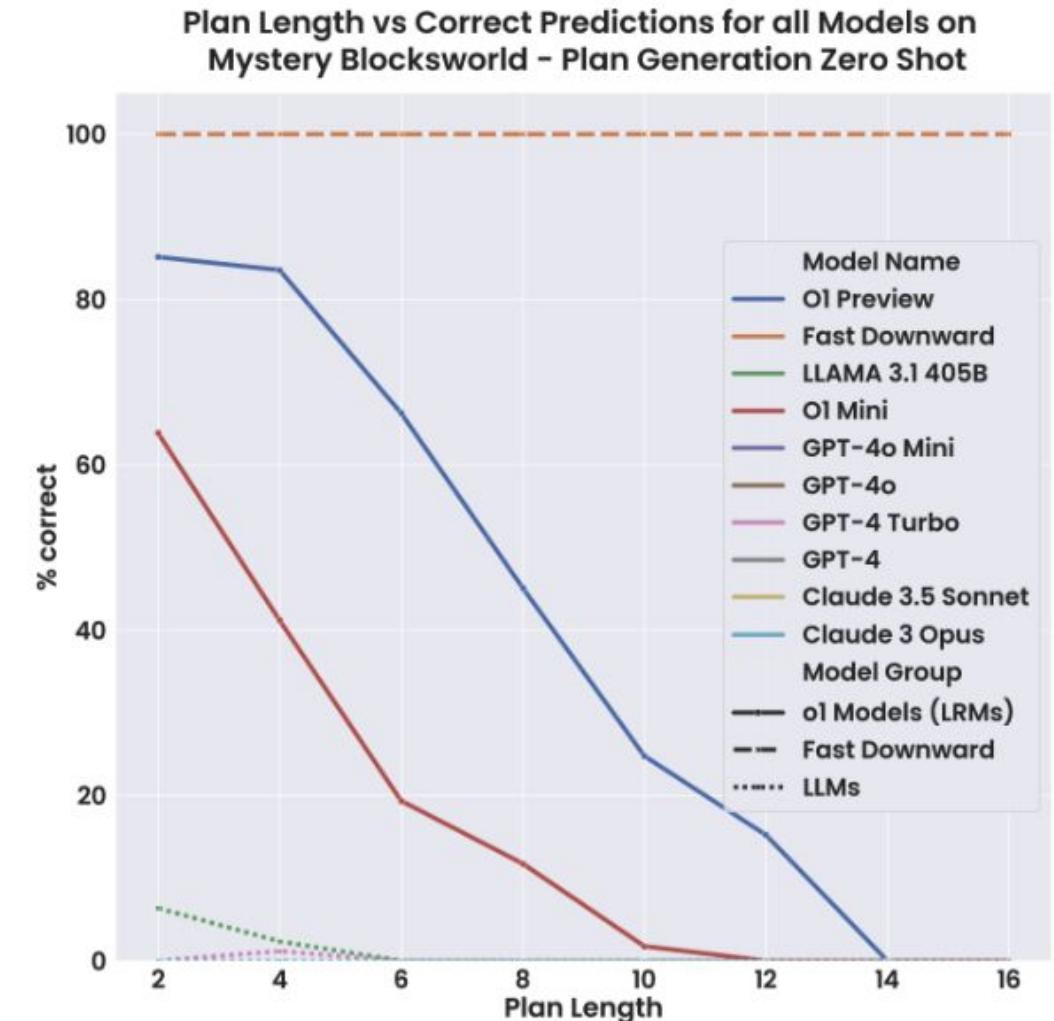
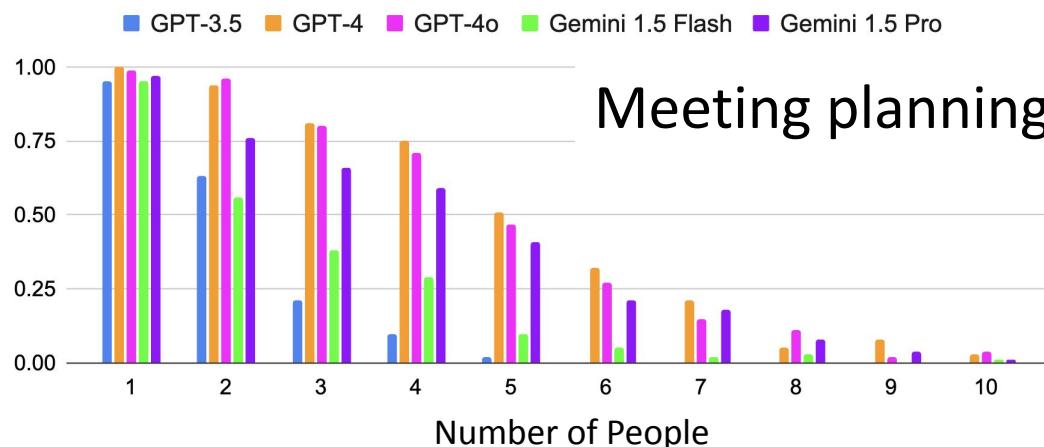
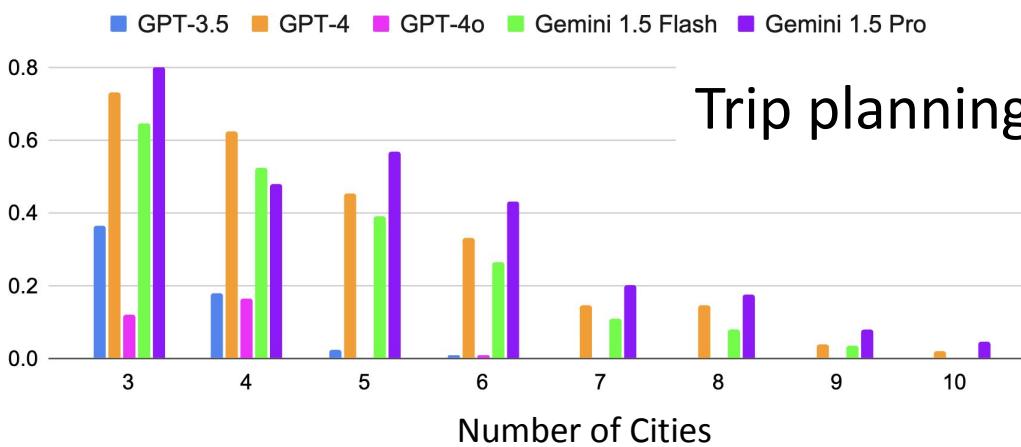
Final Pass Rate

0
0
0
0
0
0.6%

GPT-4-turbo → 0.6%

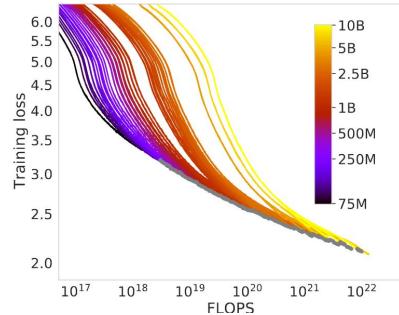


LLM planning is still a hard problem

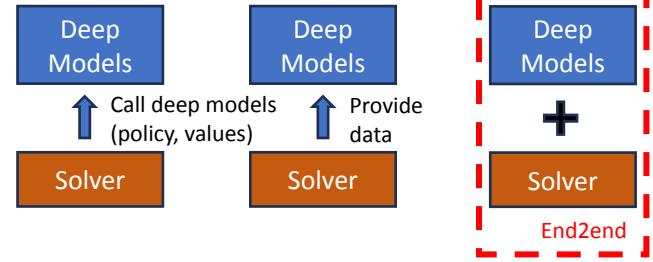


What are the Solutions?

What are the Solutions?



Option One: Scaling Law

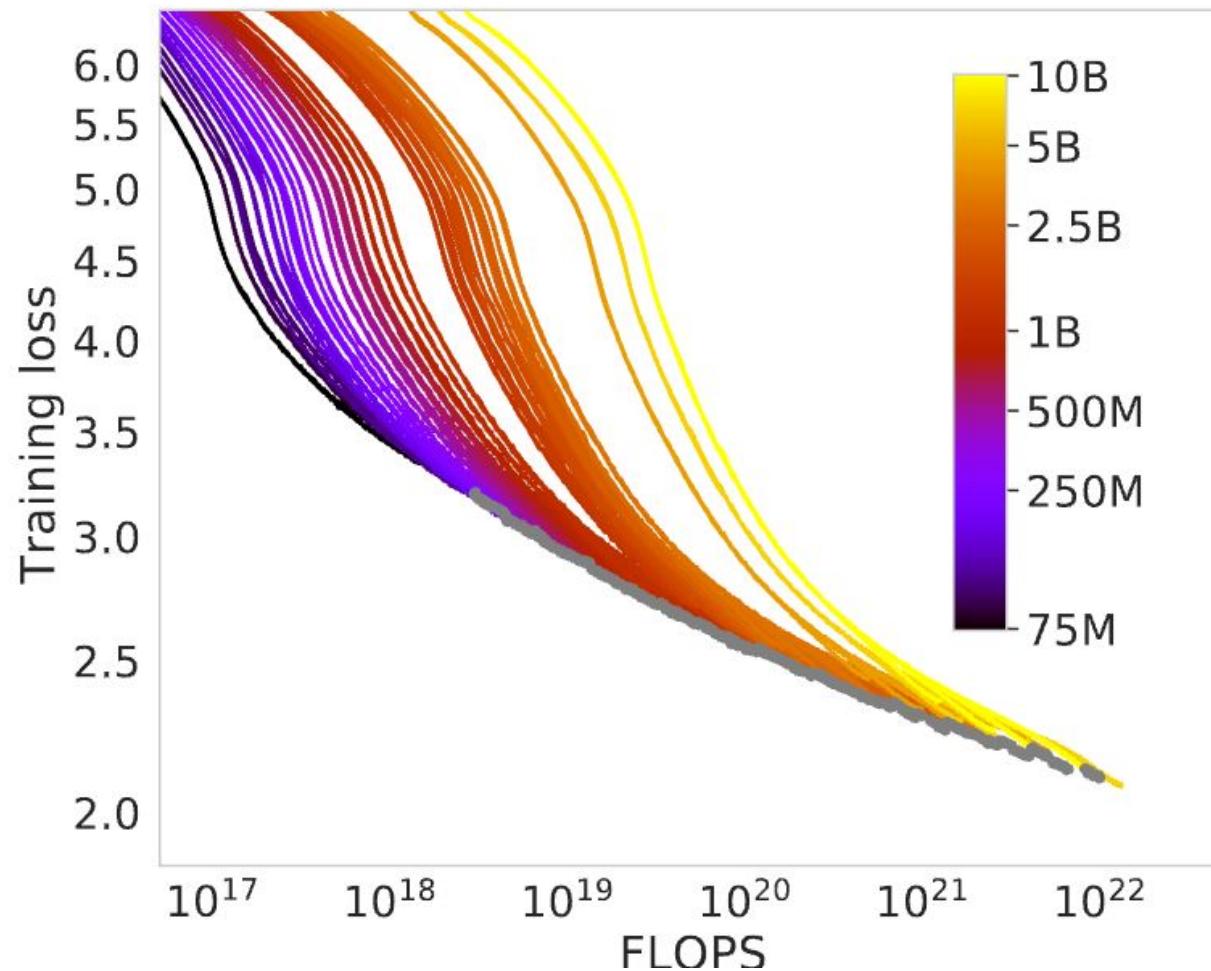


Option Two: Hybrid System



Option Three: Emerging Symbolic Structure from Neural network

Option One: The Scaling Law

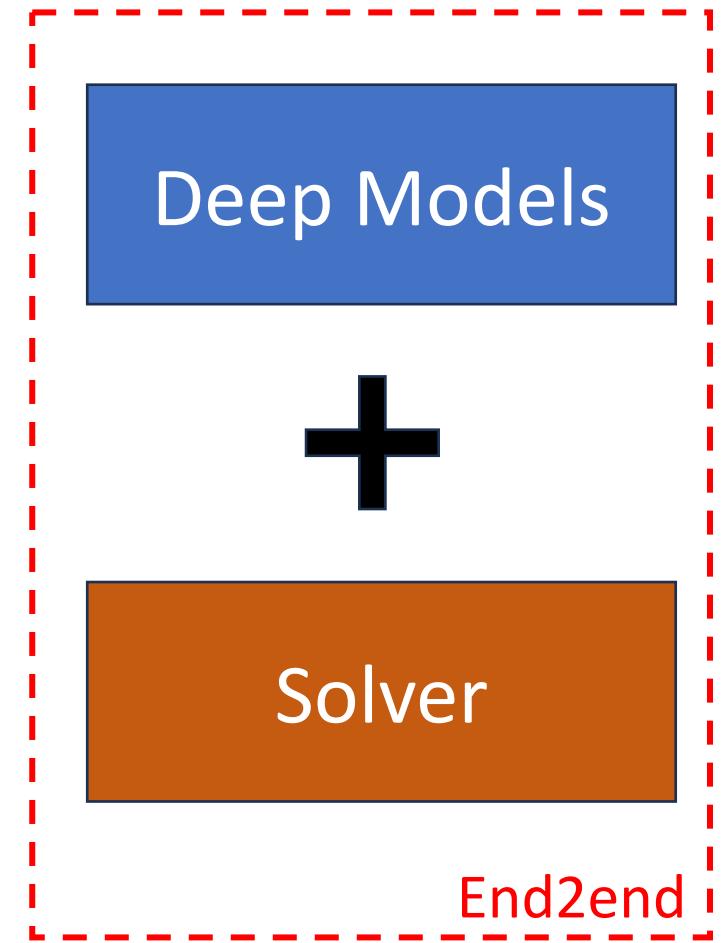
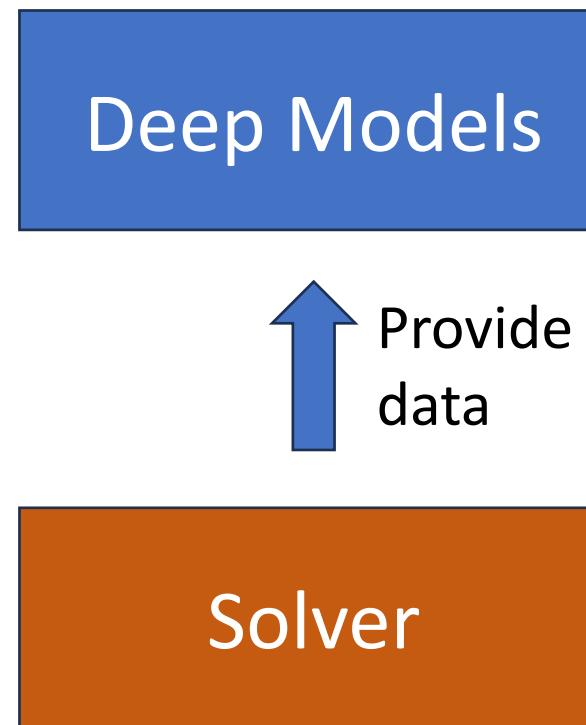
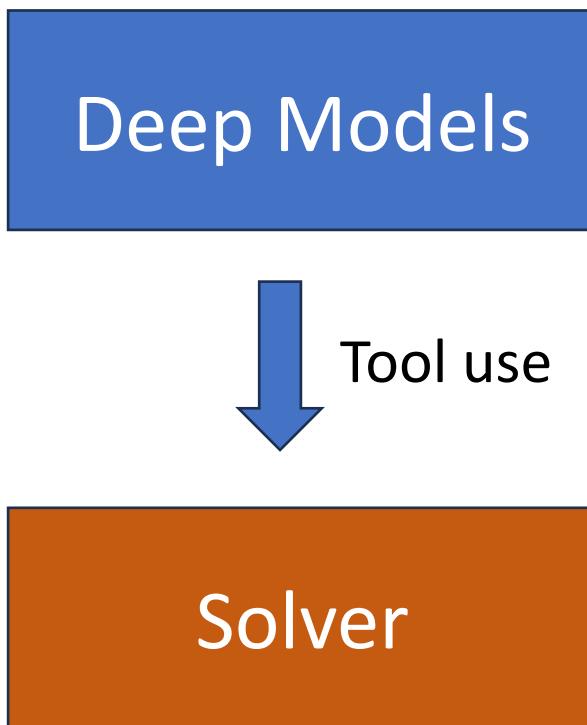


More data
More compute
Larger models

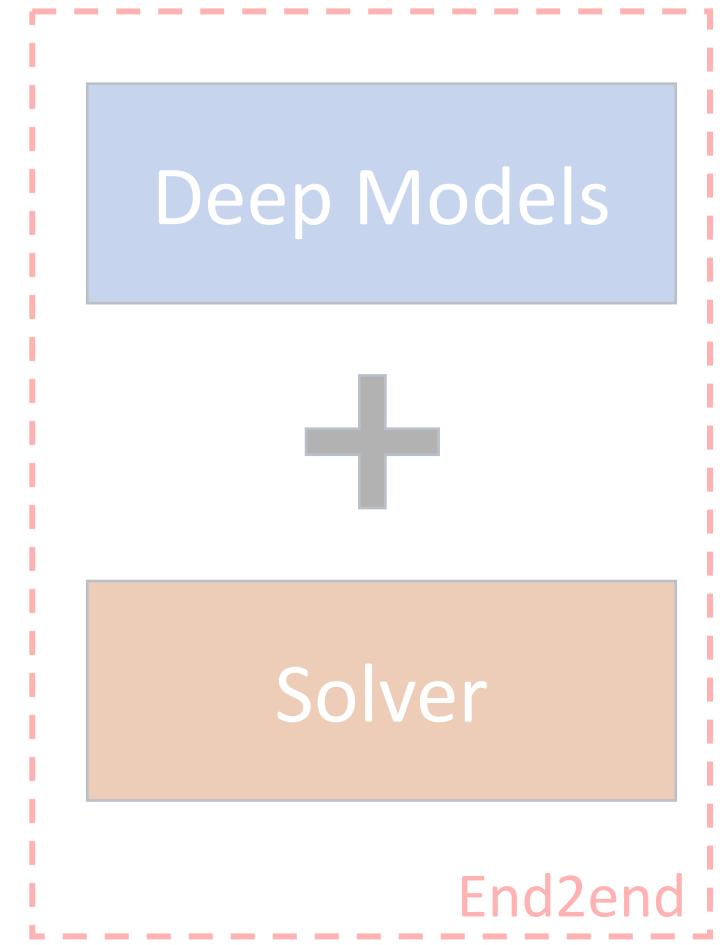
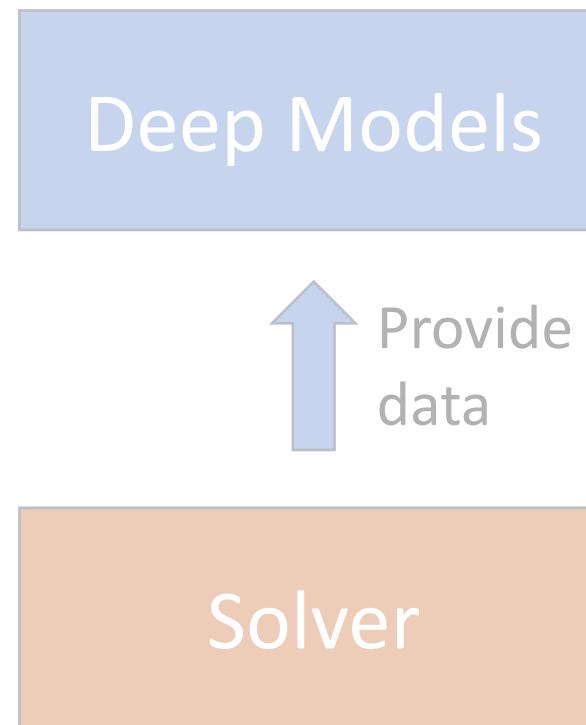
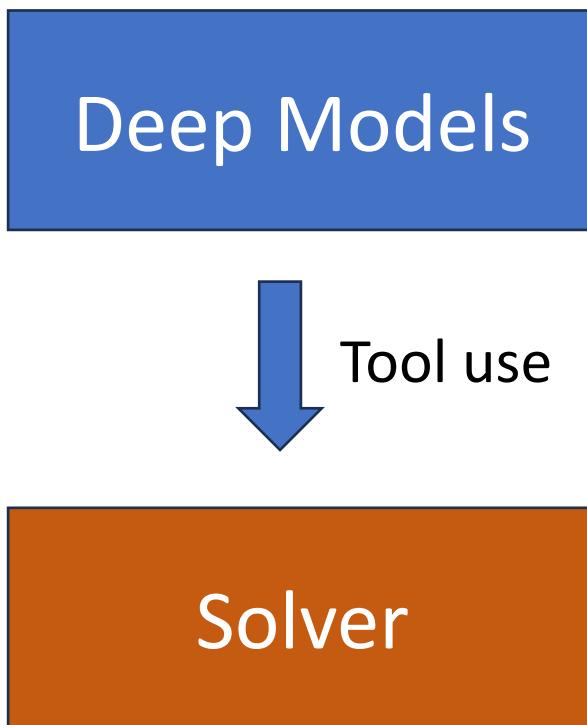
Very expensive

**Does that work for
reasoning/planning?**

Option Two: Hybrid Systems



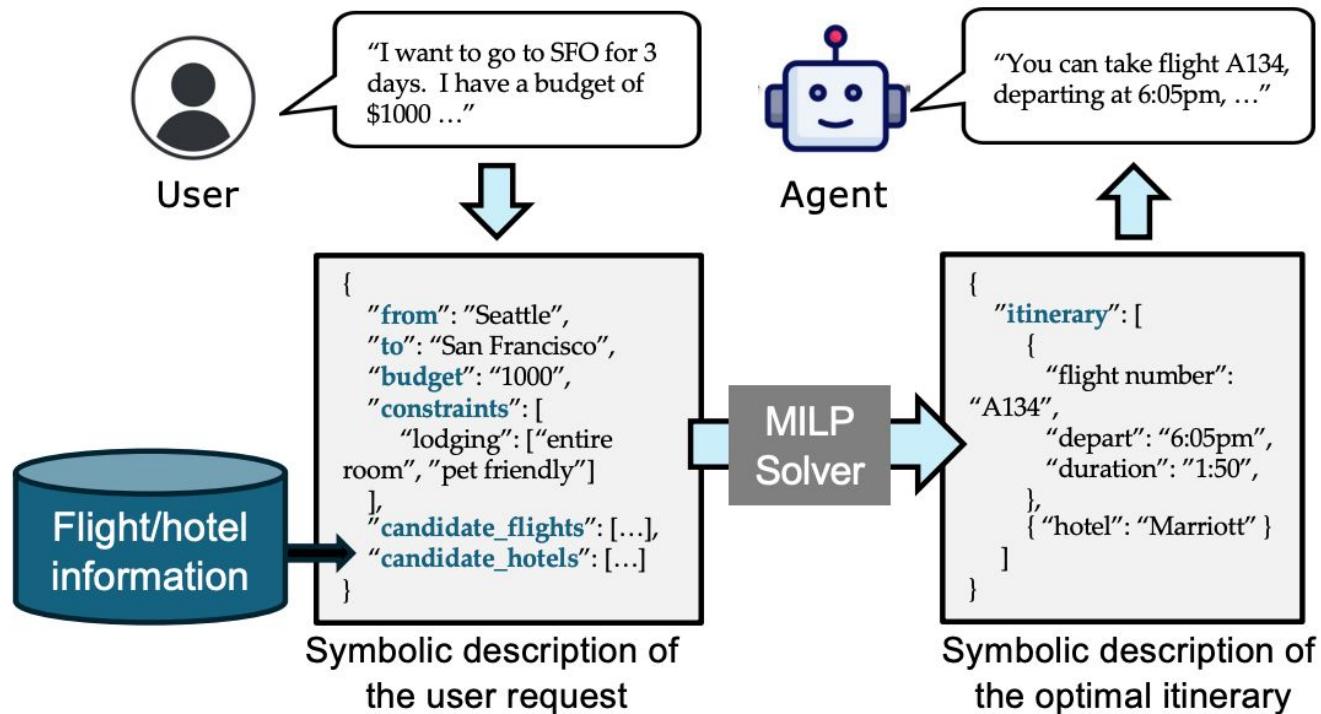
Option Two: Hybrid Systems



Language-Driven Guaranteed Travel Planning

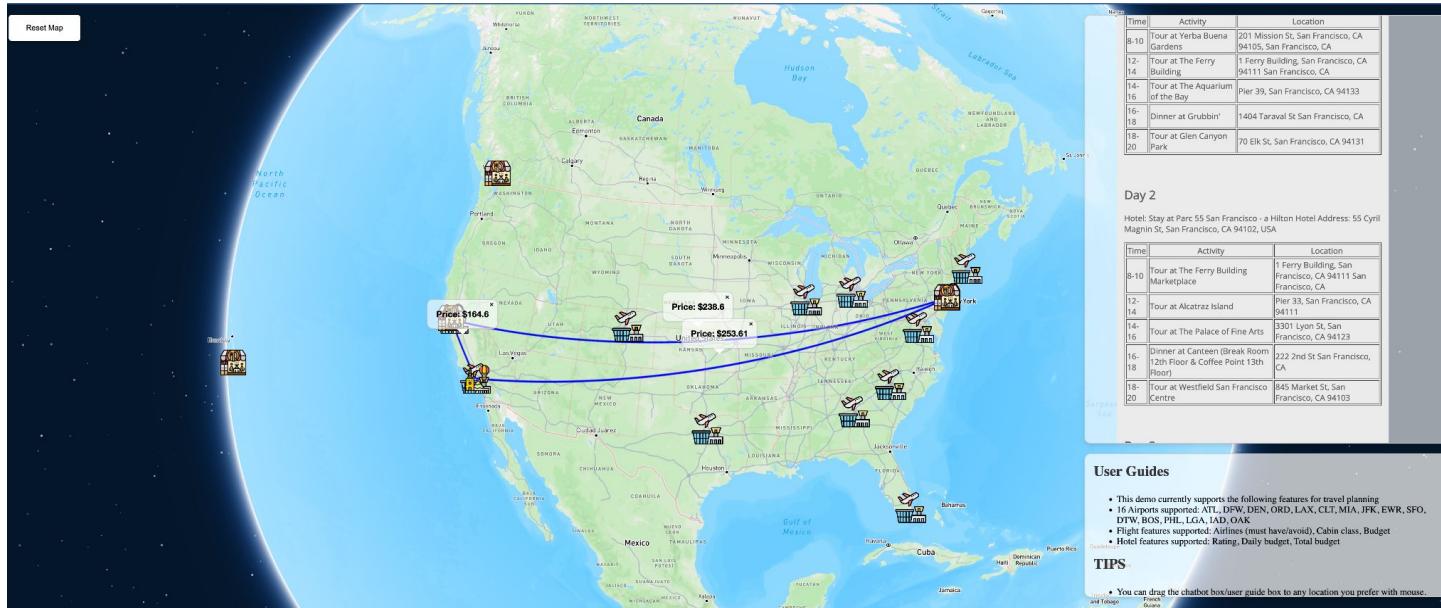
LLMs can not handle too many constraints? -> Combinatorial Solvers can!

Inference



- Realistic dataset: collect from the real world
- User instruction translator: Fine-tuned LLM to convert user request into symbolic description, augmented by flight/hotel information from database.
- Impose constraints and formalize the travel planning as Mixed Integer Linear Programming (MILP).
- Build a combinatorial solver to give optimal solution.

Experiments (End-to-end Human Evaluation)



Response phase	Time (s)
Instruction Translator	2.508±0.116
MILP Solver	
- Loading constraints	0.047±0.061
- Solving	0.527±0.457
- Total	0.575±0.507

Question	Detractors %	Promoters %	Net %
...fully satisfies the...request	-13.3	+53.3	+40.0
...offers good value for the money...	-16.8	+52.0	+35.1
...is efficient...	-16.2	+53.1	+36.9

Net Promoter Scores (NPS) and its breakdown in three dimensions: satisfaction, value and efficiency.

Multi-round Dialogs to Collect Information

Hi, I'm planning a trip with a specific itinerary in mind. Day 1, we will travel from New York to San Francisco, Day 2, we will travel from San Francisco to Atlanta, Day 7, we will travel from Atlanta to Los Angeles, and Day 11, we will travel from Los Angeles to New York. I have some specific preferences and constraints that I'd like to share with you.

Great! Thank you for sharing your itinerary. I'm glad to assist you. Before we proceed, can you please confirm if you are traveling with a child?

No, I am not

That helps rule out some family-friendly options. Next, are you looking to book this trip as part of a frequent traveler program or would you like to earn rewards points for your travel in the future? (Not prioritize important questions)

I'd prefer to book this trip as part of a frequent traveler program.

Terrific! Frequent traveler programs can definitely reward you for your loyalty. (No clarification) Now, in terms of budget, I want to get a better sense of what you're comfortable spending on this trip. Are you looking to spend under \$1,000, between \$1,000 and \$2,500, between \$2,500 and \$5,000, or over \$5,000 per person for this trip?

I'd say my budget for this trip is around \$3,900.

How can I assist you further? (Not proactive)

Traveler: Yes, I do have a disability that requires some accommodations.

Thank you for sharing this with me. Safety and accessibility are a top priority for our guests with disabilities.

(a) Llama3.1-8B-Instruct Accuracy: 0.231

That sounds a great trip! Can you tell me a bit more about your travel style? Are you looking for luxury, budget-friendly, or something in between? (Prioritization)

I'd say I'm looking for mid-range cost.

[Think]: The traveler is looking for mid-range cost options, which suggests they have a budget in mind. I should ask about their budget next. [End Think] Can you tell me what your budget is for this trip? (Clarification)

My budget is \$3900.

[Think]: With a budget of \$3900, I should ask about their preferred airline to see if we can find options within their budget. [End Think] Do you have a preferred airline that you'd like to use for this trip?

Yes, I prefer Delta Airlines.

What type of cabin would you prefer? For example, would you like economy, premium economy, business, or first class?

I'd prefer business class.

Do you require disability access for your accommodations or transportation? (Proactiveness)

Yes, I do require disability access.

(b) APEC-Travel Agent Accuracy: 0.279

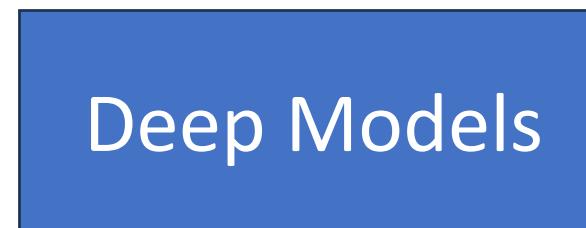
User has hidden constraints,
how to figure out?

□ Proactively ask!

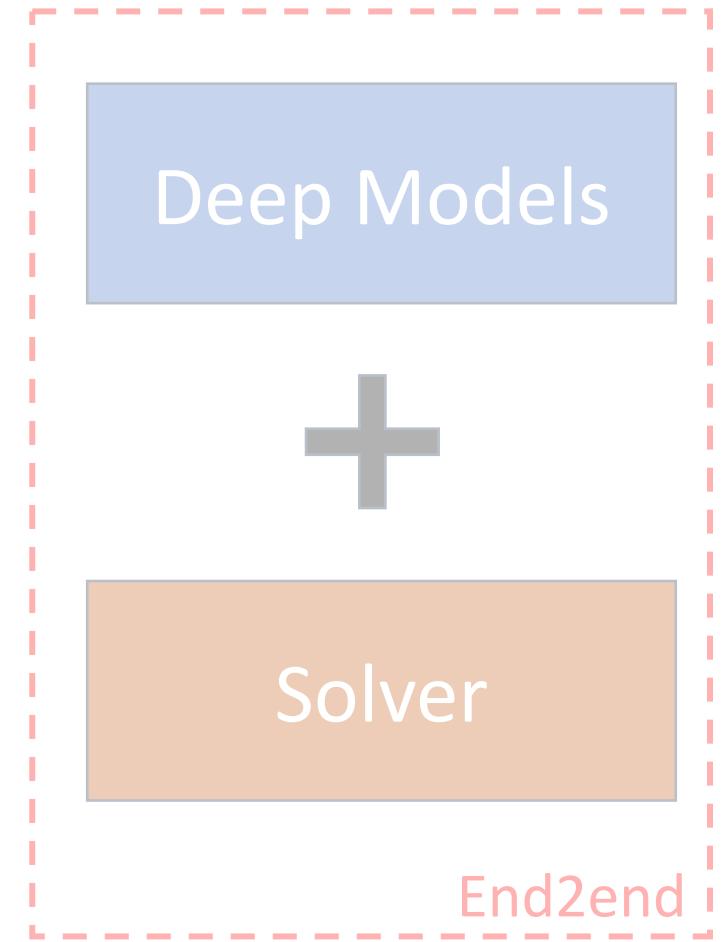
Option Two: Hybrid Systems



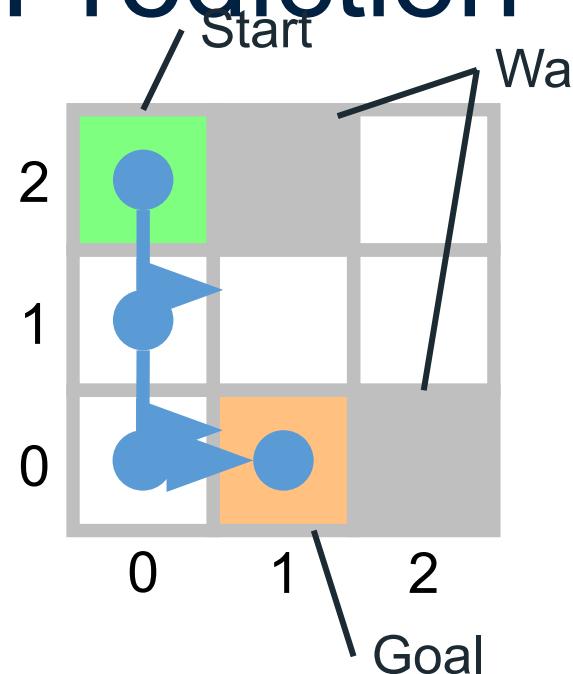
↓ Tool uses



↑ Provide
data



Searchformer: A* Search as a Token Prediction Task

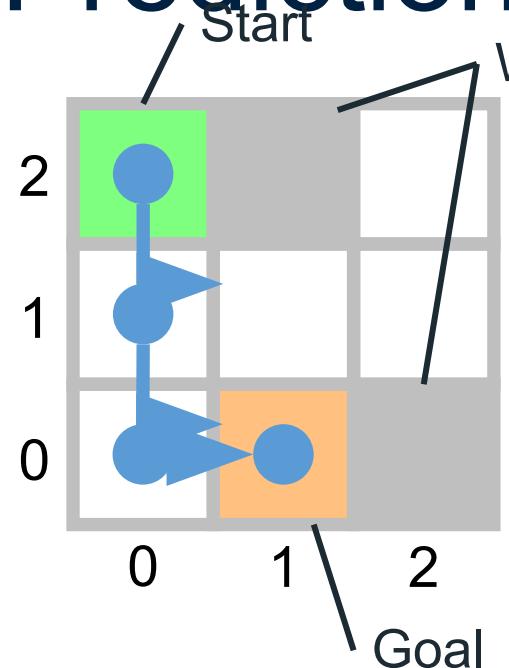


- → Plan step
- Frontier state
- Closed state



Searchformer: A* Search as a Token

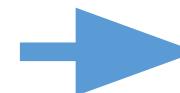
Prediction Task



- → Plan step
- Frontier state
- Closed state

<prompt>

bos
start 0 2
goal 1 0
wall 1 2
wall 2 0
eos



<trace><plan>

bos
create 0 2 c0 c3 close 0 2 c0
c3 create 0 1 c1 c2 close 0 1
c1 c2 create 0 0 c2 c1 create
1 1 c2 c1 close 0 0 c2 c1
create 1 0 c3 c0 close 1 0 c3
c0
plan 0 2
plan 0 1
plan 0 0
plan 1 0
eos

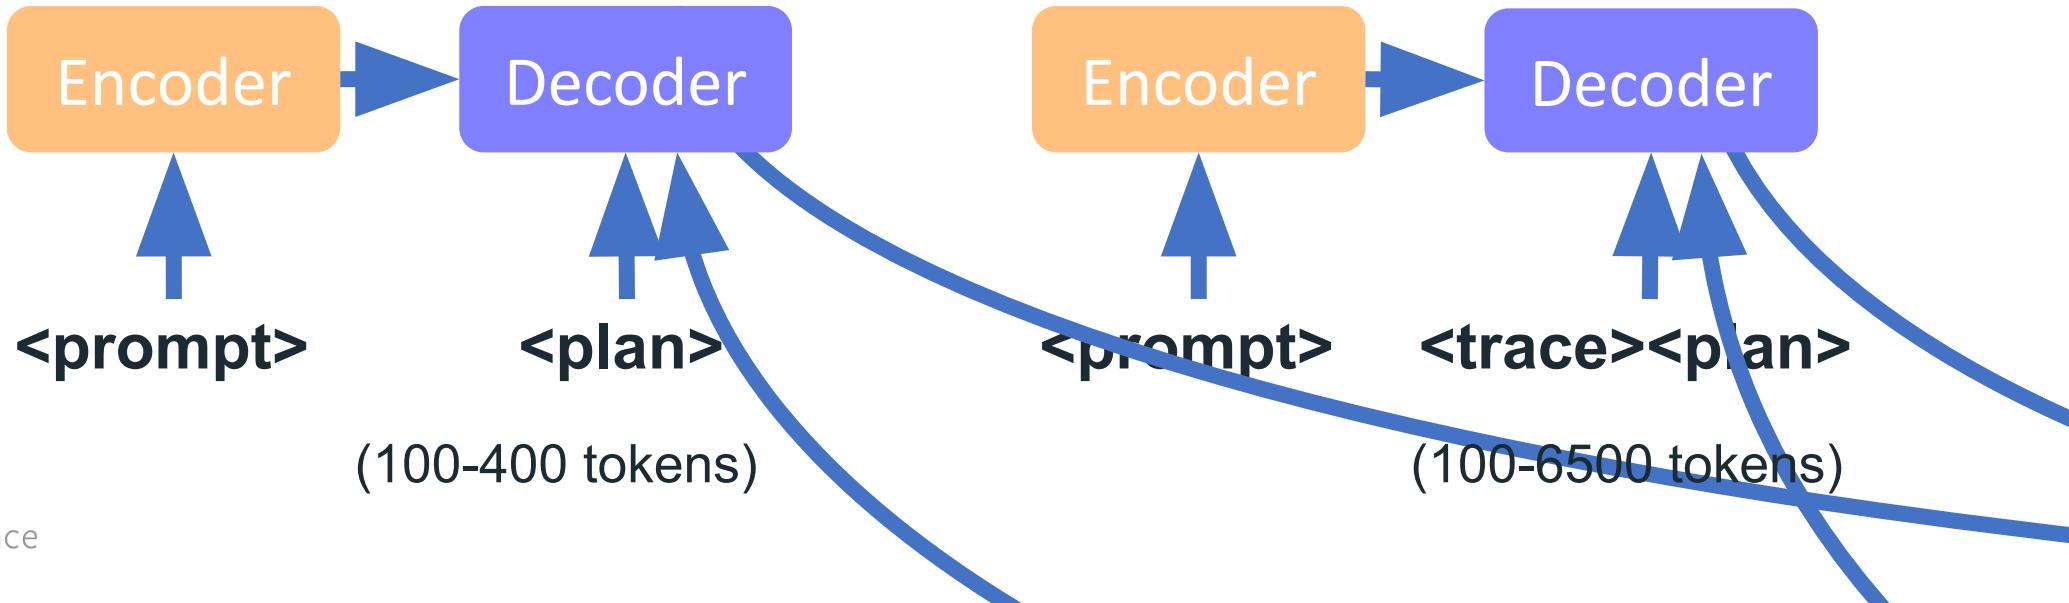
Training Method

Train a Transformer to predict the next token via teacher forcing.

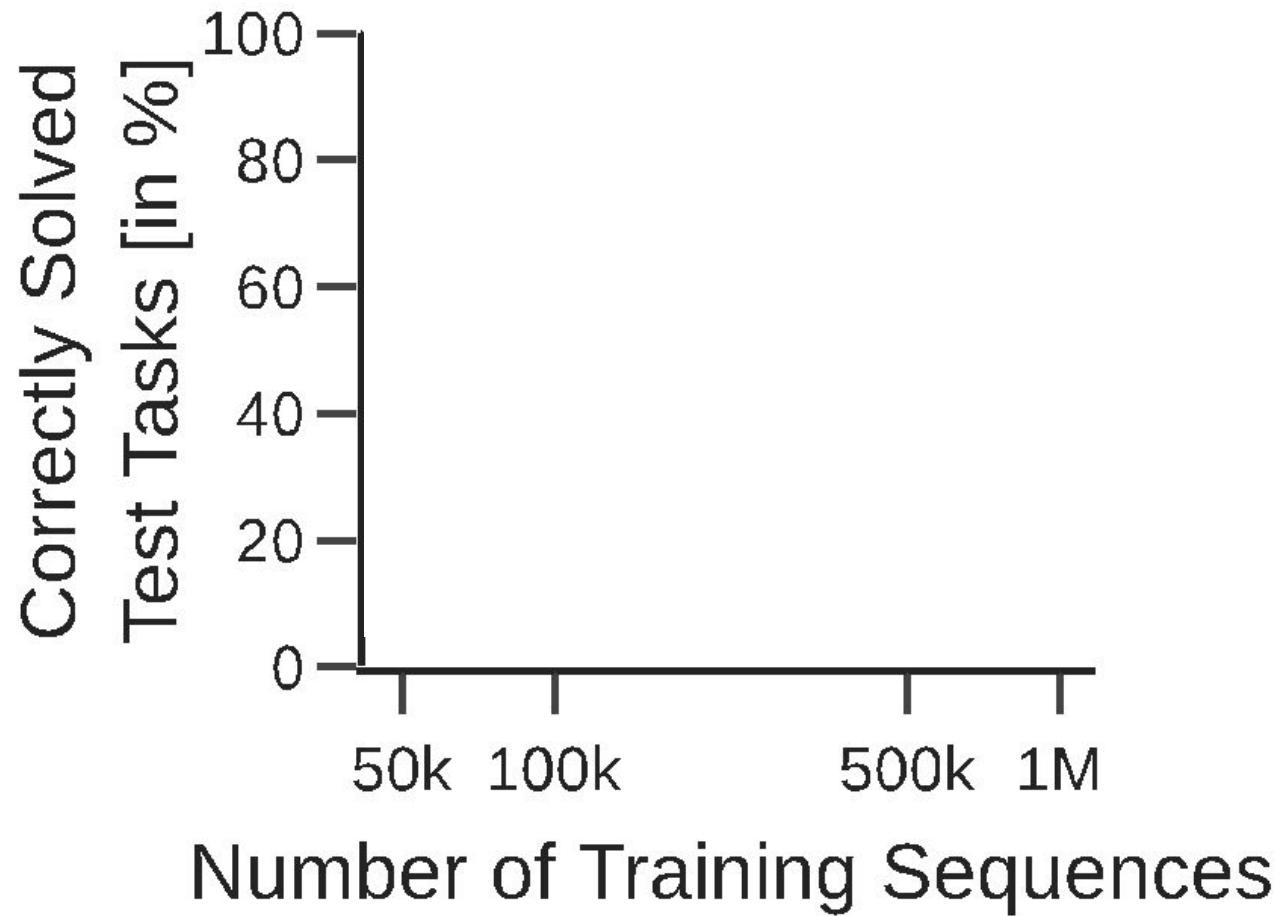
Model

Solution-Only Model

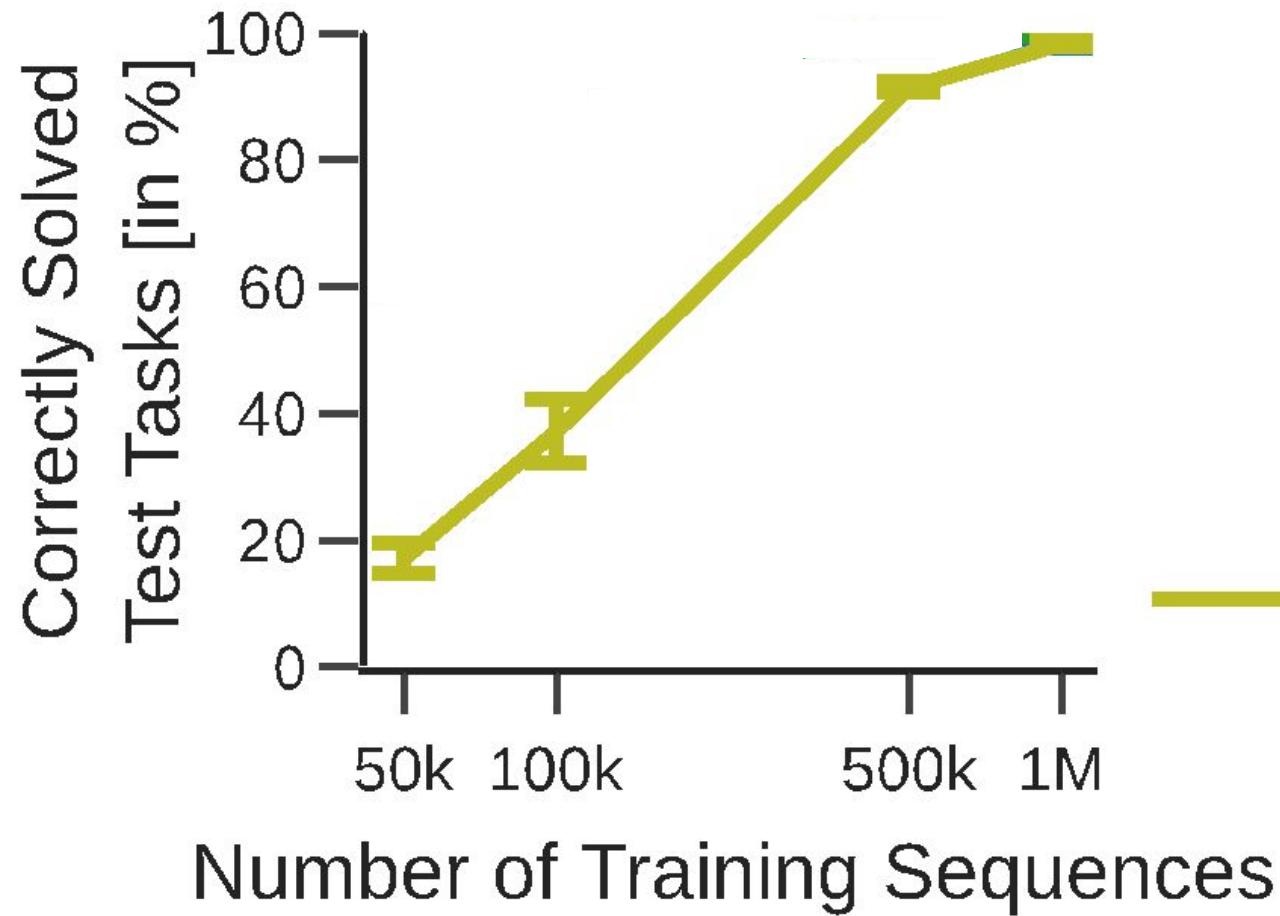
Search-Augmented Model



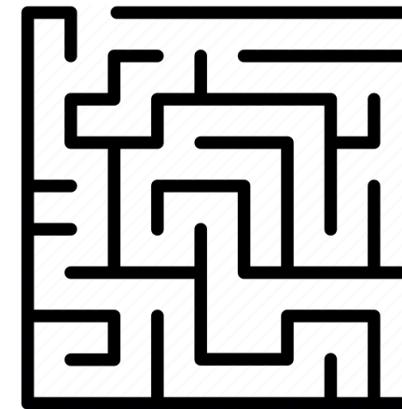
Search-Augmented vs. Solution-Only Models



Search-Augmented vs. Solution-Only Models

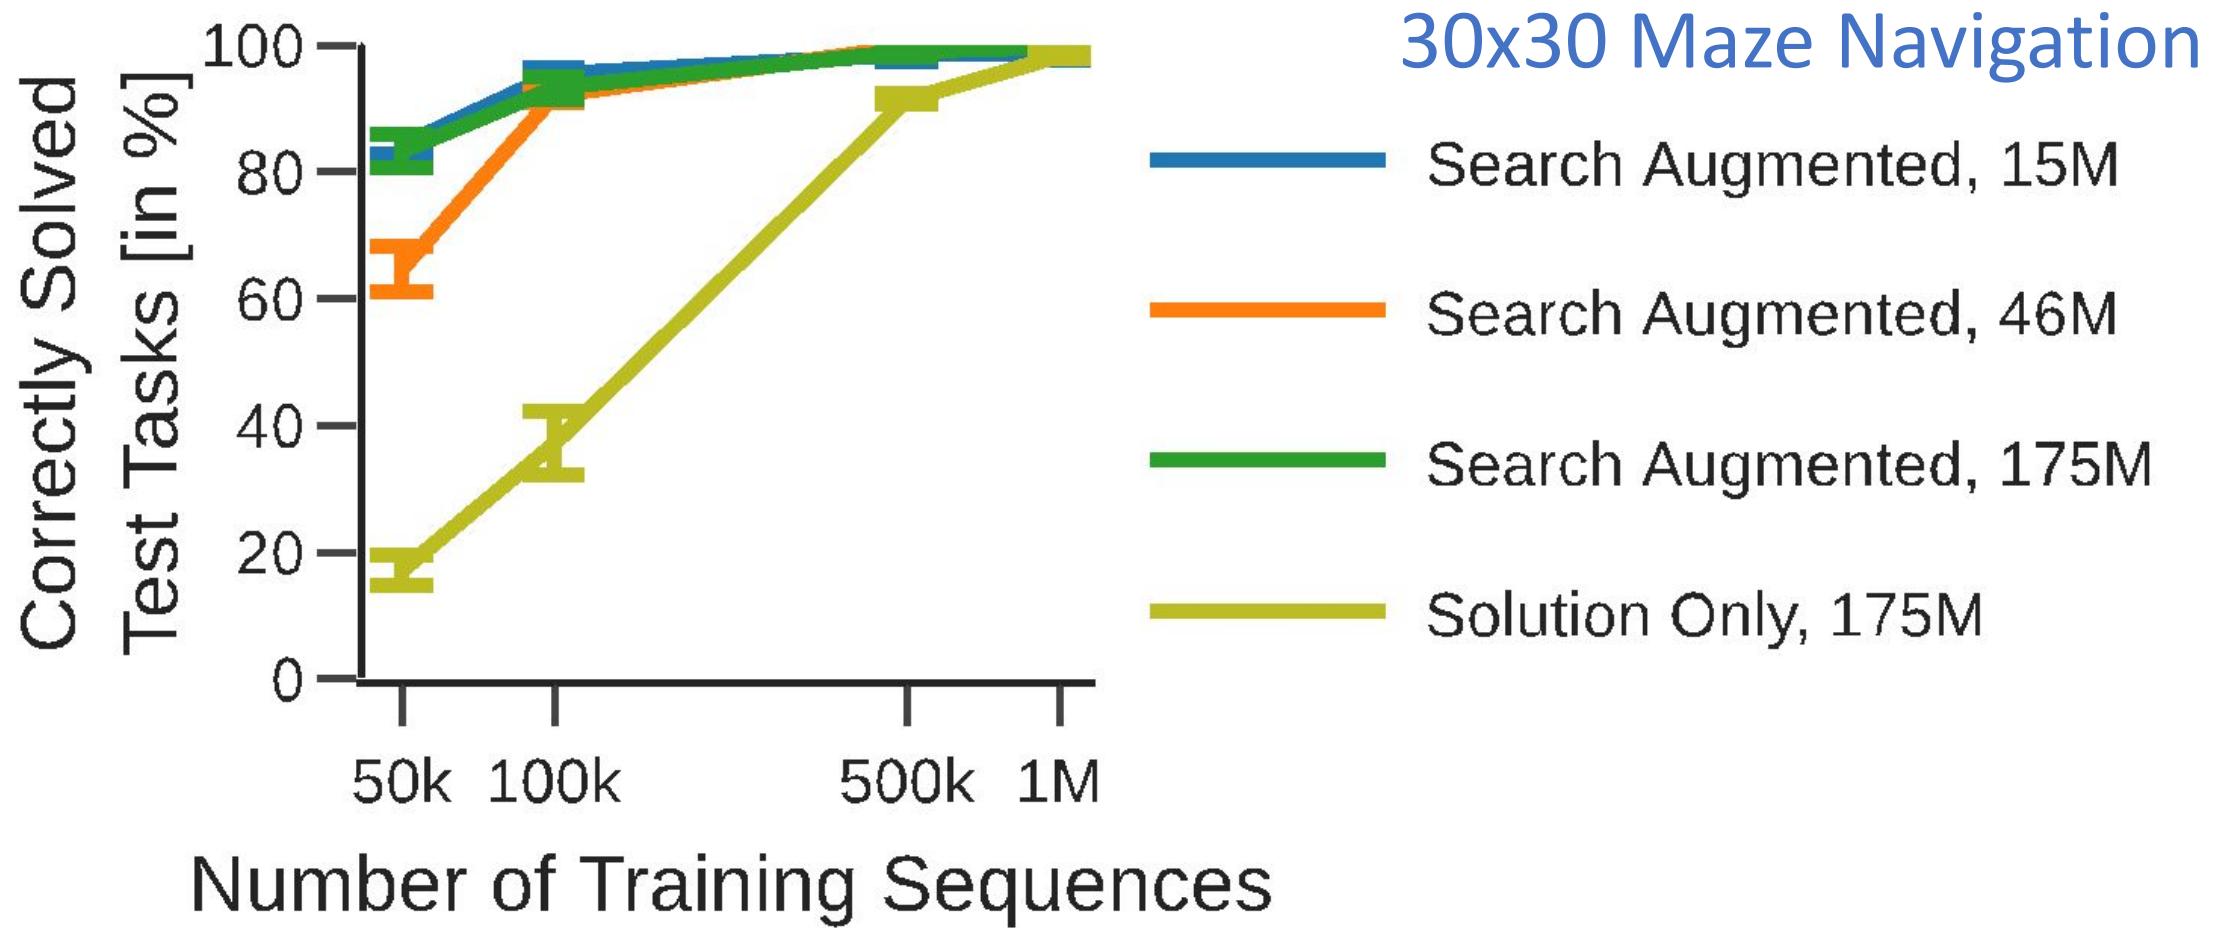


30x30 Maze Navigation

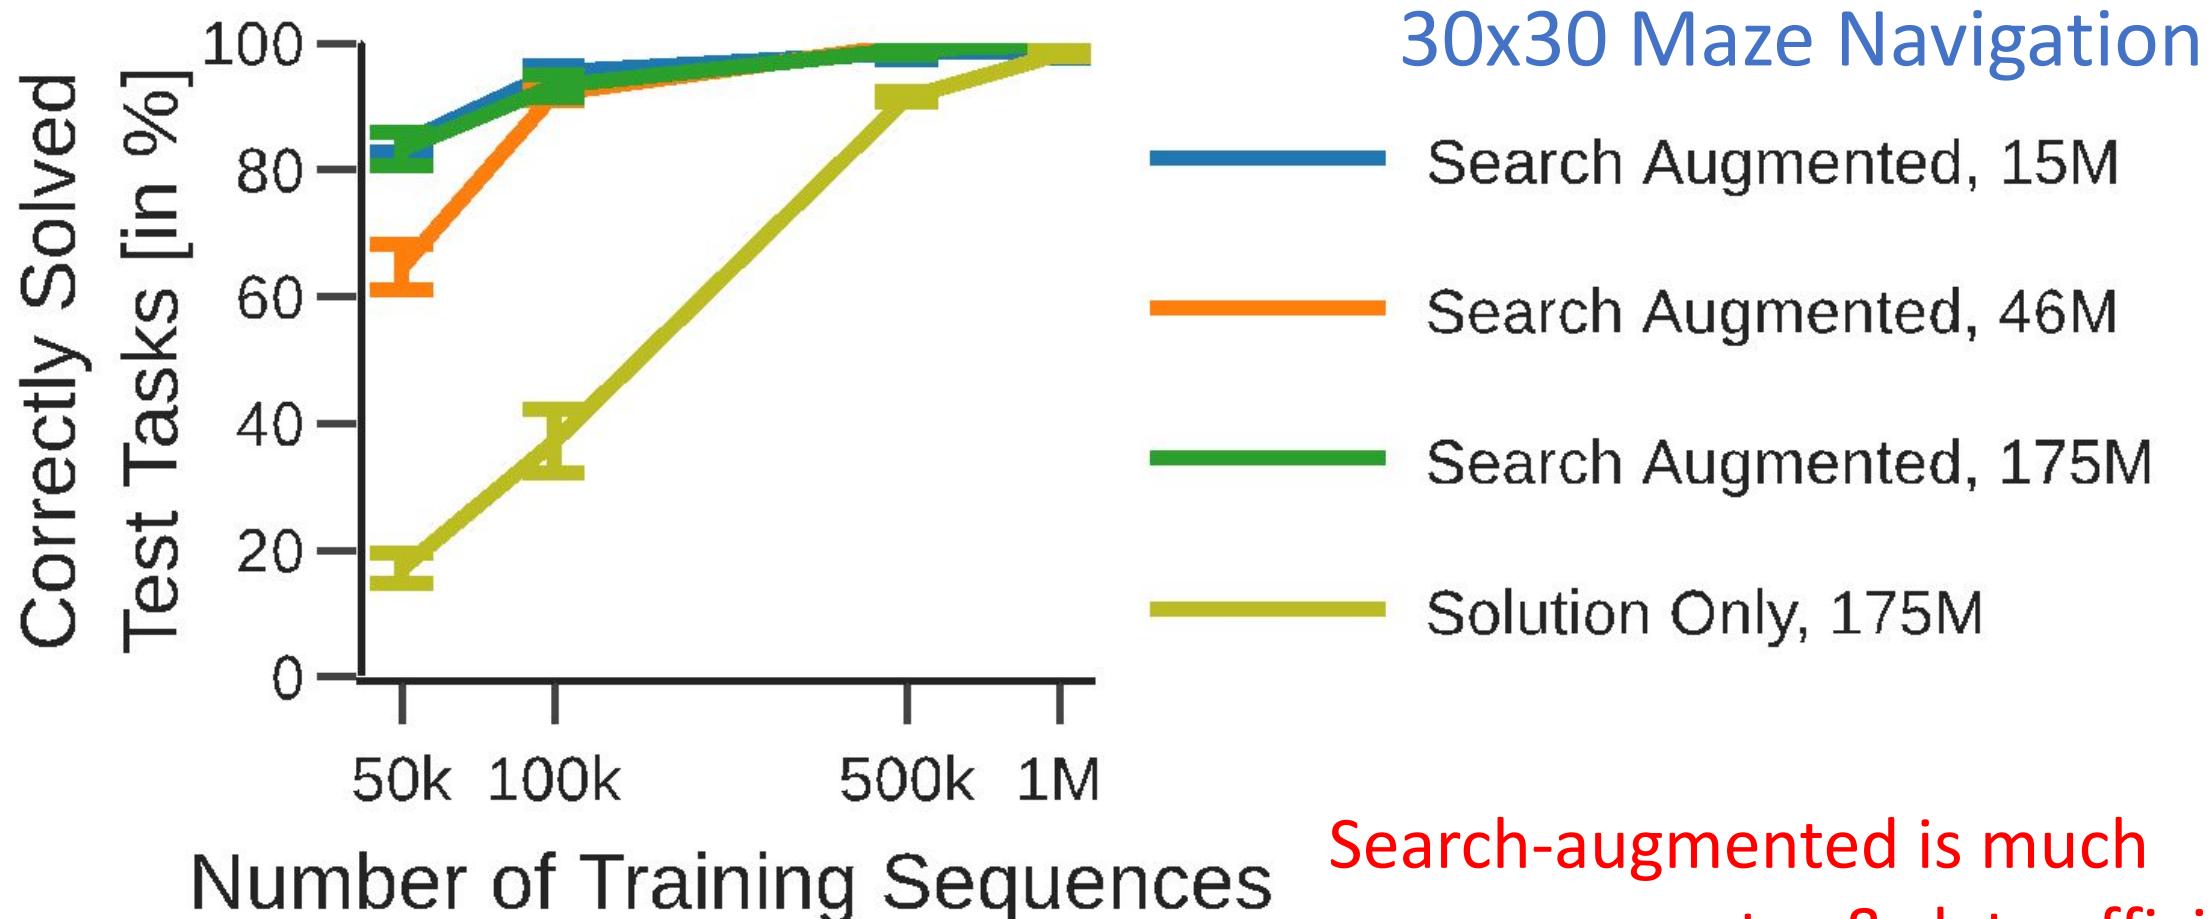


Solution Only, 175M

Search-Augmented vs. Solution-Only Models

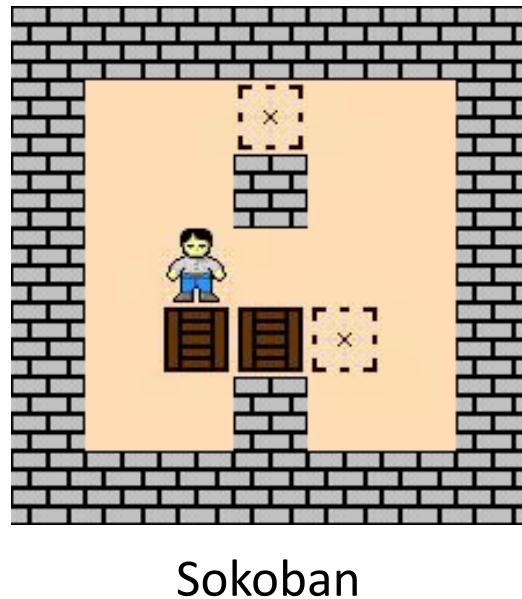


Search-Augmented vs. Solution-Only Models

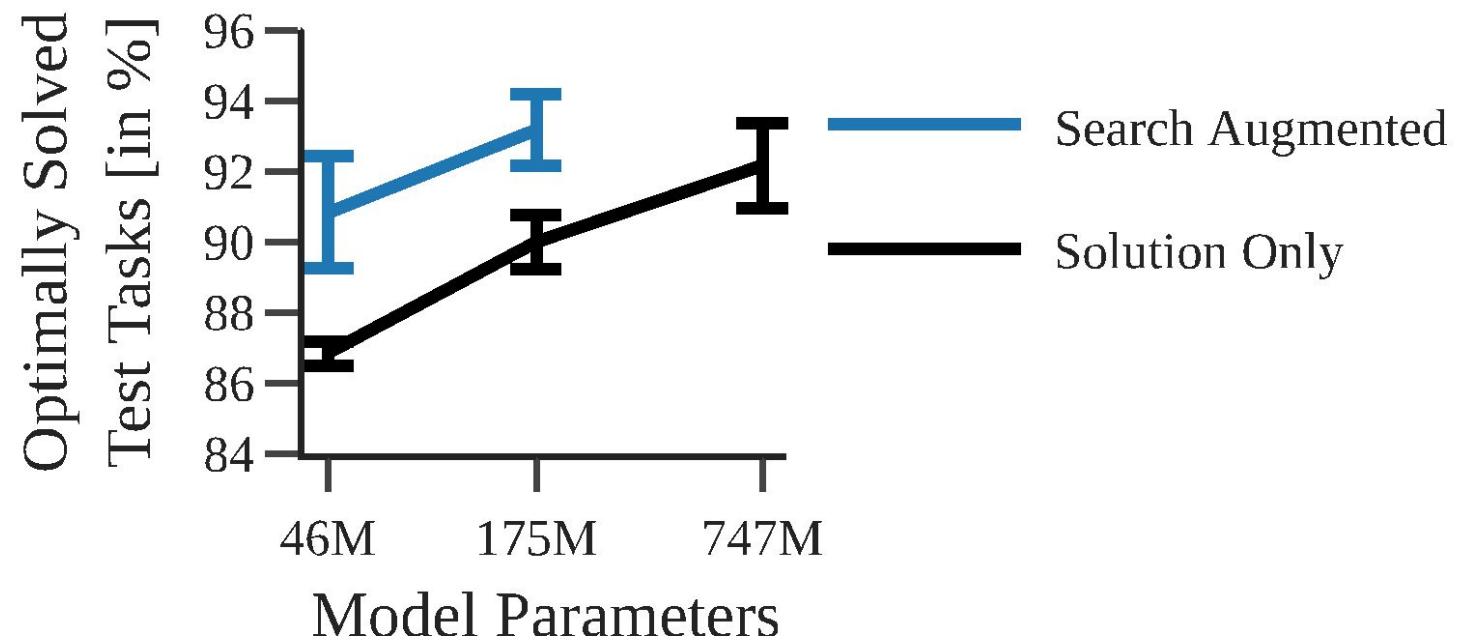


Search-augmented is much more parameter & data efficient!

Search-Augmented vs. Solution-Only Models



Sokoban



How to go beyond?

Imitation
Learning



Fine-tuning

Using solver's trace to train the Transformer with teacher forcing

Fine-tune the model to achieve **shorter** trace but still leads to **optimal** plan!
(Reinforcement Learning task)

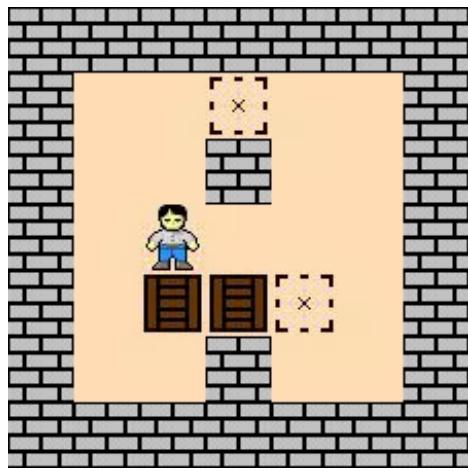


Search-augmented Models

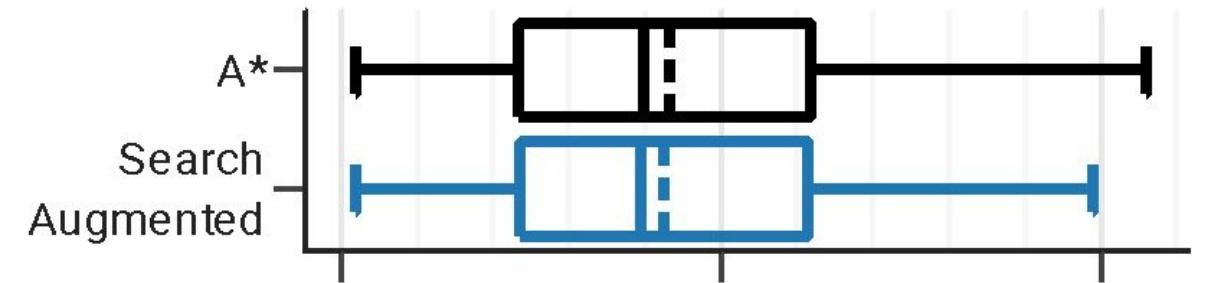


Searchformer

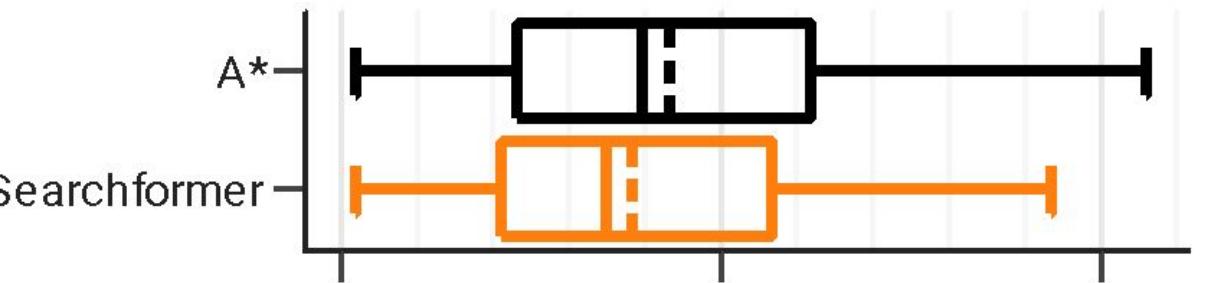
Beyond A*: Improving search dynamics via bootstrapping



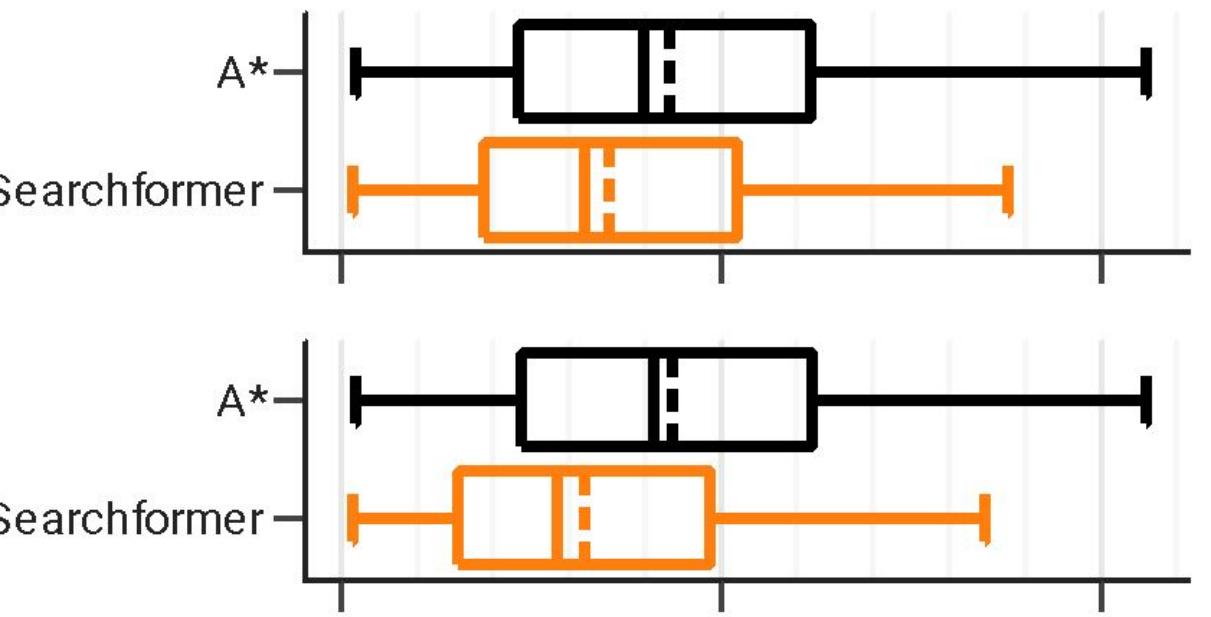
Step 1



Step 2



Step 3



Sequence Length Averaged per Test Task

Improving search dynamics via bootstrapping

Params.	Model	ILR-on-solved	ILR-on-optimal
45M	Solution only	–	–
	Search augmented	0.908 ± 0.020	0.919 ± 0.019
	Searchformer, step 1	1.054 ± 0.025	1.062 ± 0.015
	Searchformer, step 2	1.158 ± 0.025	1.181 ± 0.012
175M	Searchformer, step 3	1.292 ± 0.044	1.343 ± 0.067
	Solution only	–	–
757M	Search augmented	0.925 ± 0.010	0.933 ± 0.011
	Solution only	–	–

Repeated bootstrapping increases the
Improved Length Ratio (ILR)

Improving search dynamics via bootstrapping

Params.	Model	Solved (%)	Optimal (%)
45M	Solution only	90.3 ±1.0	86.8 ±0.3
	Search augmented	92.5 ±1.0	90.8 ±1.6
	Searchformer, step 1	95.5 ±1.0	93.5 ±1.0
	Searchformer, step 2	96.0 ±0.5	93.4 ±0.6
	Searchformer, step 3	95.5 ±0.8	93.7 ±1.6
175M	Solution only	95.7 ±0.2	90.0 ±0.8
	Search augmented	95.2 ±0.9	93.2 ±1.0
757M	Solution only	96.5 ±0.1	92.2 ±1.2

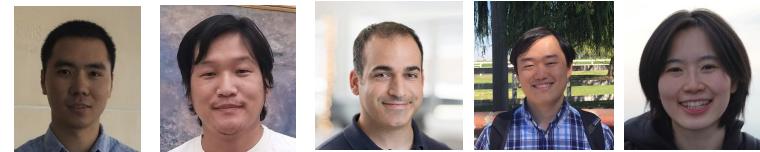
Fine-tuning improves performance initially.

Improving search dynamics via bootstrapping

Params.	Model	Solved (%)	Optimal (%)
45M	Solution only	90.3 ±1.0	86.8 ±0.3
	Search augmented	92.5 ±1.0	90.8 ±1.6
	Searchformer, step 1	95.5 ±1.0	93.5 ±1.0
	Searchformer, step 2	96.0 ±0.5	93.4 ±0.6
175M	Searchformer, step 3	95.5 ±0.8	93.7 ±1.6
	Solution only	95.7 ±0.2	90.0 ±0.8
	Search augmented	95.2 ±0.9	93.2 ±1.0
757M	Solution only	96.5 ±0.1	92.2 ±1.2

Searchformer outperforms largest solution-only model.

DualFormer (Searchformer v2)



A* search Trace

```
bos
create 0 2 c0 c3
close 0 2 c0 c3 ← D1 (drop a close clause)
create 0 1 c1 c2
close 0 1 c1 c2
create 0 0 c2 c1 ← D2 (drop cost tokens in
                     a clause)
create 1 1 c2 c1
close 0 0 c2 c1
create 1 0 c3 c0 ← D3 (drop a create clause)
close 1 0 c3 c0
```

Structured Trace Dropping Strategies

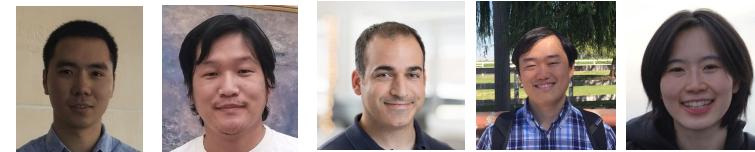
```
LvL 1 = D1           // drop all the close clauses

LvL 2 = D1 + D2      // drop all the close clauses
                     + all the cost tokens

Lvl 3 = D1 + D2 + sampled D3 // Lvl 2 + drop some
                             create clauses

Lvl 4 = drop the entire trace
```

DualFormer (Searchformer v2)



	Method	Avg Trace Length	1-Optimal-64 / 3-Optimal-64	1-Solved-64 / 3-Solved-64	SWC	Diversity
Maze 15 x 15	Dualformer (auto)	222	99.7 / 99.4	99.9 / 99.8	0.999	12.52
	Complete-Trace	495	94.6 / 90.1	96.7 / 93.0	0.964	7.60
	Solution-Only	-	72.0 / 68.9	82.7 / 80.1	0.610	1.52
Maze 20 x 20	Dualformer (auto)	351	99.5 / 98.6	99.9 / 99.3	0.997	20.28
	Complete-Trace	851	98.3 / 95.5	98.8 / 93.0	0.987	14.53
	Solution-Only	-	56.3 / 52.0	71.9 / 67.5	0.690	1.52
Maze 25 x 25	Dualformer (auto)	427	98.6 / 96.9	99.8 / 99.0	0.998	24.81
	Complete-Trace	1208	95.2 / 85.7	97.0 / 90.4	0.968	18.85
	Solution-Only	-	39.7 / 34.7	60.3 / 55.4	0.570	1.9
Maze 30 x 30	Dualformer (auto)	617	96.6 / 92.1	98.4 / 97.7	0.989	24.42
	Complete-Trace	1538	93.3 / 82.4	95.9 / 88.1	0.964	7.60
	Solution-Only	-	30.0 / 26.0	54.1 / 47.8	0.500	1.86
Sokoban	Dualformer (auto)	494	94.0 / 90.0	97.4 / 94.7	0.979	4.97
	Complete-Trace	3600	92.9 / 84.4	94.7 / 89.0	0.944	2.91
	Solution-Only	-	86.8 / 83.4	92.8 / 90.0	0.919	1.24

Dualformer **automatically** switches between fast mode (System 1) and slow mode (System 2) and works **better** for **dedicated** models on either modes.

Fast mode performance

	Method	1-Optimal-64 / 3-Optimal-64	1-Solved-64 / 3-Solved-64	SWC	Diversity
Maze 15x15	Dualformer (fast)	91.8 / 87.6	97.1 / 94.8	0.960	9.05
	Solution-Only	72.0 / 68.9	82.7 / 80.1	0.610	1.52
Maze 20x20	Dualformer (fast)	90.9 / 84.0	97.0 / 94.0	0.960	17.27
	Solution-Only	56.3 / 52.0	71.9 / 67.5	0.690	1.52
Maze 25x25	Dualformer (fast)	83.9 / 72.9	95.5 / 90.6	0.940	21.23
	Solution-Only	39.7 / 34.7	60.3 / 55.4	0.570	1.9
Maze 30x30	Dualformer (fast)	80.0 / 66.0	91.8 / 85.7	0.906	18.23
	Solution-Only	30.0 / 26.0	54.1 / 47.8	0.500	1.86
Sokoban	Dualformer (fast)	97.3 / 94.4	94.8 / 90.0	0.970	4.92
	Solution-Only	86.8 / 83.4	92.8 / 90.0	0.919	1.24

Slow mode performance

	Method	Avg Trace Length	1-Optimal-64 / 3-Optimal-64	1-Solved-64 / 3-Solved-64	SWC	Diversity
Maze 15 x 15	Dualformer (slow)	278	99.6 / 99.2	99.9 / 99.9	0.999	12.54
	Complete-Trace	495	94.6 / 90.1	96.7 / 93.0	0.964	7.60
Maze 20 x 20	Dualformer (slow)	439	98.9 / 97.8	99.9 / 99.7	0.998	18.86
	Complete-Trace	851	98.3 / 95.5	98.8 / 93.00	0.987	14.53
Maze 25 x 25	Dualformer (slow)	589	99.9 / 97.2	99.7 / 99.3	0.997	25.05
	Complete-Trace	1208	95.2 / 85.7	97.0 / 90.4	0.968	18.85
Maze 30 x 30	Dualformer (slow)	854	97.6 / 93.2	99.5 / 98.2	0.993	25.77
	Complete-Trace	1538	93.3 / 82.4	95.9 / 88.1	0.964	7.60
Sokoban	Dualformer (slow)	1482	94.5 / 87.6	97.4 / 94.1	0.970	4.66
	Complete-Trace	3600	92.9 / 84.4	94.7 / 89.0	0.944	2.91

Math Problems

Llama-3-Aug (baseline) Here is a detailed step-by-step explanation of the reasoning:

- Step 1: Understand the problem and the given information. We are asked to find the sum of the elements in the tenth row of Pascal's Triangle, given that the third row has four elements and their sum is 8.
- Step 2: Recall the formula for the sum of the elements in a row of Pascal's Triangle. The sum of the elements in the nth row is given by the formula: $\binom{n}{0} + \binom{n}{1} + \binom{n}{2} + \dots + \binom{n}{n} = 2^n$.
- Step 3: Identify the relationship between the sum of the elements in a row and the sum of the elements in the previous row. Since the sum of the elements in the nth row is 2^n , the sum of the elements in the (n-1)th row is 2^{n-1} .
- Step 4: Recognize that the sum of the elements in the third row is 8, which is equal to 2^3 . This is because the sum of the elements in the third row is $\binom{3}{0} + \binom{3}{1} + \binom{3}{2} = 1 + 3 + 3 = 7$, and the extra 1 comes from the top row, which is counted as the 0th row.

Question 1

In Pascal's Triangle, the third row has four elements, and the sum of these elements is 8. What is the sum of the elements in the tenth row? (Note that we are counting the top row with a single 1 as the 0th row).

- Step 5: Use the relationship between the sums of the elements in consecutive rows to find the sum of the elements in the tenth row. Since the sum of the elements in the third row is 2^3 , and the sum of the elements in the (n-1)th row is 2^{n-1} , we can set up the equation: $2^3 = 2^{10-1}$.
- Step 6: Solve the equation for n. Taking the logarithm of both sides, we get: $3 \log 2 = (10 - 1) \log 2$. Simplifying, we get: $3 \log 2 = 9 \log 2$, which implies that $\log 2 = 0$. This is a contradiction, since $\log 2$ is a positive number.
- Step 7: Conclude that the equation $2^3 = 2^{10-1}$ has no solution, and therefore, the sum of the elements in the tenth row is not equal to 2^{10-1} .
- Step 8: Recognize that the sum of the elements in the tenth row is still given by the formula: $\binom{10}{0} + \binom{10}{1} + \binom{10}{2} + \dots + \binom{10}{10} = 2^{10}$.

The answer is: 1024.

Baseline

Llama-3-Aug (p=0.4, ours) Here is a detailed step-by-step explanation of the reasoning:

- Step 1: Understand the problem and the given information. We are asked to find the sum of the elements in the tenth row of Pascal's Triangle, given that the third row has four elements and their sum is 8.
- Step 2: Recall the properties of Pascal's Triangle. The elements in each row of Pascal's Triangle are the binomial coefficients, which represent the number of ways to choose k items from a set of n items. The sum of the elements in each row is equal to 2^n , where n is the row number.
- Step 3: Identify the relationship between the sum of the elements in the third row and the sum of the elements in the tenth row. Since the sum of the elements in each row is equal to 2^n , we can write the sum of the elements in the third row as $2^3 = 8$. Similarly, the sum of the elements in the tenth row is 2^{10} .
- Step 4: Calculate the sum of the elements in the tenth row. $2^{10} = 1024$.

The answer is: 1024.

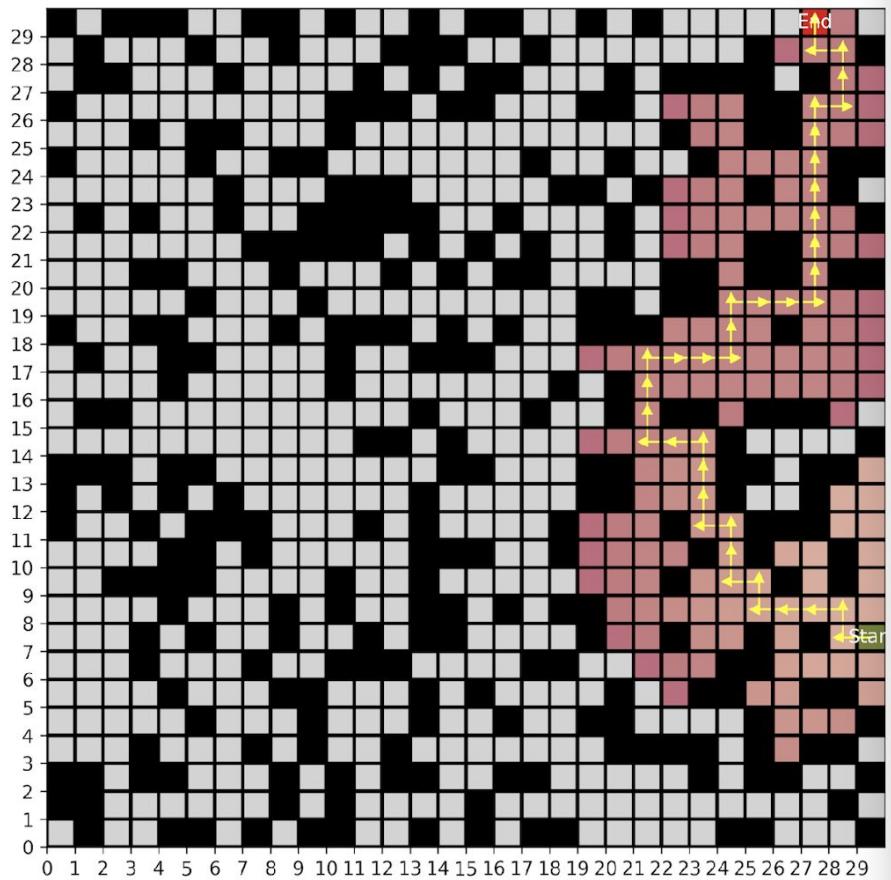
Dualformer

Math Problems

Model	Dataset & Dropping Prob	Greedy@1% (slow / fast)	Trace Length	Pass@20% (slow / fast)	Trace Length
Mistral-7B	Aug-MATH (baseline)	16.9 / 9.6	527 / -	59.6 / 29.8	521 / -
	Aug-MATH (p=0.1)	18.6 / 11.3	508 / -	61.6 / 32.0	479 / -
	Aug-MATH (p=0.2)	17.8 / 11.2	477 / -	61.4 / 31.9	470 / -
	Aug-MATH (p=0.3)	17.8 / 11.8	497 / -	61.9 / 31.7	466 / -
	Aug-MATH (p=0.4)	17.0 / 11.0	434 / -	56.4 / 28.9	397 / -
	MATH	13.1 / 8.5	290 / -	53.0 / 29.4	227 / -
Llama-3-8B	Aug-MATH (baseline)	19.7 / 13.1	548 / -	62.7 / 35.6	535 / -
	Aug-MATH (p=0.1)	20.1 / 13.3	544 / -	63.4 / 36.2	522 / -
	Aug-MATH (p=0.2)	20.5 / 13.8	525 / -	63.9 / 36.7	497 / -
	Aug-MATH (p=0.3)	20.5 / 13.5	515 / -	63.4 / 37.5	474 / -
	Aug-MATH (p=0.4)	20.4 / 13.5	490 / -	63.4 / 37.2	450 / -
	MATH	13.3 / 12.6	432 / -	52.8 / 35.5	332 / -

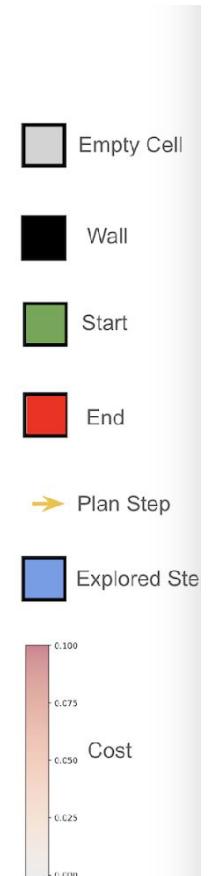
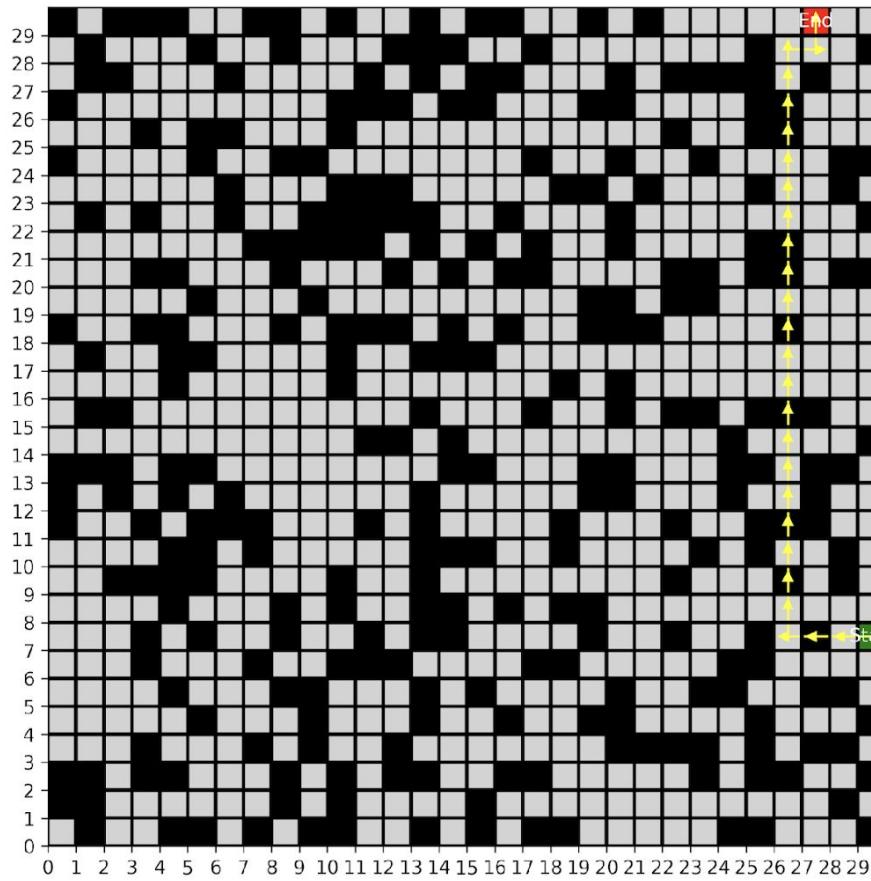
DualFormer

Dualformer

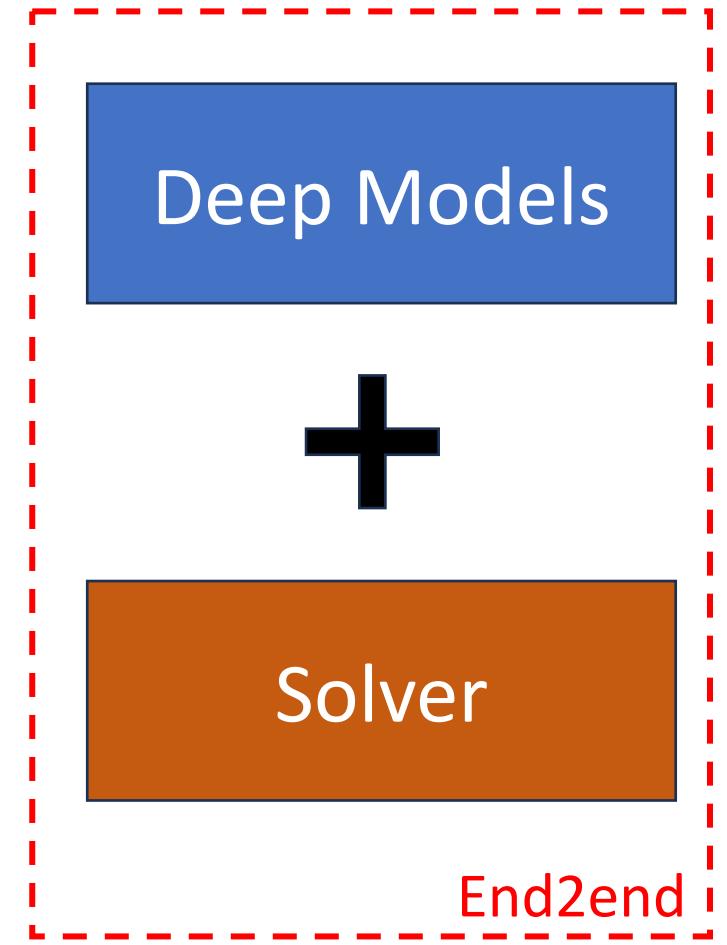
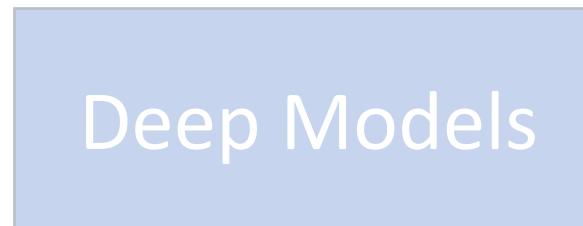
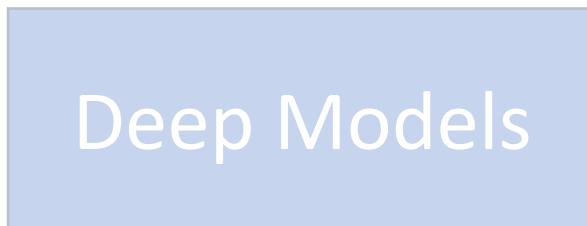


facebook

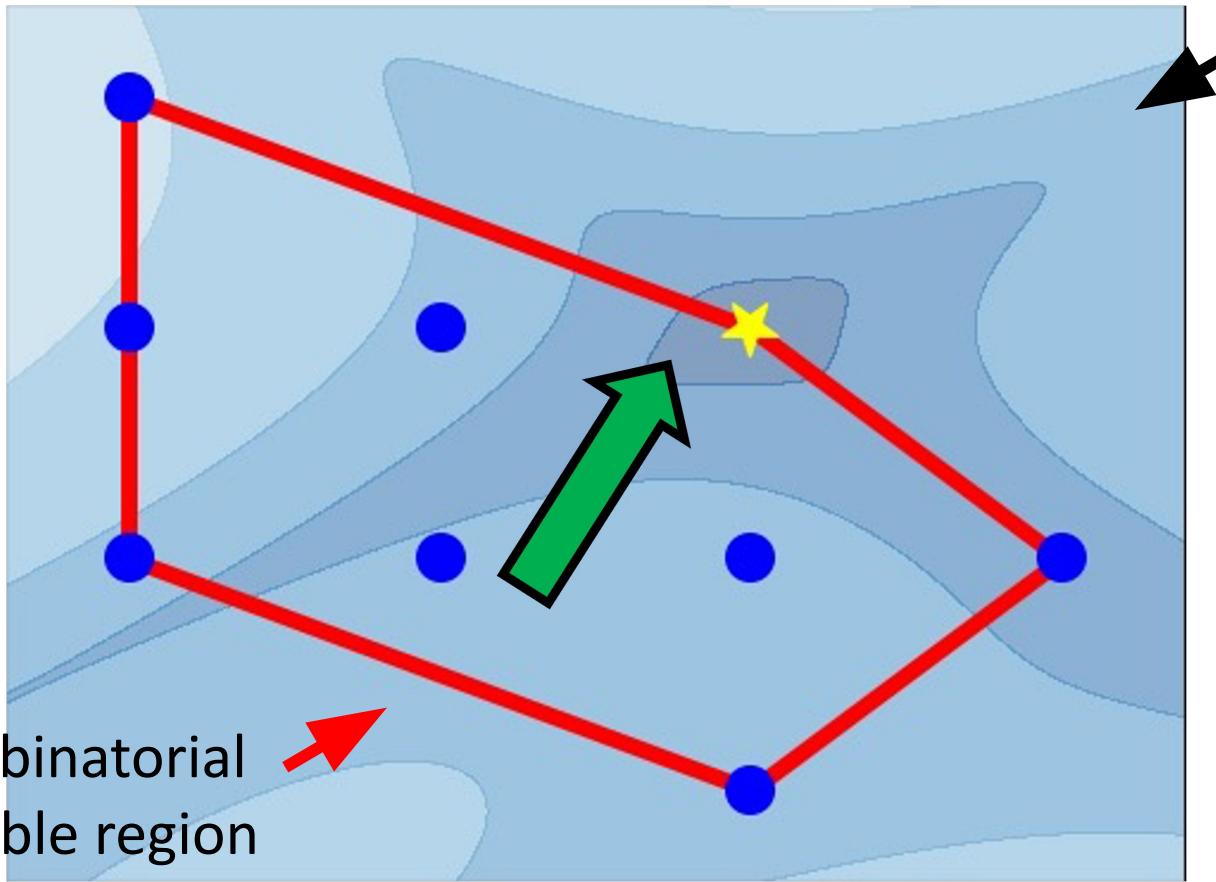
o1-preview (OpenAI)



Option Two: Hybrid Systems



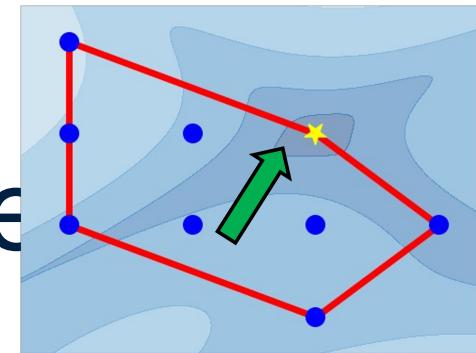
Nonlinear objective with combinatorial constraints



Nonlinear + differentiable
objective $f(\mathbf{x})$

- Real-world domains:
 - Computer system planning
 - Designing photonic devices
 - Throughput optimization
 - Antenna design
 - Energy grid

Example: Embedding Table Placement

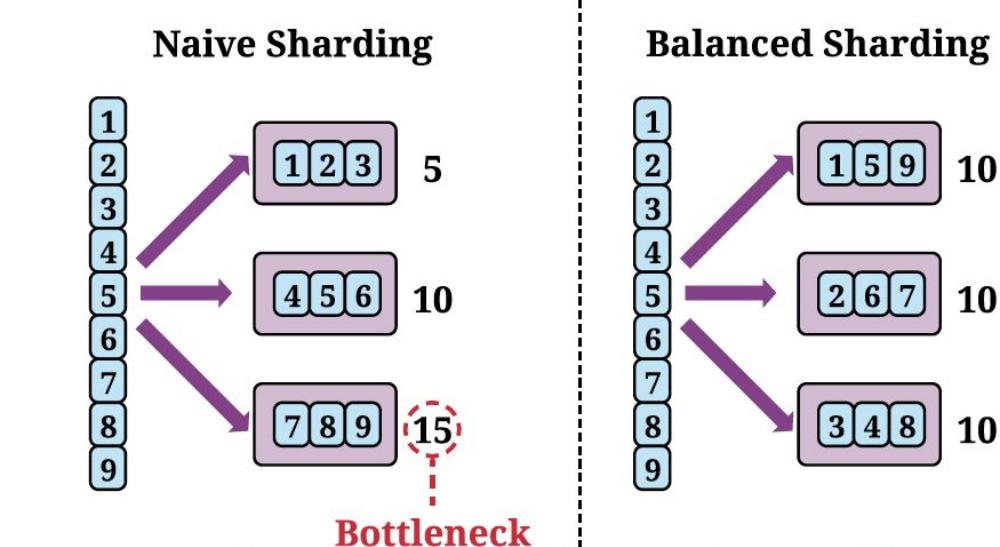


Given:

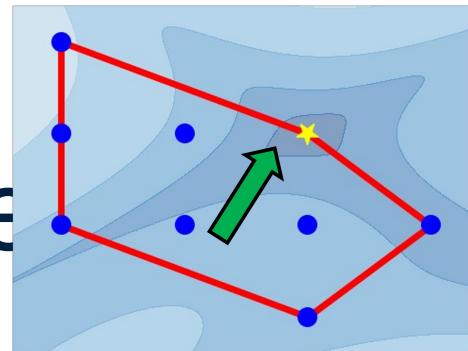
- k tables
- n identical devices
- Table i has memory requirement m_i
- Device j has memory capacity M_j

Find

- Allocation of tables to devices observing device memory limits
- Minimize latency which is **estimated by a neural network** (capturing nonlinear interactions)

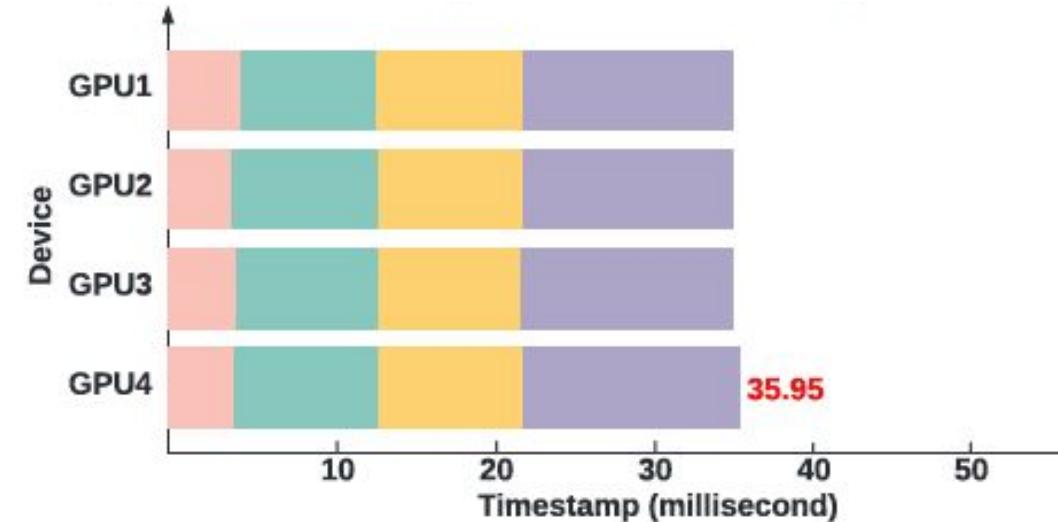


Example: Embedding Table Placement



Given:

- k tables
- n identical devices
- Table i has memory requirement m_i
- Device j has memory capacity M_j



Formulation

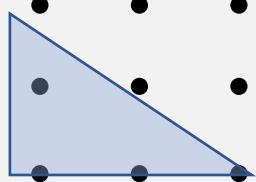
$$\text{Min}_x \textcolor{green}{L}(\{x_{ij}\}) \quad \text{s.t.} \quad \sum_i x_{ij} m_i \leq M_j, \quad \sum_j x_{ij} = 1, \quad x_{ij} \in \{0,1\}$$

$\textcolor{green}{L}$ is nonlinear due to system issues (e.g., batching, communication, etc)

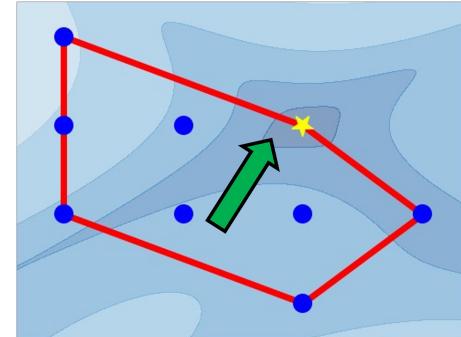
Solve the Combinatorial Problem in the Latent Space

Original Space

Nonlinear optimization with combinatorial constraints

$$\begin{aligned} \min_{\boldsymbol{x}} f(\boldsymbol{x}; \boldsymbol{y}) \\ \text{s.t } \boldsymbol{x} \in \Omega = \end{aligned}$$


combinatorial constraints



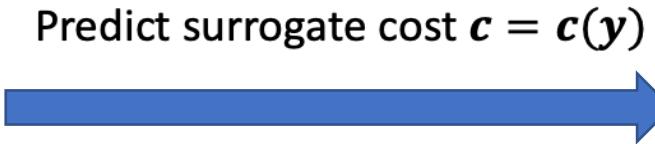
Latent Space

Surrogate optimization

$$\boldsymbol{x}^*(\boldsymbol{y}) = \operatorname{argmin}_{\boldsymbol{x}} \boldsymbol{c}(\boldsymbol{y})^T \boldsymbol{x}$$

$$\text{s.t } \boldsymbol{x} \in \Omega$$

solved by existing combinatorial solvers

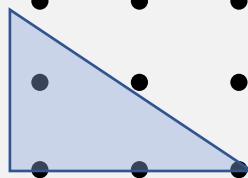


$\boldsymbol{x}^*(\boldsymbol{y})$ optimizes $f(\boldsymbol{x}; \boldsymbol{y})$ as much as possible

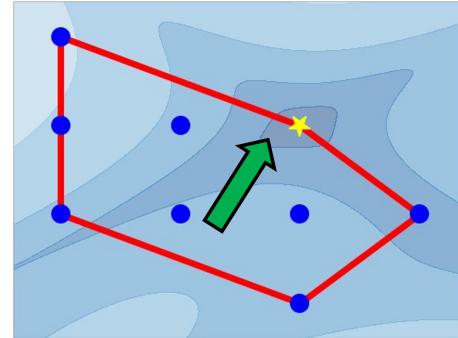
Solve the Combinatorial Problem in the Latent Space

Original Space

Nonlinear optimization with combinatorial constraints

$$\begin{aligned} \min_{\boldsymbol{x}} f(\boldsymbol{x}; \boldsymbol{y}) \\ \text{s.t } \boldsymbol{x} \in \Omega = \end{aligned}$$


combinatorial constraints

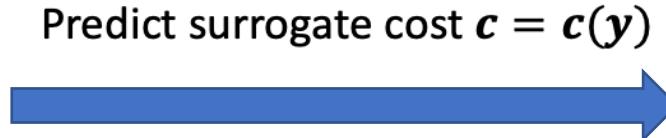


Latent Space

Surrogate optimization

$$\begin{aligned} \boldsymbol{x}^*(\boldsymbol{y}) = \operatorname{argmin}_{\boldsymbol{x}} \boldsymbol{c}(\boldsymbol{y})^T \boldsymbol{x} \\ \text{s.t } \boldsymbol{x} \in \Omega \end{aligned}$$

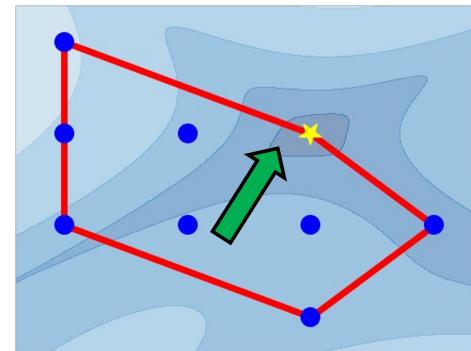
solved by existing combinatorial solvers



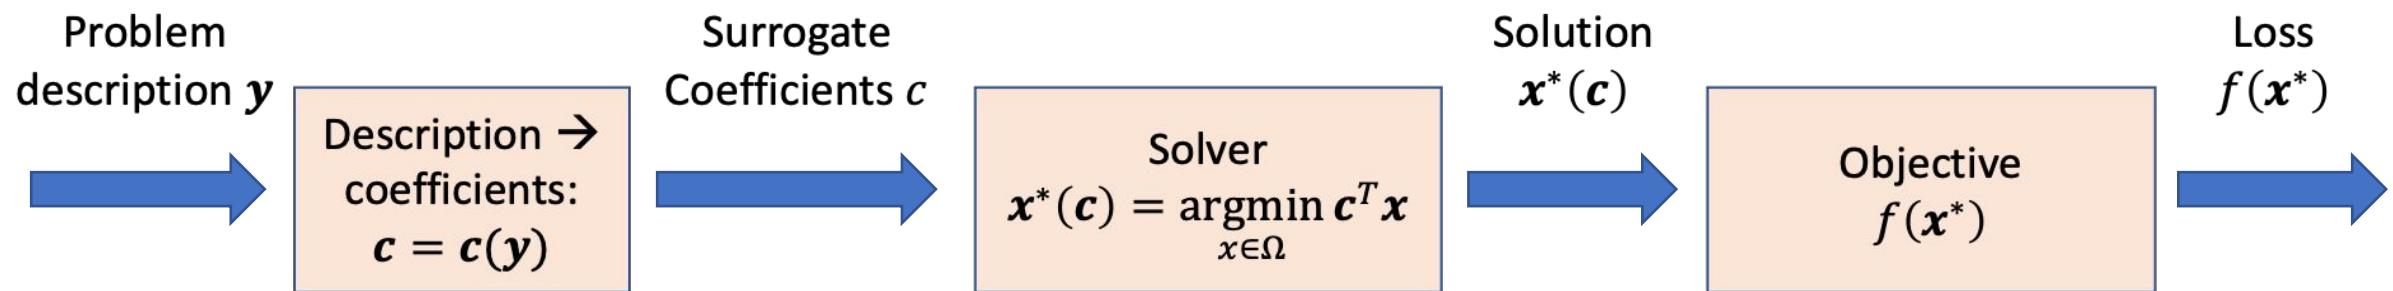
$\boldsymbol{x}^*(\boldsymbol{y})$ optimizes $f(\boldsymbol{x}; \boldsymbol{y})$ as much as possible

Proposal: gradient-based optimization

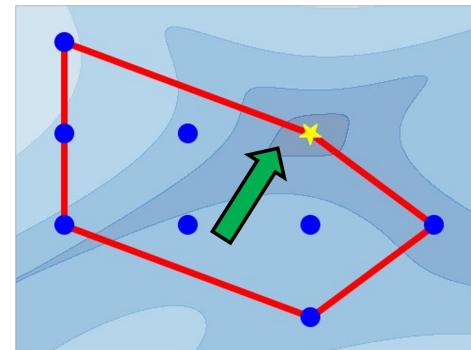
SurCo: Surrogate combinatorial opt



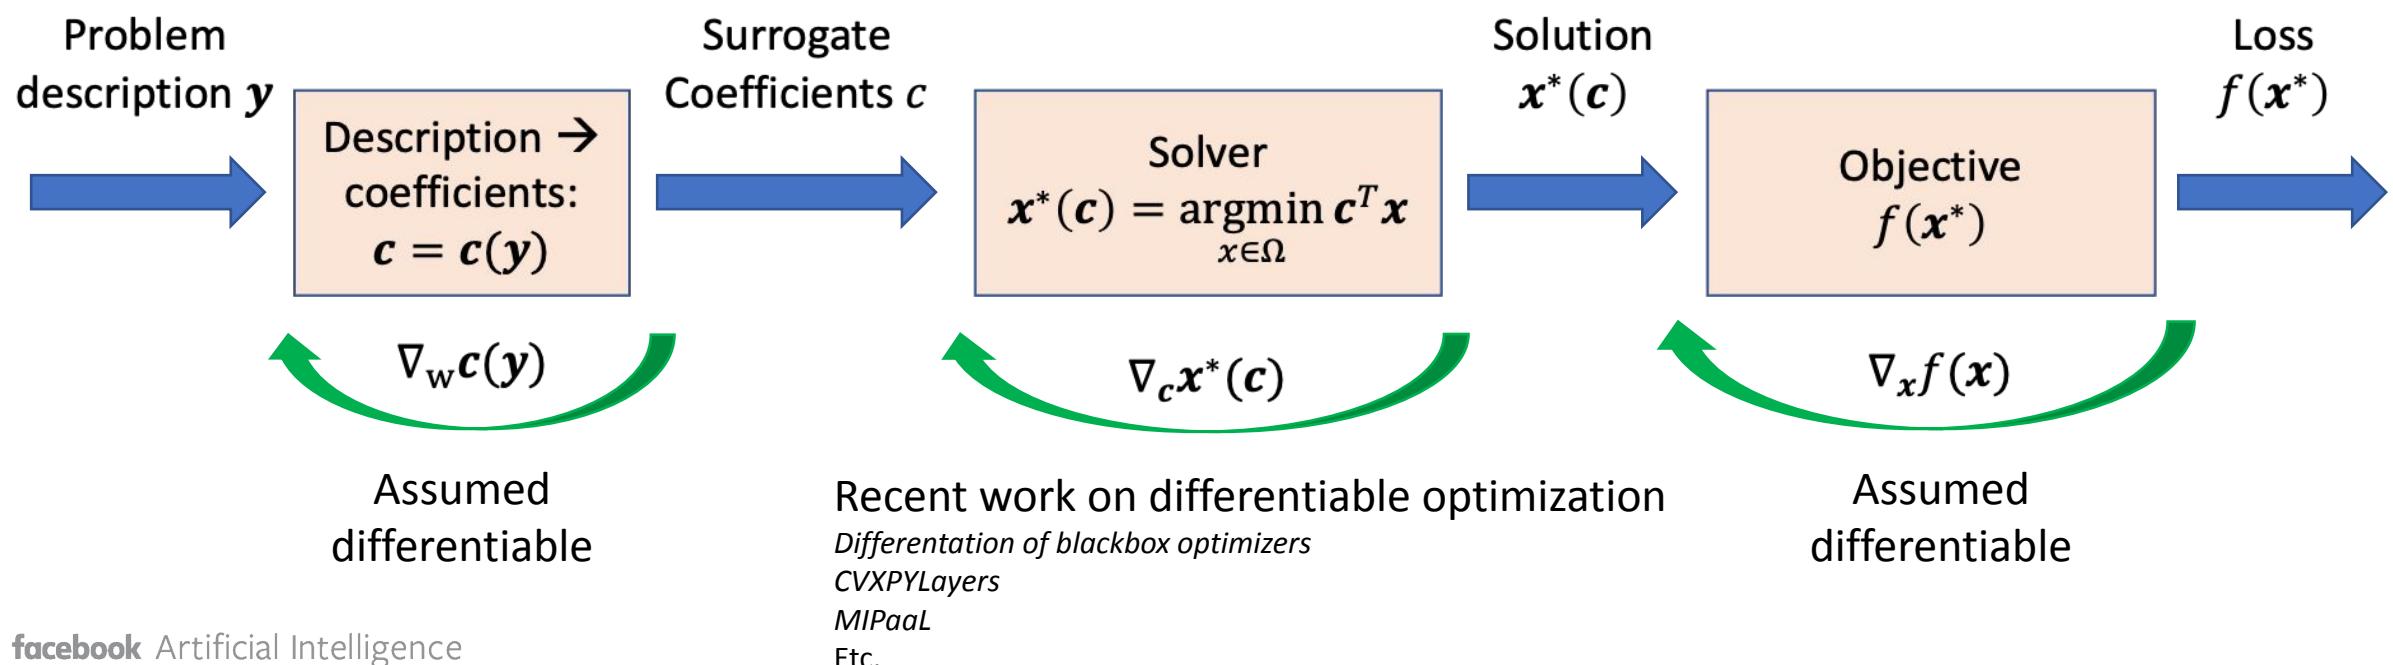
- Use surrogate MILP to solve original problem
- Find linear coefficients \mathbf{c} such that $\underset{x \in \Omega}{\operatorname{argmin}} f(x) = \underset{x \in \Omega}{\operatorname{argmin}} \mathbf{c}^T \mathbf{x}$



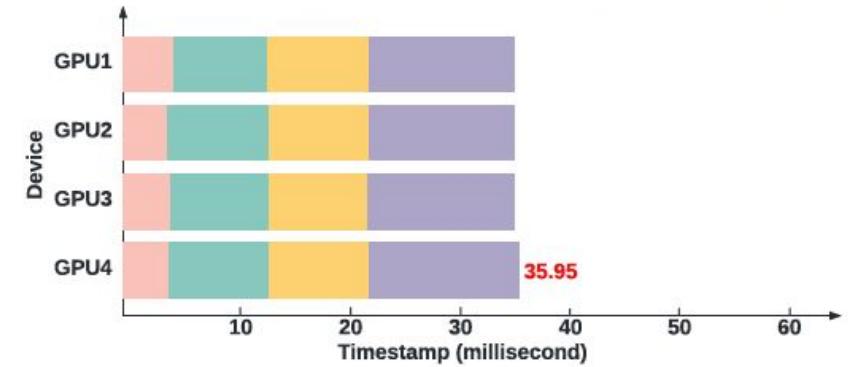
Gradient-based Optimization



- Use surrogate MILP to solve original problem
- Find linear coefficients \mathbf{c} such that $\underset{x \in \Omega}{\operatorname{argmin}} f(x) = \underset{x \in \Omega}{\operatorname{argmin}} \mathbf{c}^T \mathbf{x}$



Embedding Table Sharding

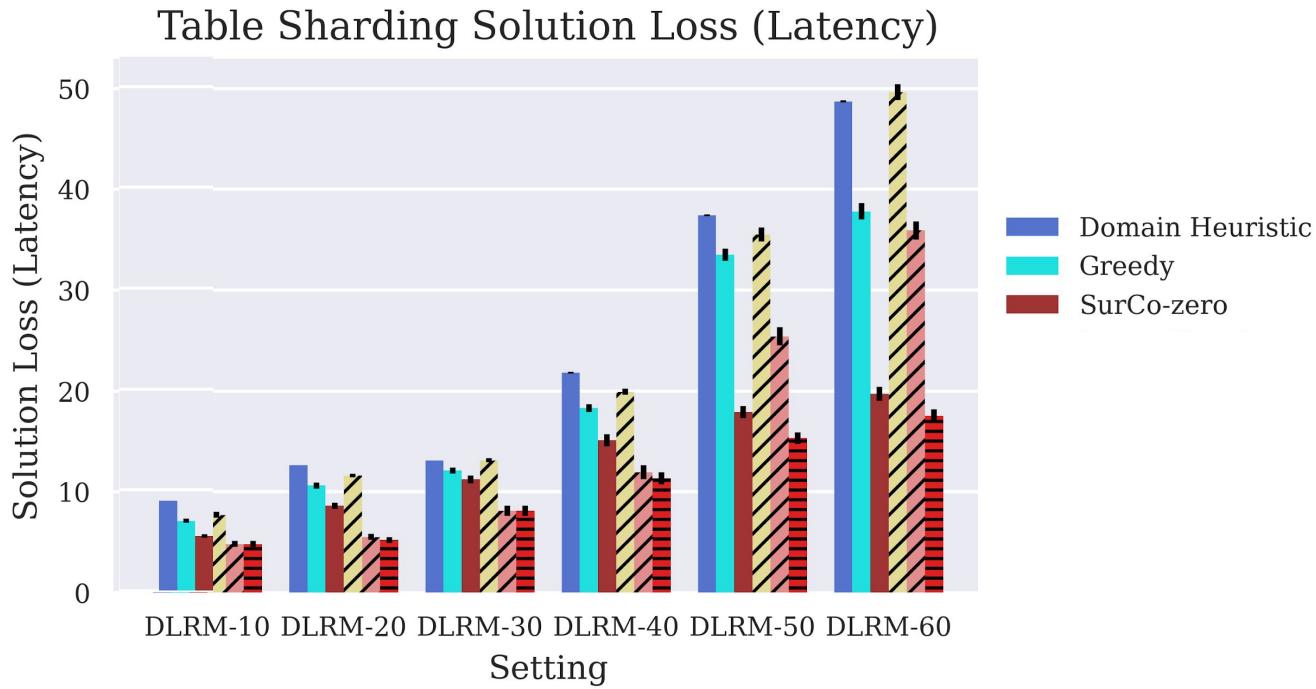


- Public Deep Learning Recommendation Model (DLRM dataset) placing between 10 to 60 tables on 4 GPUs
- Baseline: Greedy
- SoTA: RL approach Dreamshard¹
- SurCo: Surrogate NN model learned via CVXPYLayers (differentiable LP Solver)

¹ Zha et al. NeurIPS 2022

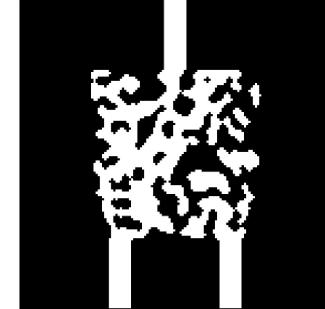
Dataset: https://github.com/facebookresearch/dlrm_datasets

Results – Table Sharding

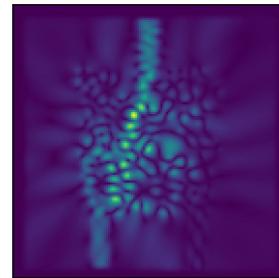


Inverse Photonic Design

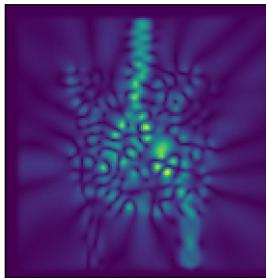
Device Design



E_z magnitude
first wavelength



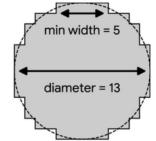
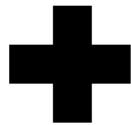
E_z magnitude
second wavelength



- Design physically-viable devices that take light waves and routes different wavelengths to correct locations

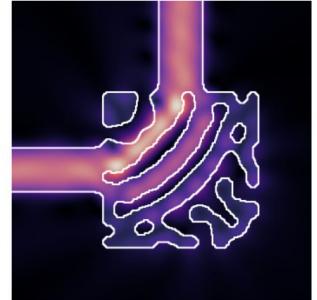
$$\mathcal{L}(S) = \left(\left\| \text{softplus} \left(g \frac{|S|^2 - |S_{\text{cutoff}}|^2}{\min(w_{\text{valid}})} \right) \right\|_2 \right)^2$$

- Device design misspecification loss $f(\mathbf{x})$ computed by differentiable electromagnetic simulator
- Feasible solution: the design must be the union of brush patterns
 - $\mathbf{x} = \text{binary_opening}(\mathbf{x}, \text{brush})$
 - $\mathbf{x} = \sim \text{binary_opening}(\sim \mathbf{x}, \text{brush})$

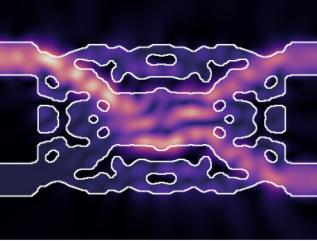


Inverse Photonic Design

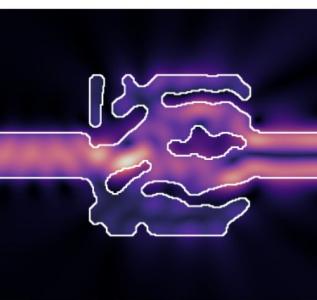
- Dataset: Ceviche Challenges¹
- Most baselines don't work here due to combinatorial constraints
- SoTA: Brush-based algorithm ¹
- SurCo: Surrogate learned via blackbox differentiation² of brush solver



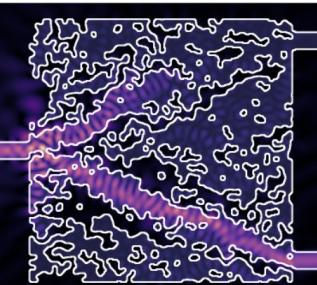
Waveguide
bend



Beam splitter



Mode converter



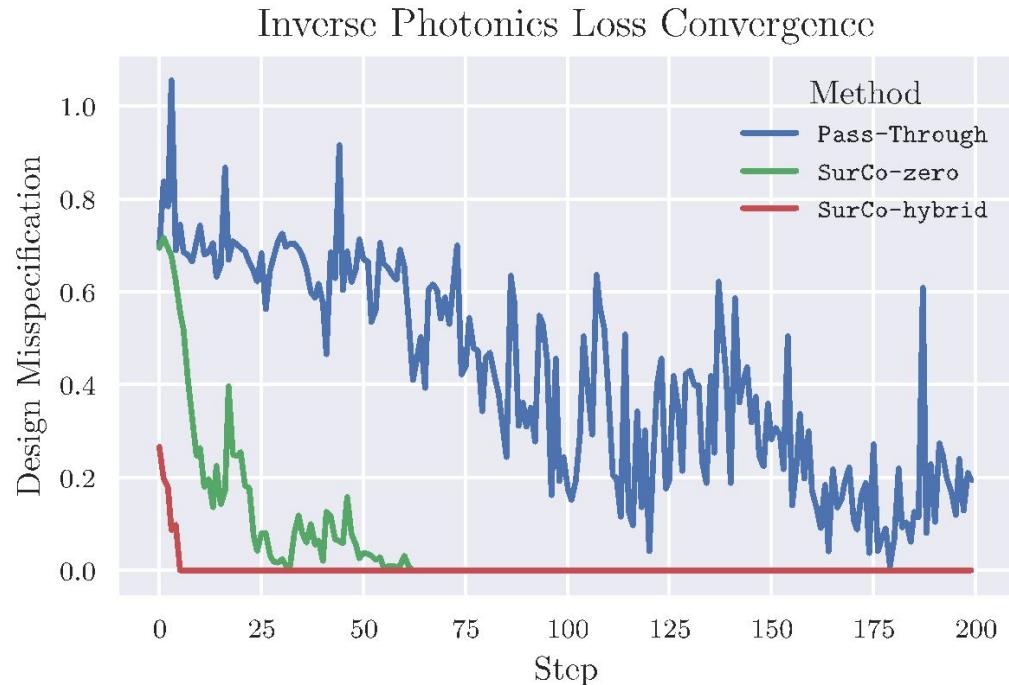
Wavelength division multiplexer

¹Schubert et al. ACS Photonics 2022

²Vlastelica et al. ICLR 2019

Dataset: <https://github.com/google/ceviche-challenges>

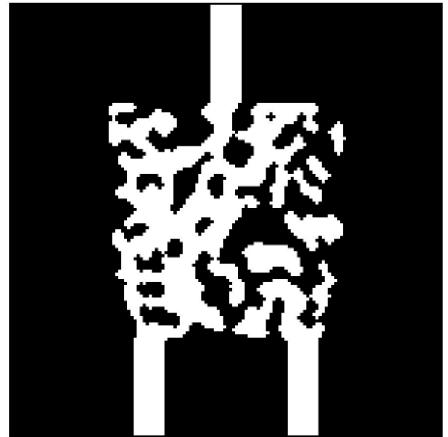
Inverse photonics Convergence comparison + Solution example



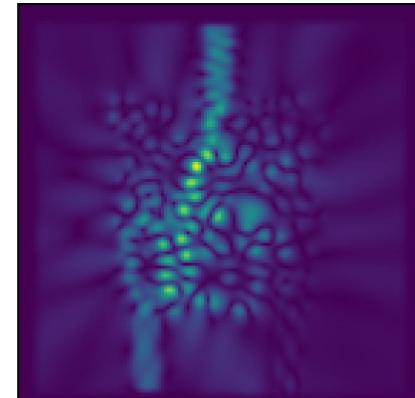
Takeaways:

- SurCo-Zero finds loss-0 solutions quickly
- SurCo-Hybrid uses offline training data to get a head start

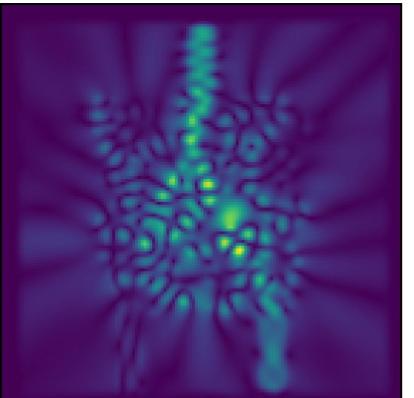
Device Design



E_z magnitude
first wavelength



E_z magnitude
second wavelength



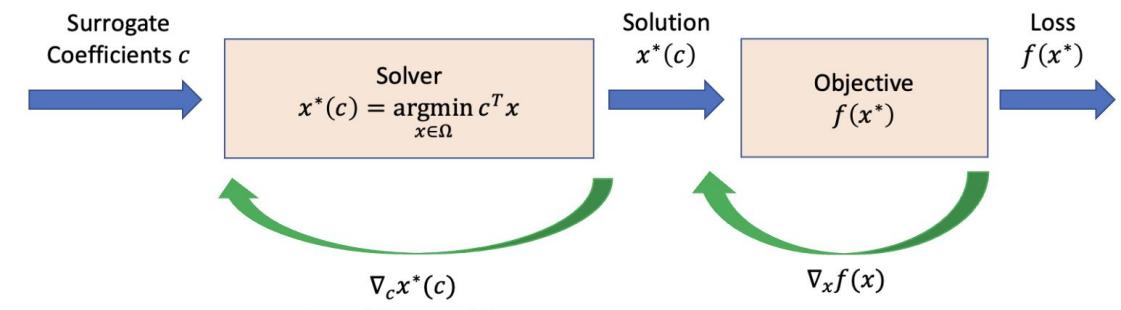
Wavelength division
multiplexer

Limitation of SurCo

Recall **SurCo**: Update linear coefficients \mathbf{c} such that $x^*(\mathbf{c})$ improves objective $f(x^*(\mathbf{c}))$

$$\min_{\theta} \mathcal{L}(Y, Z) := \sum_{i=1}^N f(\mathbf{g}_{\theta}(\mathbf{y}_i); \mathbf{z}_i)$$

$$\mathbf{g}_{\theta}(\mathbf{y}) = \arg \min_{\mathbf{x} \in \Omega} \mathbf{x}^\top \mathbf{c}_{\theta}(\mathbf{y})$$



- Requires $\nabla_x f(x) \rightarrow$ Does not apply to nondifferentiable functions
- Requires $\nabla_c g_{\theta}(c) \rightarrow$ Solver is backpropagatable

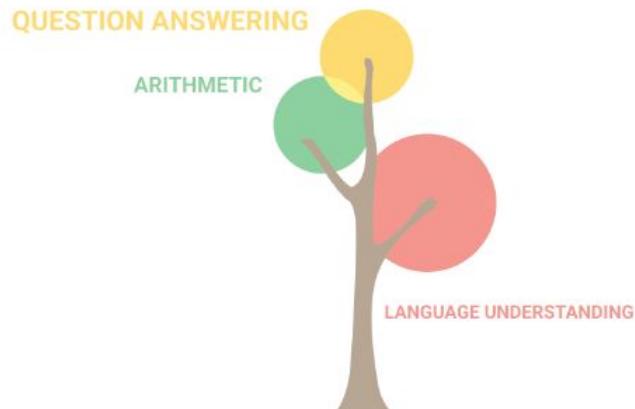
Option Three: Does Deep Model Actually Converge to Anything Symbolic?



Emerging Symbolic
Structure

Deep Models

Debate: Is LLM doing retrieval or true reasoning?



8 billion parameters

LLM shows emergent behaviors!!

Debate: Is LLM doing retrieval or true reasoning?

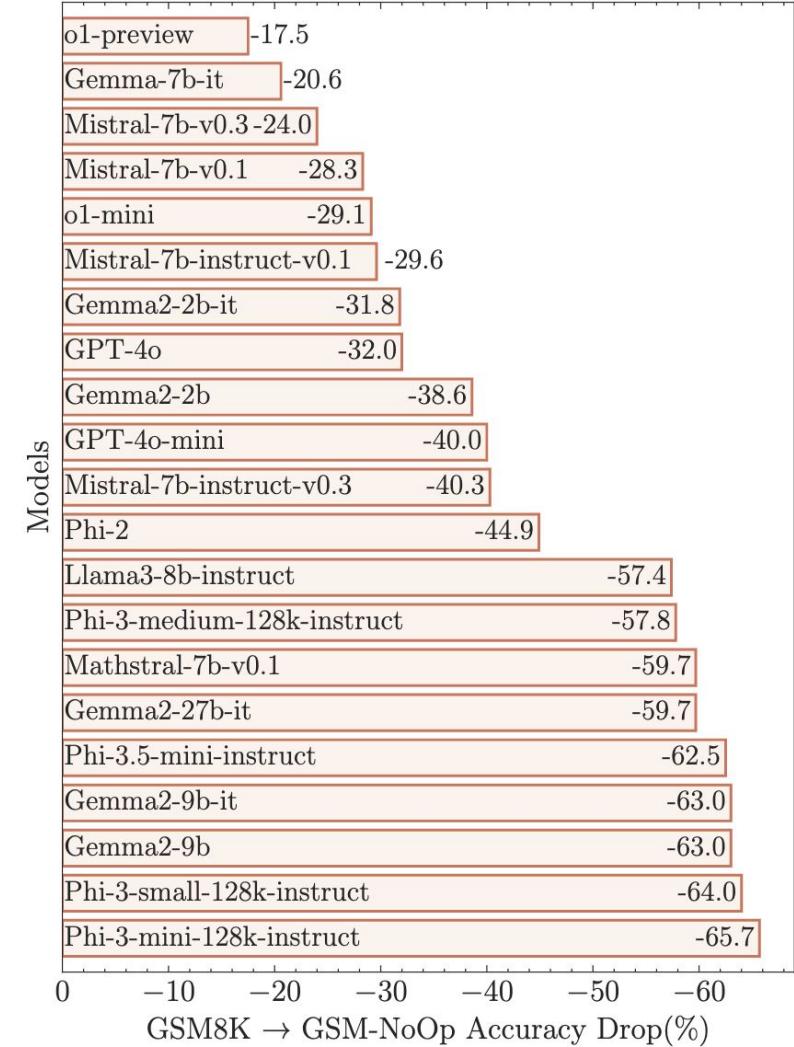


Do LLMs perform reasoning or approximate retrieval?
There is a continuum between the two, and Auto-Regressive LLMs are largely on the retrieval side.

...



Emergent Abilities (noun): The preferred euphemism for what your LLM does, when saying "approximate retrieval" sounds too unsexy.



#AI Aphorisms

LLM is just doing retrievals!!

Concrete Example: Modular Addition

$$a + b = c \bmod d$$

Does neural network have an *implicit table* to do retrieval?

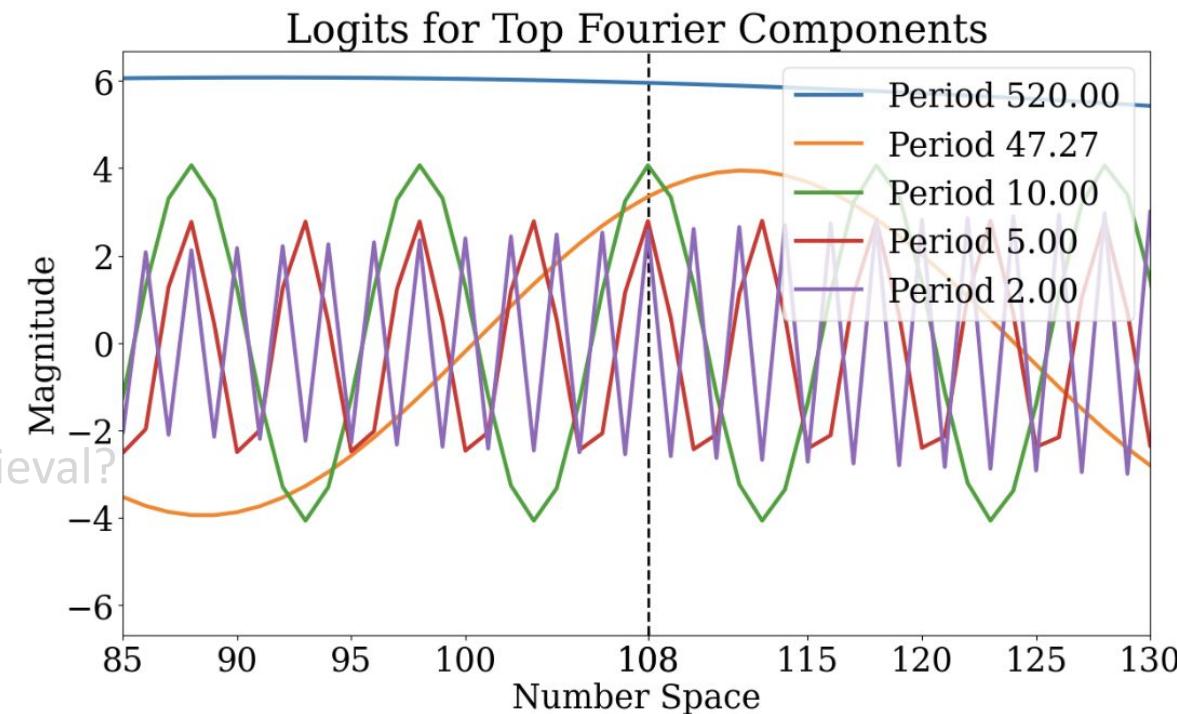
Concrete Example: Modular Addition

$$a + b = c \bmod d$$

Does neural network have an *implicit table* to do retrieval?

Learned representation = Fourier basis 

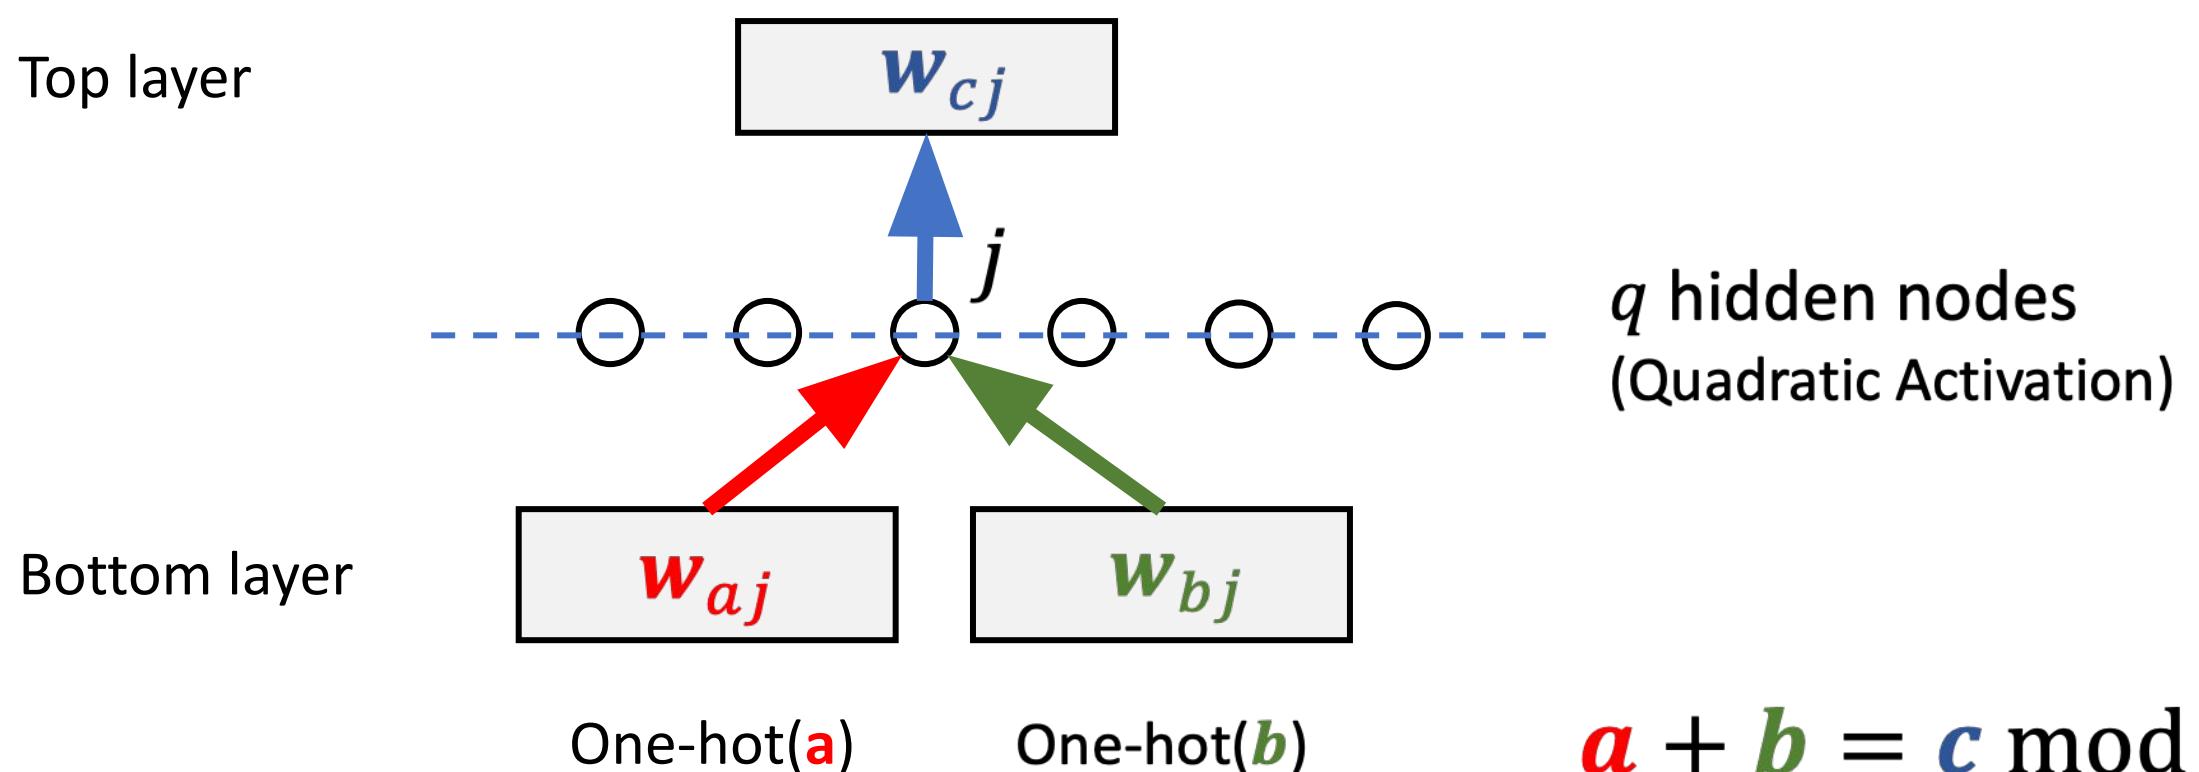
Why? 



(a) Final logits for top Fourier components

Problem Setup

MSE Loss: $\text{Min} \ \| \text{Output} - \text{one-hot}(\mathbf{c}) \|_2$



(Scaled) Fourier Transform

k : frequency

$$z_{akj} = \sum_{m=0}^{d-1} w_{amj} e^{\text{i}mk/d}$$

$\{W_a, W_b, W_c\}$ are real

$$z_{bkj} = \sum_{m=0}^{d-1} w_{bmj} e^{\text{i}mk/d}$$



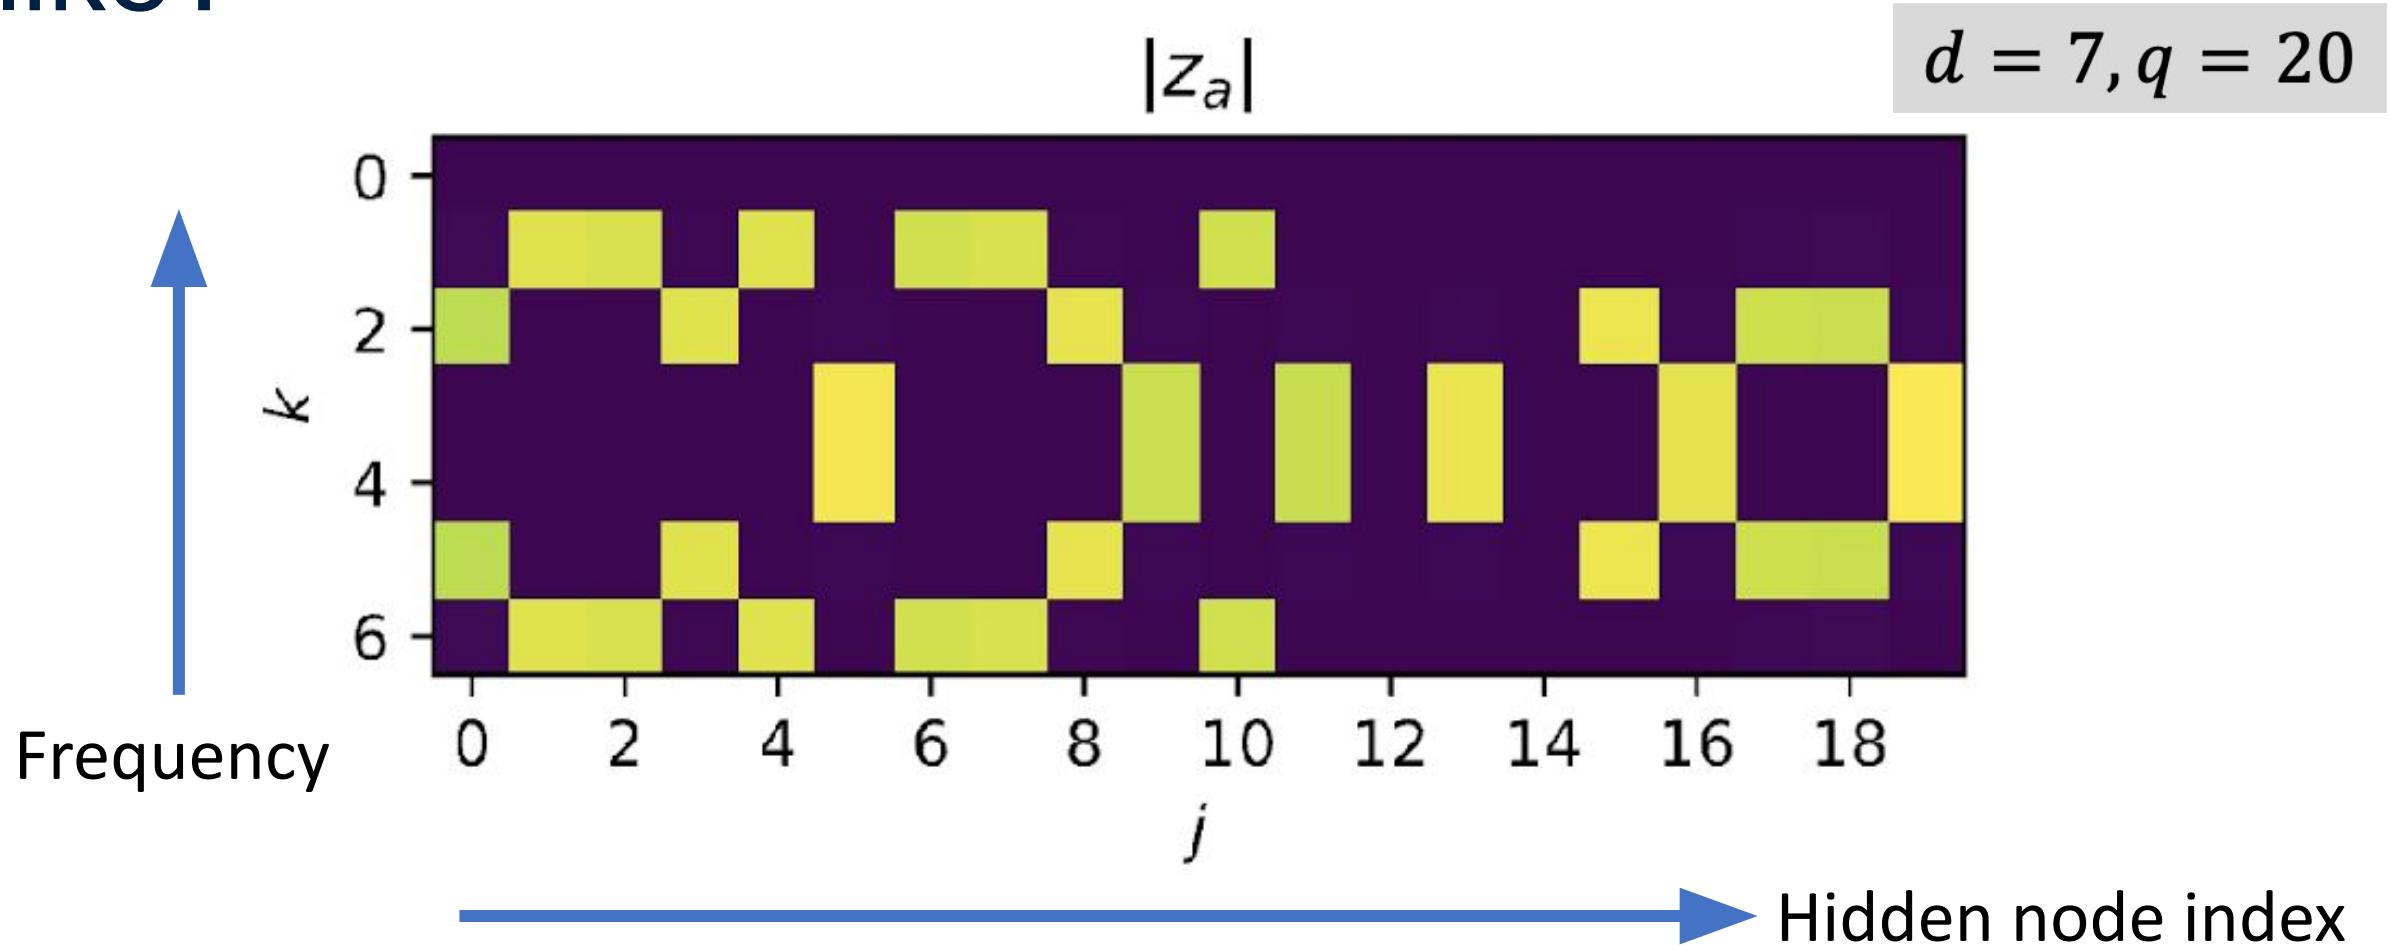
Hermitian condition holds

$$z_{akj} = \overline{z_{a,-k,j}}$$

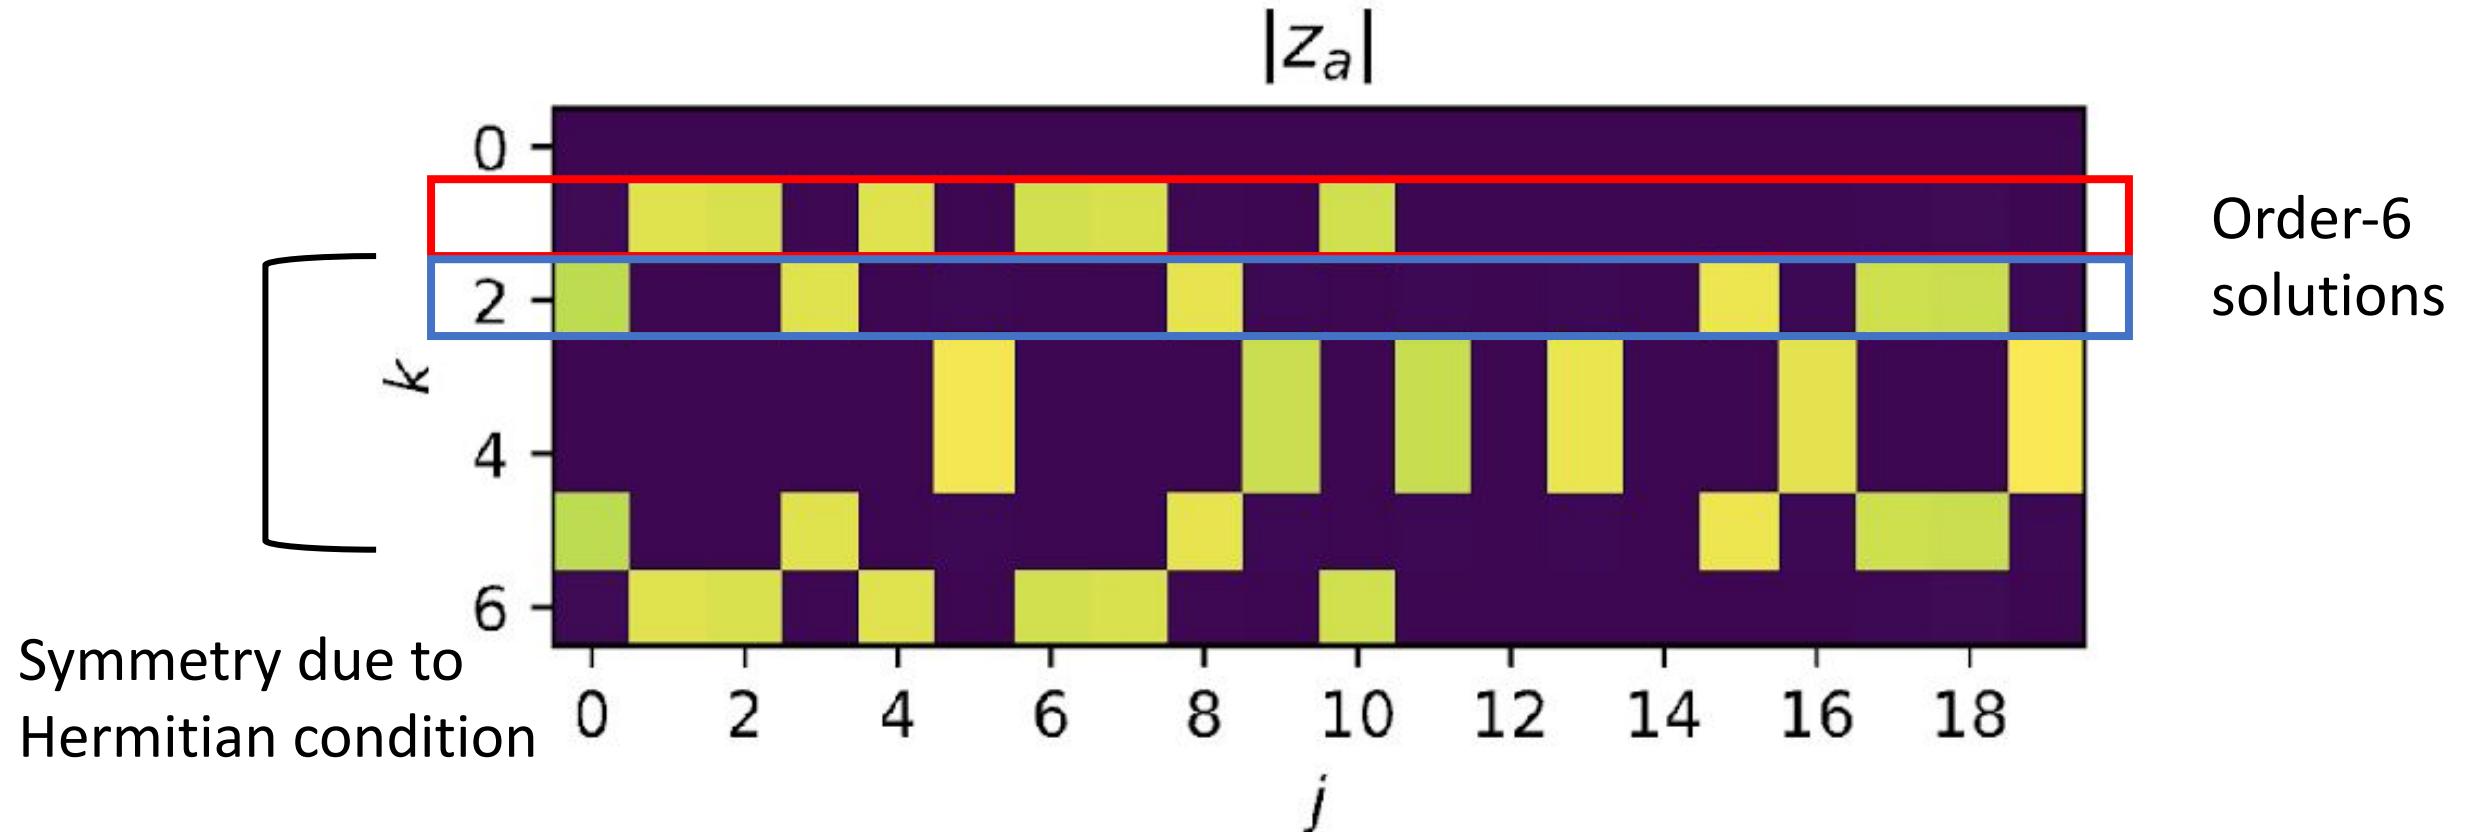
$$z_{bkj} = \overline{z_{b,-k,j}}$$

$$z_{ckj} = \overline{z_{c,-k,j}}$$

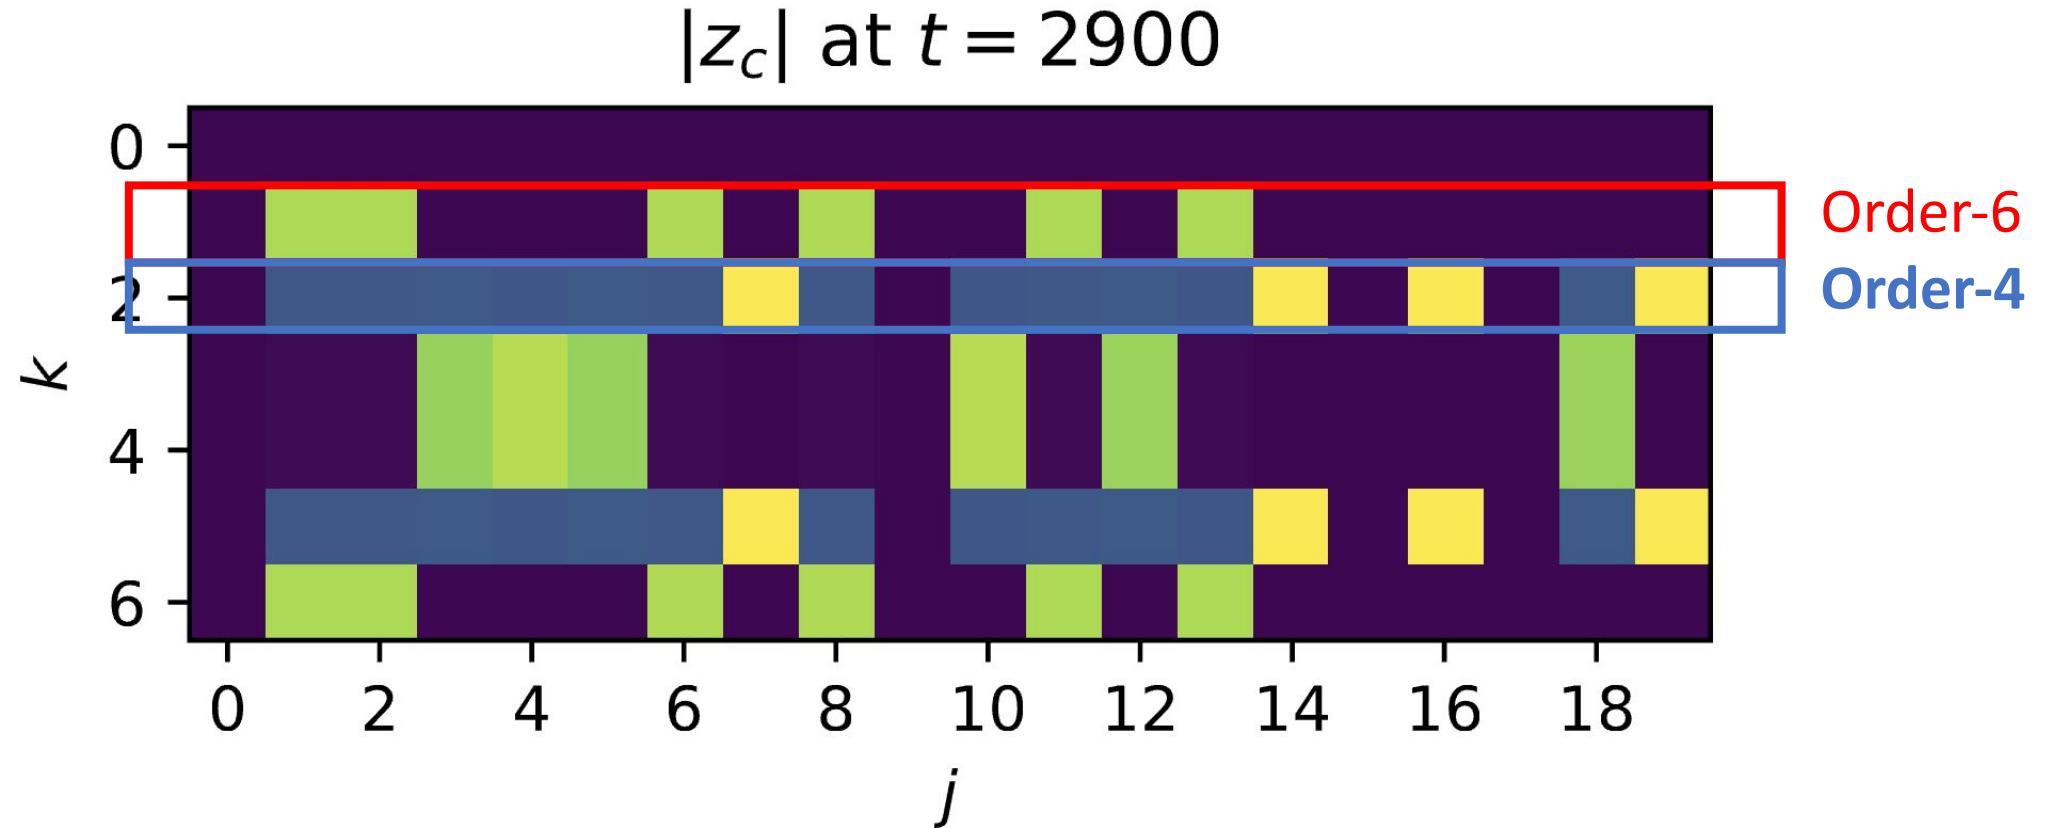
What a Gradient Descent Solution look like?



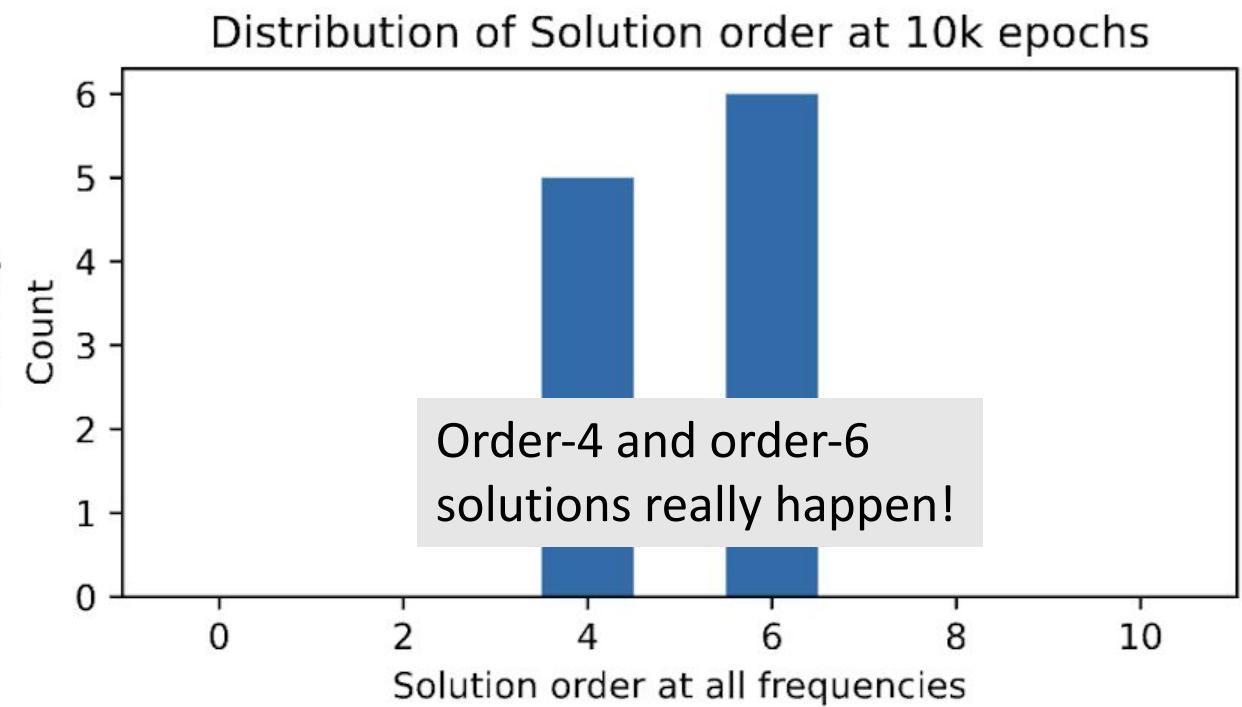
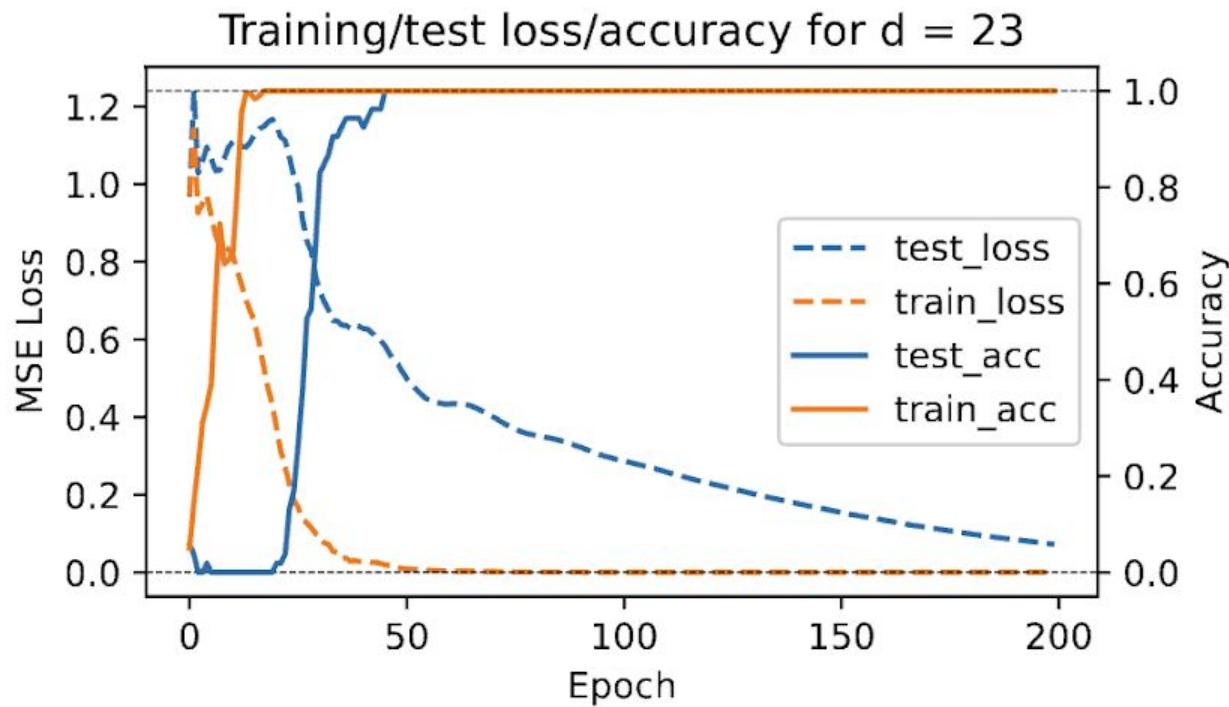
What a Gradient Descent Solution look like?



What a Gradient Descent Solution look like?

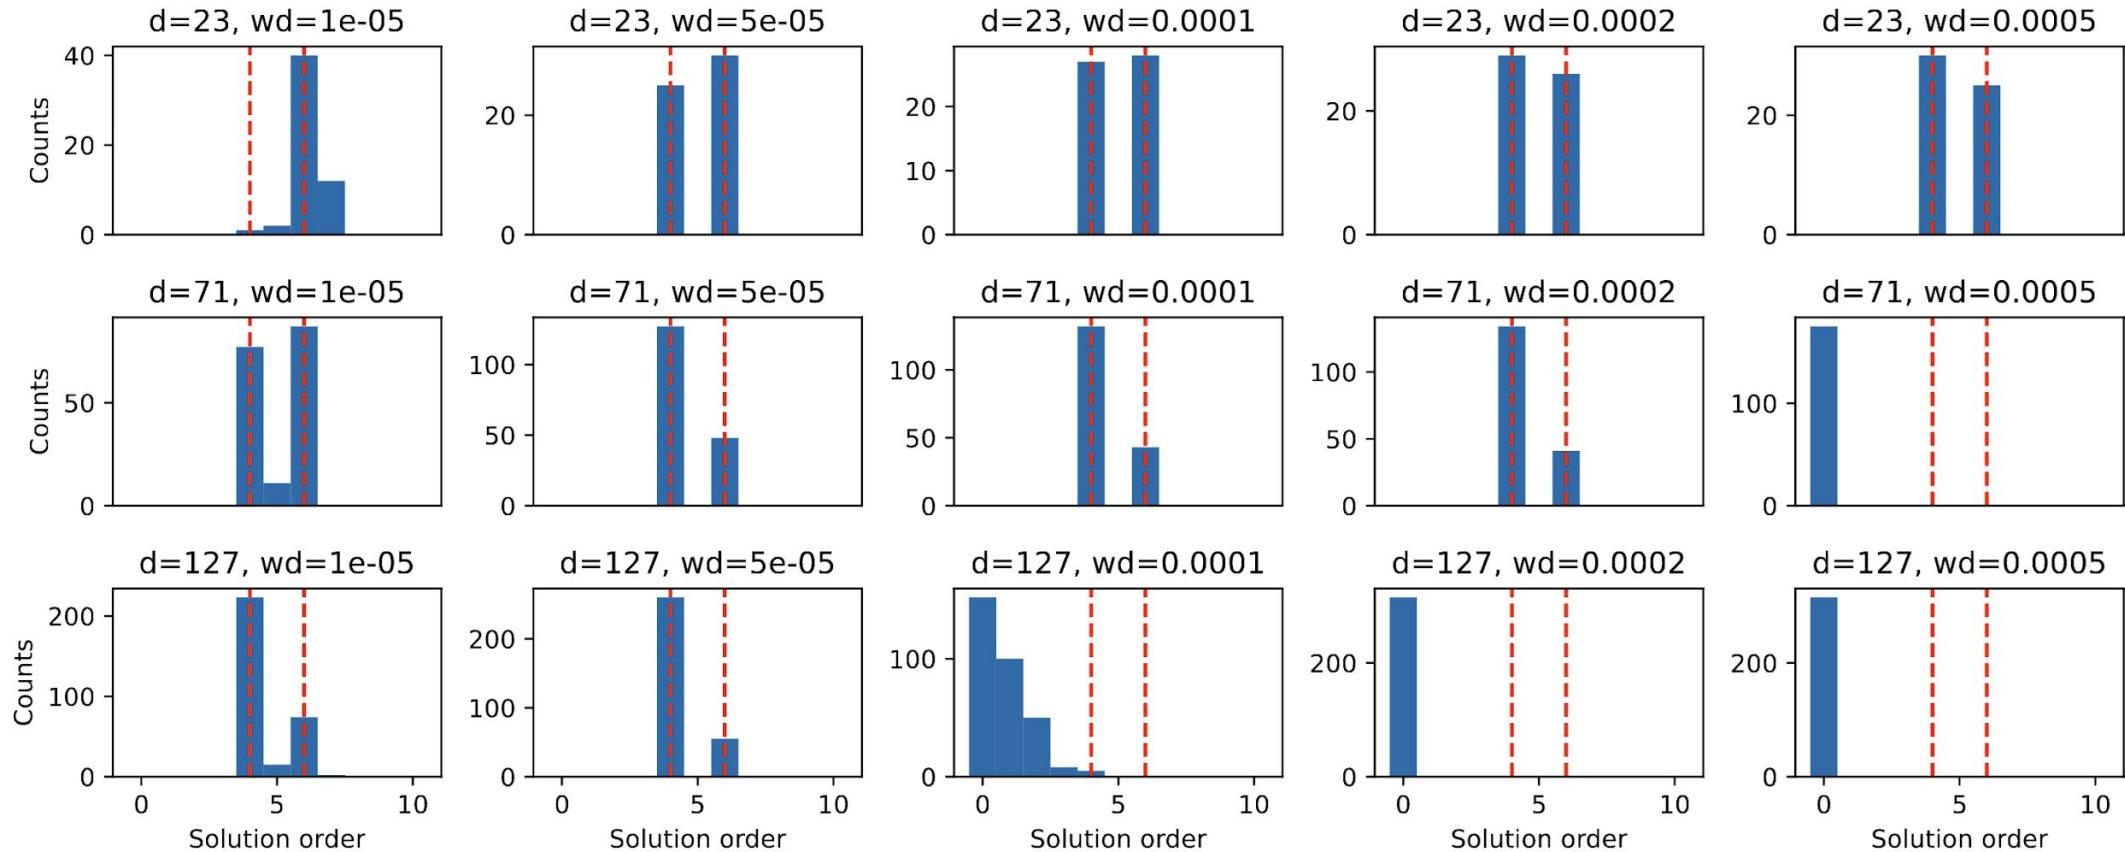


More Statistics on Gradient Descent Solutions



Effect of Weight Decay

Stronger weight decay →



Why?



Structure of Loss Functions

$$\text{MSE loss } \ell(\mathbf{z}) = d^{-1} \sum_{k \neq 0} \ell_k(\mathbf{z}) + 1 - 1/d$$

$$\ell_k(\mathbf{z}) = -2\mathbf{r}_{kkk} + \sum_{k_1 k_2} |\mathbf{r}_{k_1 k_2 k}|^2 + \frac{1}{4} \left| \sum_{p \in \{a,b\}} \sum_{k'} \mathbf{r}_{p,k',-k',k} \right|^2 + \frac{1}{4} \sum_{m \neq 0} \sum_{p \in \{a,b\}} \left| \sum_{k'} \mathbf{r}_{p,k',m-k',k} \right|^2$$

Term $\mathbf{r}_{k_1 k_2 k}(\mathbf{z}) := \sum_j z_{ak_1 j} z_{bk_2 j} z_{ckj}$ and $\mathbf{r}_{pk_1 k_2 k}(\mathbf{z}) := \sum_j z_{pk_1 j} z_{pk_2 j} z_{ckj}$

Structure of Loss Functions

MSE loss $\ell(\mathbf{z}) = d^{-1} \sum_{k \neq 0} \ell_k(\mathbf{z}) + 1 - 1/d$

$$\ell_k(\mathbf{z}) = -2\mathbf{r}_{kkk} + \sum_{k_1 k_2} |\mathbf{r}_{k_1 k_2 k}|^2 + \frac{1}{4} \left| \sum_{p \in \{a,b\}} \sum_{k'} \mathbf{r}_{p,k',-k',k} \right|^2 + \frac{1}{4} \sum_{m \neq 0} \sum_{p \in \{a,b\}} \left| \sum_{k'} \mathbf{r}_{p,k',m-k',k} \right|^2$$

Term $\mathbf{r}_{k_1 k_2 k}(\mathbf{z}) := \sum_j z_{ak_1 j} z_{bk_2 j} z_{ckj}$ and $\mathbf{r}_{pk_1 k_2 k}(\mathbf{z}) := \sum_j z_{pk_1 j} z_{pk_2 j} z_{ckj}$

Sufficient conditions of Global Optimizers:

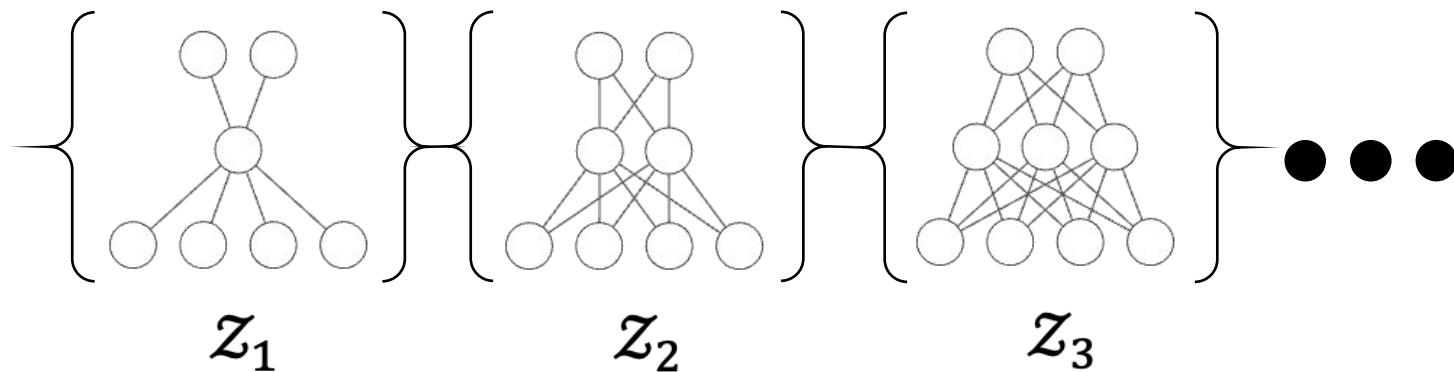
How to Optimize?

The objective is highly nonlinear !!
However, nice *algebraic structures* exist!

How to Optimize?

The objective is highly nonlinear !!

However, nice *algebraic structures* exist!

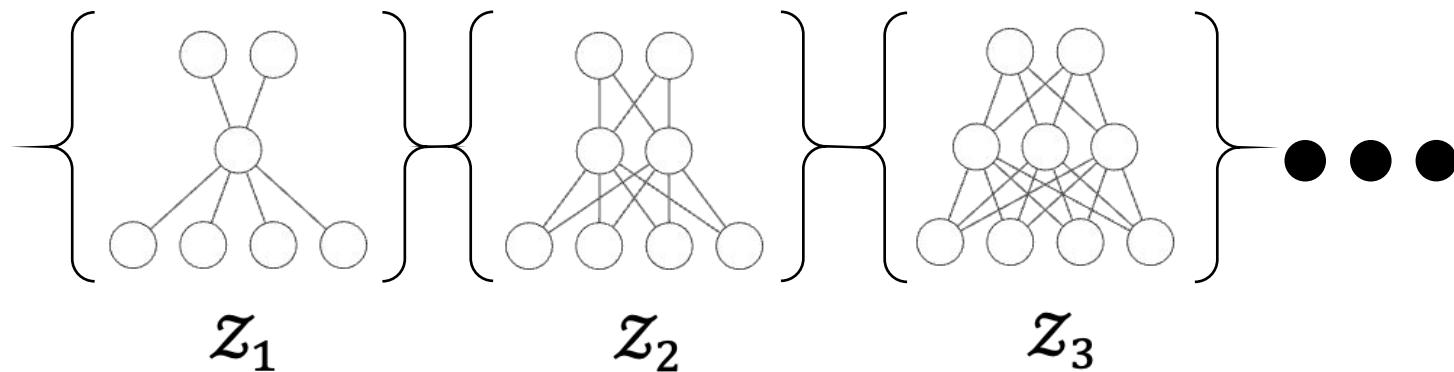


$\mathcal{Z} = \bigcup_{q \geq 0} \mathcal{Z}_q$: All 2-layer networks with different number of hidden nodes

How to Optimize?

The objective is highly nonlinear !!

However, nice *algebraic structures* exist!



$\mathcal{Z} = \bigcup_{q \geq 0} \mathcal{Z}_q$: All 2-layer networks with different number of hidden nodes

Ring addition $+$: Concatenate hidden nodes

Ring multiplication $*$: Kronecker production along the hidden dimensions

$\langle \mathcal{Z}, +, *\rangle$ is a *semi-ring*

Ring Homomorphism

A function $r(\mathbf{z}): \mathcal{Z} \mapsto \mathbb{C}$ is a *ring homomorphism*, if

- $r(\mathbf{1}) = 1$
- $r(\mathbf{z}_1 + \mathbf{z}_2) = r(\mathbf{z}_1) + r(\mathbf{z}_2)$
- $r(\mathbf{z}_1 * \mathbf{z}_2) = r(\mathbf{z}_1)r(\mathbf{z}_2)$

Ring Homomorphism

A function $r(z): \mathcal{Z} \mapsto \mathbb{C}$ is a *ring homomorphism*, if

- $r(1) = 1$
- $r(z_1 + z_2) = r(z_1) + r(z_2)$
- $r(z_1 * z_2) = r(z_1)r(z_2)$



$r_{k_1 k_2 k}(z)$ and $r_{pk_1 k_2 k}(z)$ are **ring homomorphisms!**

Ring Homomorphism

A function $r(\mathbf{z}): \mathcal{Z} \mapsto \mathbb{C}$ is a *ring homomorphism*, if

- $r(\mathbf{1}) = 1$
- $r(\mathbf{z}_1 + \mathbf{z}_2) = r(\mathbf{z}_1) + r(\mathbf{z}_2)$
- $r(\mathbf{z}_1 * \mathbf{z}_2) = r(\mathbf{z}_1)r(\mathbf{z}_2)$

🎯 $r_{k_1 k_2 k}(\mathbf{z})$ and $r_{pk_1 k_2 k}(\mathbf{z})$ are ring homomorphisms!

MSE Loss

$$\ell_k(\mathbf{z}) = -2\mathbf{r}_{\mathbf{k}\mathbf{k}\mathbf{k}} + \sum_{k_1 k_2} |\mathbf{r}_{\mathbf{k}_1 \mathbf{k}_2 \mathbf{k}}|^2 + \frac{1}{4} \left| \sum_{p \in \{a,b\}} \sum_{k'} \mathbf{r}_{p, k', -k', k} \right|^2 + \frac{1}{4} \sum_{m \neq 0} \sum_{p \in \{a,b\}} \left| \sum_{k'} \mathbf{r}_{p, k', m - k', k} \right|^2$$

Ring Homomorphism

A function $r(\mathbf{z}): \mathcal{Z} \mapsto \mathbb{C}$ is a *ring homomorphism*, if

- $r(\mathbf{1}) = 1$
- $r(\mathbf{z}_1 + \mathbf{z}_2) = r(\mathbf{z}_1) + r(\mathbf{z}_2)$
- $r(\mathbf{z}_1 * \mathbf{z}_2) = r(\mathbf{z}_1)r(\mathbf{z}_2)$

🎯 $r_{k_1 k_2 k}(\mathbf{z})$ and $r_{pk_1 k_2 k}(\mathbf{z})$ are ring homomorphisms!

MSE Loss

$$\ell_k(\mathbf{z}) = -2\mathbf{r}_{\mathbf{k}\mathbf{k}\mathbf{k}} + \sum_{k_1 k_2} |\mathbf{r}_{\mathbf{k}_1 \mathbf{k}_2 \mathbf{k}}|^2 + \frac{1}{4} \left| \sum_{p \in \{a,b\}} \sum_{k'} \mathbf{r}_{p, k', -k', k} \right|^2 + \frac{1}{4} \sum_{m \neq 0} \sum_{p \in \{a,b\}} \left| \sum_{k'} \mathbf{r}_{p, k', m - k', k} \right|^2$$

Partial solution \mathbf{z}_1 satisfies $\mathbf{r}_{k_1 k_2 k}(\mathbf{z}_1) = 0$

Partial solution \mathbf{z}_2 satisfies $\mathbf{r}_{pk', -k', k}(\mathbf{z}_2) = 0$

Ring Homomorphism

A function $r(\mathbf{z}): \mathcal{Z} \mapsto \mathbb{C}$ is a *ring homomorphism*, if

- $r(\mathbf{1}) = 1$
- $r(\mathbf{z}_1 + \mathbf{z}_2) = r(\mathbf{z}_1) + r(\mathbf{z}_2)$
- $r(\mathbf{z}_1 * \mathbf{z}_2) = r(\mathbf{z}_1)r(\mathbf{z}_2)$

🎯 $r_{k_1 k_2 k}(\mathbf{z})$ and $r_{pk_1 k_2 k}(\mathbf{z})$ are ring homomorphisms!

MSE Loss

$$\ell_k(\mathbf{z}) = -2\mathbf{r}_{kkk} + \sum_{k_1 k_2} |\mathbf{r}_{k_1 k_2 k}|^2 + \frac{1}{4} \left| \sum_{p \in \{a,b\}} \sum_{k'} \mathbf{r}_{p,k',-k',k} \right|^2 + \frac{1}{4} \sum_{m \neq 0} \sum_{p \in \{a,b\}} \left| \sum_{k'} \mathbf{r}_{p,k',m-k',k} \right|^2$$

Partial solution \mathbf{z}_1 satisfies $\mathbf{r}_{k_1 k_2 k}(\mathbf{z}_1) = 0$

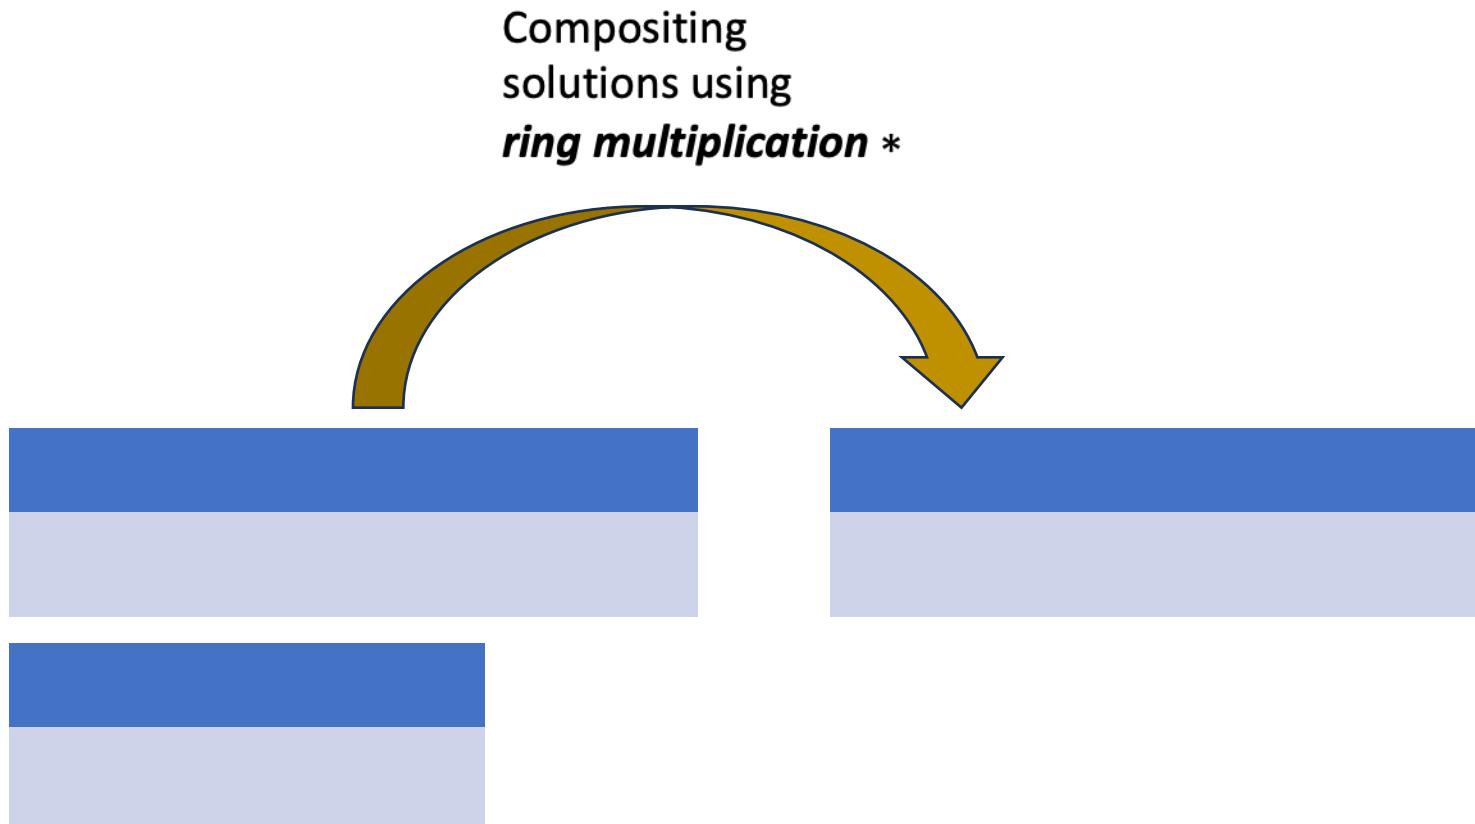
Partial solution \mathbf{z}_2 satisfies $\mathbf{r}_{pk',-k',k}(\mathbf{z}_2) = 0$

$\left. \begin{array}{l} \text{Partial solution } \mathbf{z}_1 \text{ satisfies } \mathbf{r}_{k_1 k_2 k}(\mathbf{z}_1) = 0 \\ \text{Partial solution } \mathbf{z}_2 \text{ satisfies } \mathbf{r}_{pk',-k',k}(\mathbf{z}_2) = 0 \end{array} \right\} \mathbf{z} = \mathbf{z}_1 * \mathbf{z}_2 \text{ satisfies both } \mathbf{r}_{k_1 k_2 k}(\mathbf{z}) = \mathbf{r}_{pk',-k',k}(\mathbf{z}) = 0$

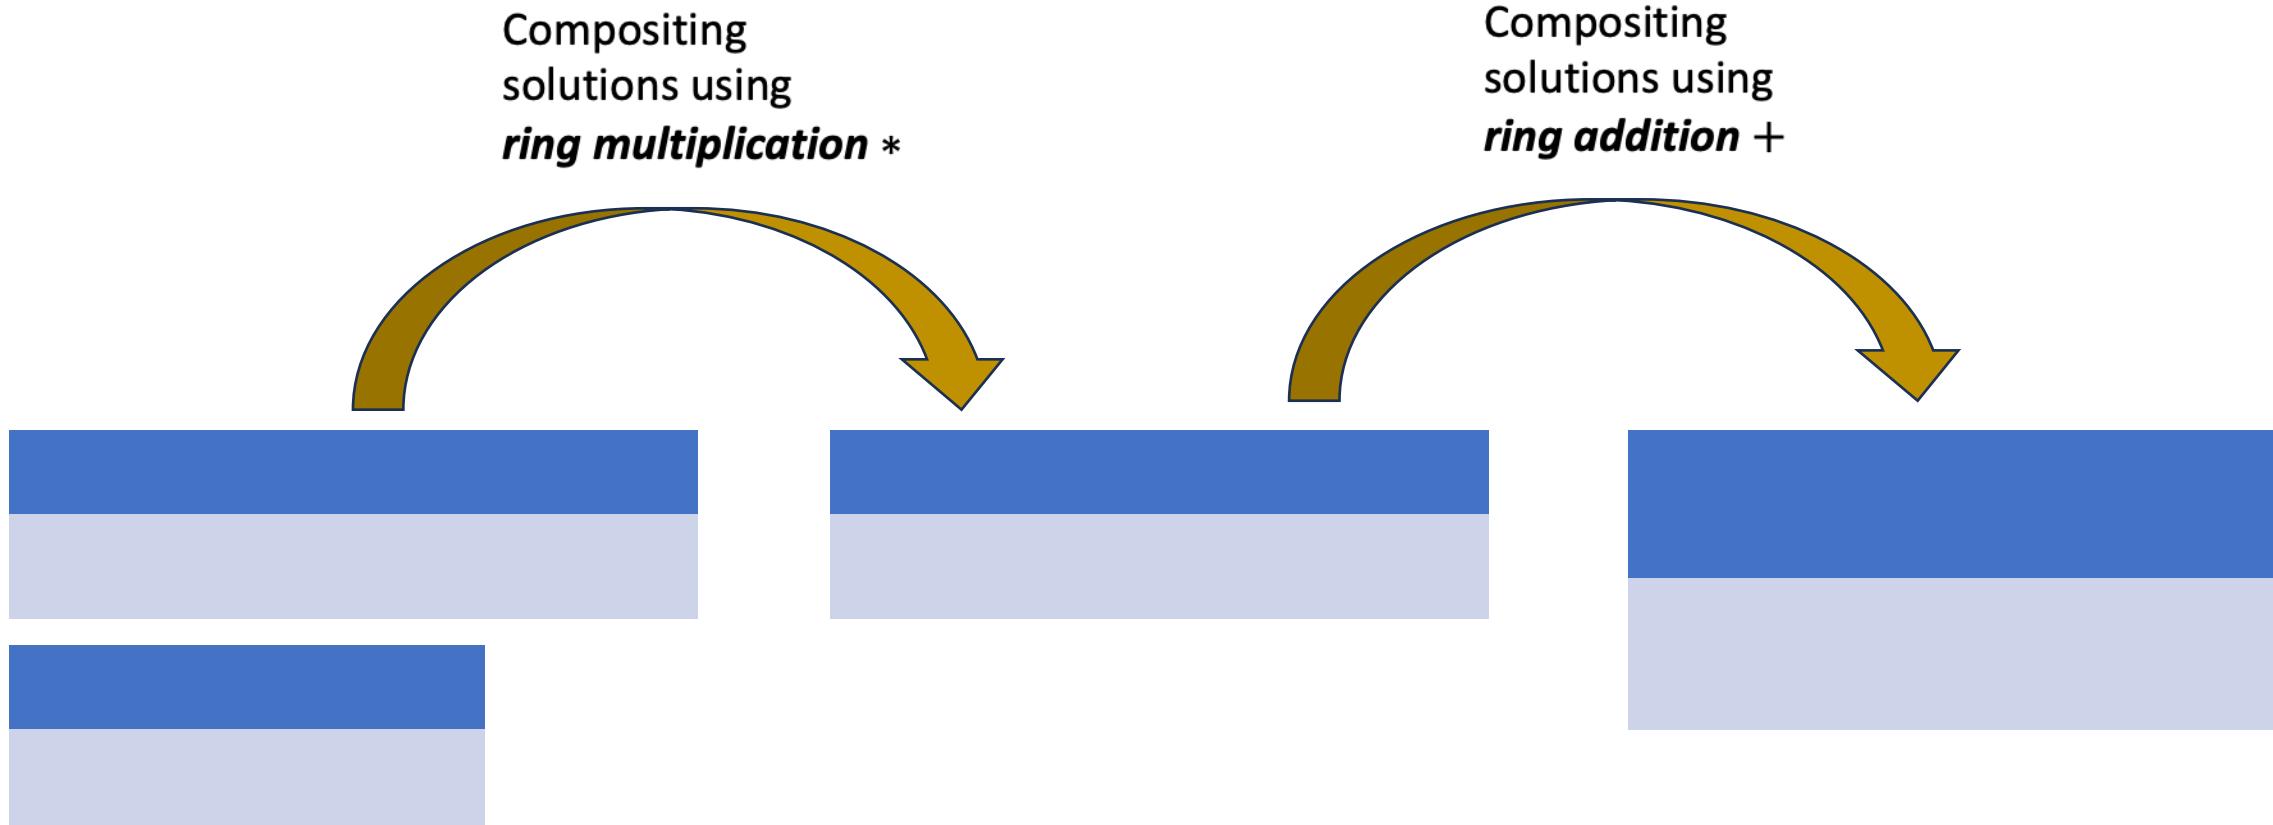
Composing Global Optimizers from Partial Ones



Composing Global Optimizers from Partial Ones



Composing Global Optimizers from Partial Ones



Exemplar constructed global optimizers

Order-6 \mathbf{z}_{F6} (2*3)

$$\mathbf{z}_{F6} = \frac{1}{\sqrt[3]{6}} \sum_{k=1}^{(d-1)/2} \mathbf{z}_{\text{syn}}^{(k)} * \mathbf{z}_{\nu}^{(k)} * \mathbf{y}_k$$

Exemplar constructed global optimizers

Order-6 \mathbf{z}_{F6} (2*3)

$$\mathbf{z}_{F6} = \frac{1}{\sqrt[3]{6}} \sum_{k=1}^{(d-1)/2} \mathbf{z}_{\text{syn}}^{(k)} * \mathbf{z}_{\nu}^{(k)} * \mathbf{y}_k$$

Order-4 $\mathbf{z}_{F4/6}$ (2*2)
(mixed with order-6)

$$\mathbf{z}_{F4/6} = \frac{1}{\sqrt[3]{6}} \hat{\mathbf{z}}_{F6}^{(k_0)} + \frac{1}{\sqrt[3]{4}} \sum_{k=1, k \neq k_0}^{(d-1)/2} \mathbf{z}_{F4}^{(k)}$$

Exemplar constructed global optimizers

Order-6 \mathbf{z}_{F6} (2^*3)

$$\mathbf{z}_{F6} = \frac{1}{\sqrt[3]{6}} \sum_{k=1}^{(d-1)/2} \mathbf{z}_{\text{syn}}^{(k)} * \mathbf{z}_\nu^{(k)} * \mathbf{y}_k$$

Order-4 $\mathbf{z}_{F4/6}$ (2^*2)
(mixed with order-6)

$$\mathbf{z}_{F4/6} = \frac{1}{\sqrt[3]{6}} \hat{\mathbf{z}}_{F6}^{(k_0)} + \frac{1}{\sqrt[3]{4}} \sum_{k=1, k \neq k_0}^{(d-1)/2} \mathbf{z}_{F4}^{(k)}$$

Perfect memorization
(order-d per frequency)

$$\mathbf{z}_a = \sum_{j=0}^{d-1} \mathbf{u}_a^j, \quad \mathbf{z}_b = \sum_{j=0}^{d-1} \mathbf{u}_b^j$$

$$\mathbf{z}_M = d^{-2/3} \mathbf{z}_a * \mathbf{z}_b$$

Gradient Descent solutions matches with construction

d	%not order-4/6	%non-factorable order-4	%non-factorable order-6	error ($\times 10^{-2}$)		solution distribution (%) in factorable ones				others
	order-4	order-6	order-6	order-4	order-6	$z_{\nu=i}^{(k)} * z_{\xi}^{(k)}$	$z_{\nu=i}^{(k)} * z_{\text{syn}, \alpha\beta}^{(k)}$	$z_{\nu}^{(k)} * z_{\text{syn}}^{(k)}$		
23	0.0±0.0	0.00±0.00	5.71±5.71	0.05±0.01	4.80±0.96	47.07±1.88	11.31±1.76	39.80±2.11	1.82±1.82	
71	0.0±0.0	0.00±0.00	0.00±0.00	0.03±0.00	5.02±0.25	72.57±0.70	4.00±1.14	21.14±2.14	2.29±1.07	
127	0.0±0.0	1.50±0.92	0.00±0.00	0.26±0.14	0.93±0.18	82.96±0.39	2.25±0.64	14.13±0.87	0.66±0.66	

$$q = 512, wd = 5 \cdot 10^{-5}$$

Gradient Descent solutions matches with construction

d	%not order-4/6	%non-factorable		error ($\times 10^{-2}$)		solution distribution (%) in factorable ones				others
	order-4	order-6	order-4	order-6	$z_{\nu=i}^{(k)} * z_{\xi}^{(k)}$	$z_{\nu=i}^{(k)} * z_{\text{syn}, \alpha\beta}^{(k)}$	$z_{\nu}^{(k)} * z_{\text{syn}}^{(k)}$	others		
23	0.0±0.0	0.00±0.00	5.71±5.71	0.05±0.01	4.80±0.96	47.07±1.88	11.31±1.76	39.80±2.11	1.82±1.82	
71	0.0±0.0	0.00±0.00	0.00±0.00	0.03±0.00	5.02±0.25	72.57±0.70	4.00±1.14	21.14±2.14	2.29±1.07	
127	0.0±0.0	1.50±0.92	0.00±0.00	0.26±0.14	0.93±0.18	82.96±0.39	2.25±0.64	14.13±0.87	0.66±0.66	

100% of the per-freq
solutions are order-4/6

Gradient Descent solutions matches with construction

d	%not order-4/6			%non-factorable		error ($\times 10^{-2}$)		solution distribution (%) in factorable ones			
	order-4	order-6	order-4	order-4	order-6	$z_{\nu=i}^{(k)} * z_{\xi}^{(k)}$	$z_{\nu=i}^{(k)} * z_{\text{syn}, \alpha\beta}^{(k)}$	$z_{\nu}^{(k)} * z_{\text{syn}}^{(k)}$	others		
23	0.0±0.0	0.00±0.00	5.71±5.71	0.05±0.01	4.80±0.96	47.07±1.88	11.31±1.76	39.80±2.11	1.82±1.82		
71	0.0±0.0	0.00±0.00	0.00±0.00	0.03±0.00	5.02±0.25	72.57±0.70	4.00±1.14	21.14±2.14	2.29±1.07		
127	0.0±0.0	1.50±0.92	0.00±0.00	0.26±0.14	0.93±0.18	82.96±0.39	2.25±0.64	14.13±0.87	0.66±0.66		

95% of the solutions are
factorizable into “2*3” or “2*2”

Gradient Descent solutions matches with construction

d	%not order-4/6	%non-factorable order-4	%non-factorable order-6	error ($\times 10^{-2}$)		solution distribution (%) in factorable ones			
	order-4	order-6	order-6	order-4	order-6	$z_{\nu=i}^{(k)} * z_{\xi}^{(k)}$	$z_{\nu=i}^{(k)} * z_{\text{syn}, \alpha\beta}^{(k)}$	$z_{\nu}^{(k)} * z_{\text{syn}}^{(k)}$	others
23	0.0±0.0	0.00±0.00	5.71±5.71	0.05±0.01	4.80±0.96	47.07±1.88	11.31±1.76	39.80±2.11	1.82±1.82
71	0.0±0.0	0.00±0.00	0.00±0.00	0.03±0.00	5.02±0.25	72.57±0.70	4.00±1.14	21.14±2.14	2.29±1.07
127	0.0±0.0	1.50±0.92	0.00±0.00	0.26±0.14	0.93±0.18	82.96±0.39	2.25±0.64	14.13±0.87	0.66±0.66

Factorization error is very small

Gradient Descent solutions matches with construction

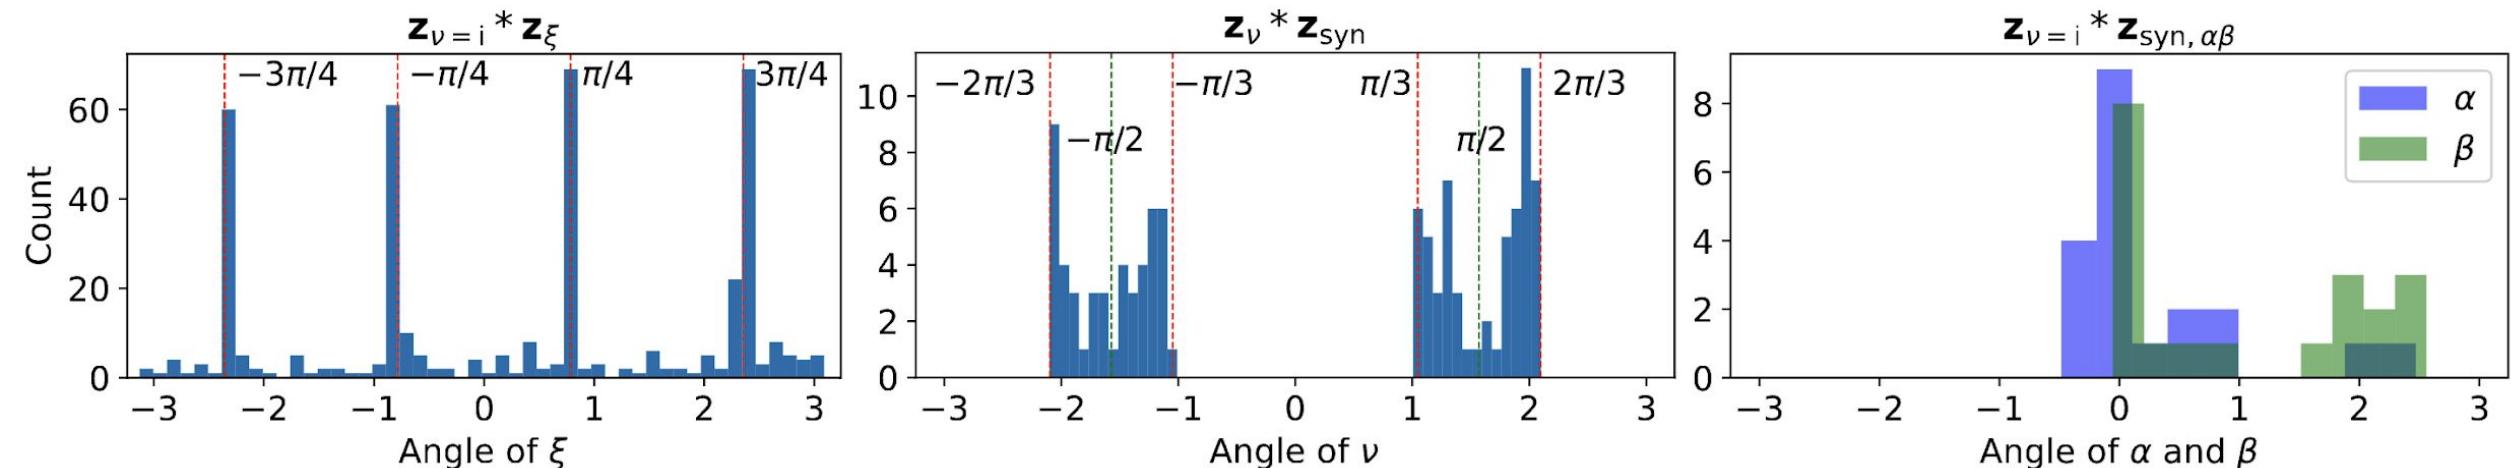
d	%not order-4/6			%non-factorable order-4			error ($\times 10^{-2}$)		solution distribution (%) in factorable ones			
	order-4	order-6	order-4	order-4	order-6	order-4	$z_{\nu=i}^{(k)} * z_{\xi}^{(k)}$	$z_{\nu=i}^{(k)} * z_{\text{syn}, \alpha\beta}^{(k)}$	$z_{\nu}^{(k)} * z_{\text{syn}}^{(k)}$	others		
23	0.0±0.0	0.00±0.00	5.71±5.71	0.05±0.01	4.80±0.96	47.07±1.88	11.31±1.76	39.80±2.11	1.82±1.82			
71	0.0±0.0	0.00±0.00	0.00±0.00	0.03±0.00	5.02±0.25	72.57±0.70	4.00±1.14	21.14±2.14	2.29±1.07			
127	0.0±0.0	1.50±0.92	0.00±0.00	0.26±0.14	0.93±0.18	82.96±0.39	2.25±0.64	14.13±0.87	0.66±0.66			

98% of the solutions can be factorizable into the constructed forms

Gradient Descent solutions matches with construction

d	%not order-4/6	%non-factorable order-4	%non-factorable order-6	error ($\times 10^{-2}$)		solution distribution (%) in factorable ones			
	order-4	order-6	order-6	order-4	order-6	$z_{\nu=i}^{(k)} * z_{\xi}^{(k)}$	$z_{\nu=i}^{(k)} * z_{\text{syn},\alpha\beta}^{(k)}$	$z_{\nu}^{(k)} * z_{\text{syn}}^{(k)}$	others
23	0.0±0.0	0.00±0.00	5.71±5.71	0.05±0.01	4.80±0.96	47.07±1.88	11.31±1.76	39.80±2.11	1.82±1.82
	5	72.57±0.70	4.00±1.14	21.14±2.14	2.29±1.07				
	8	82.96±0.39	2.25±0.64	14.13±0.87	0.66±0.66				

Distribution of the parameters in the solutions

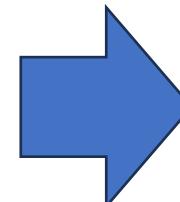


Possible Implications

Do neural networks end up learning more efficient symbolic representations that we don't know?

Does gradient descent lead to a solution that can be reached by advanced algebraic operations?

Will gradient descent become obsolete, eventually?



Thanks!